



DEPARTMENT OF MASTER OF COMPUTER APPLICATIONS

LABORATORY MANUAL

III Semester

Batch:2024-26

Name: Goli Shiva Dhanush

USN: 1MS24MC031

Course: Cloud Computing

Course code: 24MCASS3

Course Credits: 0 : 1: 2

RAMAIAH INSTITUTE OF TECHNOLOGY

(Autonomous Institute, Affiliated to VTU)

Accredited by National Board of Accreditation & NAAC with 'A+' Grade, MSR

Nagar, MSRIT Post, Bangalore-560054

www.msrit.edu

Table of Lab Programs

Sl. No.	Programs/Exercise/Topic.
1.	AWS Account Setup and Configuration. AWS Console Overview. Enable MFA for root account. Create AWS budget alert .
2.	AWS Identity Access Management (IAM) User and Group creation. Enable AWS IAM MFA. Create an AWS Account Alias (for Alternate Sign-in URL)
3.	Amazon S3 – Introduction Bucket Creation and upload objects (files).
4.	Amazon S3 – Static Website Hosting (Multi-Page website) Enable Versioning . Cross-Region Replication rule .
5.	Overview of EC2. To Launch a Windows EC2 Instance and Connect via RDP Client .
6.	Launch a Linux EC2 Instance and Connect using SSH through PowerShell on Windows machine. [pem file] Launch a Linux EC2 Instance and Connect using SSH through PuTTY on Windows machine. [ppk file] Launch a Linux EC2 Instance and Connect using SSH through Linux terminal on Linux machine. [pem file] Launch a Linux EC2 Instance and Connect through EC2 connect option on the console.
7.	Hosting a static website on EC2 instance: - Manual Installation of Apache (httpd) Web Server on EC2 for Static Website Hosting. - Launching an EC2 Instance with User Data Script to Automatically Install Apache and Host a Static Website.
8.	Custom AMI - Create a Custom AMI from a Working EC2 Instance. - Launch an EC2 Instance using a Custom AMI. - Delete the custom AMI [Deregister and delete snapshot]
9.	Mini-Project: Hosting a Multi-Page Website using EC2 and S3 Submission: 11-11-2025, Marks: 5
10.	Creation and Configuration of a Custom AWS Virtual Private Cloud (VPC) . [Including Public and Private Subnets, Internet Gateway, NAT Gateway, Route Tables]
11.	Launching and Configuring EC2 Instances for Web Server in Public Subnet and Database Server in Private Subnet Using NAT Gateway for Outbound Access.
12.	Deploying and Testing a Load-Balanced Web Application By Creating a Web Server, Building a Custom AMI, Configuring a Target Group, Launching an Application Load Balancer (ALB), Registering the Instance in the Target Group, and Testing the ALB DNS.

13.	Stress Testing a Linux EC2 Instance (CPU Load Test)
14.	Mini Project: Deploying a Load-Balanced Web Application using Application Load Balancer (ALB), Auto Scaling Group (ASG), Custom AMI, and Target Group on AWS with CloudWatch Alarms & Stress Testing for Auto Scaling Validation. Submission: 22-11-2025, Marks: 5
15.	Hosting a WordPress Content Management Website on Amazon Lightsail . (only demo)
16.	Building a Basic Python Flask Web Application — Fundamentals of Web Requests & Form Handling (Pre-Requisite for AWS RDS Connectivity Lab)
17.	Flask Web App with Registration & Login Using AWS RDS (MySQL) (Full Deployment on AWS EC2 + AWS RDS)
18.	Creating and Operating a NoSQL Key-Value Database using Amazon DynamoDB
19.	ElastiCache (Redis) as an In-Memory Cache
20.	Deploy a simple Flask application on AWS Elastic Beanstalk and verify it using the public URL.
21.	Deploy a Flask App on AWS Elastic Beanstalk and Perform CRUD on DynamoDB (Using AWS SDK/boto3)
22.	CloudFormation – Launch EC2 (Amazon Linux 2023) + Install Apache using UserData
23.	Static Content Pulled from S3 using CloudFormation [EC2 + S3]
24.	AWS Lambda : Input Processing, Business Logic Execution, and CloudWatch Logging.
25.	Event-Driven Notification System using AWS Lambda and Amazon SNS
26.	Analyze data in a CSV File stored in S3 Using Amazon Athena (No Server)
27.	Smart Sensor Monitoring System using AWS IoT Core . Cloud-Based IoT Data Ingestion and Processing using AWS
28.	Accelerate an S3 Static Website Using Amazon CloudFront (CDN)
29.	Mini Project – Global Static Website Delivery using S3 + CloudFront with Cache Update
30.	Mini Project – Deploying a Containerized App on AWS using Amazon EKS

Date: 8-10-2025

Exercise: AWS Account Setup and Configuration. AWS Console Overview. Enable MFA. Create AWS budget alert.

AWS Console overview / AWS Home page/ AWS Dash board

Widgets – default view/ Add or remove widgets - small panels on the dashboard showing metrics or shortcuts; users can add or remove them as needed.

Region – Specifies the geographical data center location where your AWS resources are deployed.

Services – A categorized list of all AWS offerings such as Compute, Storage, Database, etc.

Search bar – A quick-access bar to search and pin frequently used services for faster access.

pin the most used services to console by clicking on star next to the service name.

Enable MFA

Notes

- Make sure your phone is unlocked, Bluetooth is on, and it uses a screen lock (fingerprint/PIN).

Option 1: Add a Passkey for Easier Login

Step1: Go to Security Credentials

- **Sign in to AWS console.**
- Go to **your username** → **Security credentials.**
- Under **Multi-factor authentication (MFA)** click “**Assign MFA device.**”
- Choose “**Passkeys and security keys**” → **Next.**
- On the next screen choose “**Phone or tablet**”.
- AWS will show a **browser pop-up** asking to use a device.
 1. Select **your phone** (or “Use another device” if it prompts).
- Look at your phone — you should get a “**Use passkey**” or **biometric prompt.**
- Approve using **fingerprint or phone PIN.**
- Back in AWS, click **Finish.** The passkey is now your MFA method.

Next time you sign in, just choose “**Sign in with a passkey**” → **approve on phone.**

Step 2: Test the Login

- Sign out of AWS.
- Go to the login page.
- Choose “Sign in with a passkey” → Select your phone.
- Approve the prompt on your phone — you should be signed in without any MFA codes.

Option 2: Authenticator App

This uses a 6-digit code from Google Authenticator / Authy / Microsoft Authenticator.

1. In **Security credentials**, click “**Assign MFA device.**”
2. Select “**Authenticator app**” → **Next.**
3. A QR code appears.
4. Open your authenticator app on your phone → **Add account** → **Scan QR code.**
5. The app shows a 6-digit code.
6. Enter that code back in AWS → **Assign MFA.** You’ll use the 6-digit code from the app each time you log in.

The screenshot shows the AWS IAM console interface. The top navigation bar includes the AWS logo, a search bar, and the user's account information (Account ID: 6293-0494-1888, Shiva Dhanush). The left sidebar shows the 'IAM' menu with 'Security credentials' selected. The main content area is titled 'My security credentials' and includes a 'Root user' label. It contains three sections: 'Account details', 'Multi-factor authentication (MFA) (1)', and 'Access keys (0)'. The 'Account details' section shows the account name 'Shiva Dhanush', email address 'dhanushshiva2003@gmail.com', and AWS account ID '629304941888'. The 'Multi-factor authentication (MFA) (1)' section shows a single MFA device of type 'Virtual' with identifier 'arn:aws:iam::629304941888:mfa/SD-Device'. The 'Access keys (0)' section shows no access keys are currently configured.

My security credentials Root user Info

The root user has access to all AWS resources in this account, and we recommend following [best practices](#). To learn more about the types of AWS credentials and how they're used, see [AWS Security Credentials](#) in AWS General Reference.

Account details Actions

Account name
Shiva Dhanush

Email address
dhanushshiva2003@gmail.com

AWS account ID
 629304941888

Canonical user ID
 b7bdad4c31550749e4e804a5d1c9d28274c118e47c38df89eb6b9df0a62af657

Multi-factor authentication (MFA) (1) Remove Resync Assign MFA device

Use MFA to increase the security of your AWS environment. Signing in with MFA requires an authentication code from an MFA device. Each user can have a maximum of 8 MFA devices assigned. [Learn more](#)

Type	Identifier	Certifications	Created on
☐ Virtual	arn:aws:iam::629304941888:mfa/SD-Device	Not Applicable	Mon Oct 13 2025

Access keys (0) Create access key

Use access keys to send programmatic calls to AWS from the AWS CLI, AWS Tools for PowerShell, AWS SDKs, or direct AWS API calls. You can have a maximum of two access keys (active or inactive) at a time. [Learn more](#)

© 2025, Amazon Web Services, Inc. or its affiliates. [Privacy](#) [Terms](#) [Cookie preferences](#)

Create AWS budget alert

Allows to create a simple budget and to send alarms to registered email.

Example: if you are close to or exceeding your designated budget.


By setting a budget you can monitor budget threshold from the start.

Creating a budget

Step 1:

In the search bar type budgets and under the search results:

Select 'Budgets' from the Features group, which is essentially a feature of Billing and Cost Management service.



Step 2: After the 'Budgets' page loads

Click on **Create Budget** button

Under Budget setup select 'Customize' option

Under Budget types select 'Cost budget' option

Click on **Next** button.

The screenshot shows the AWS Management Console interface for creating a budget. The top navigation bar includes the AWS logo, a search bar, and account information (Account ID: 6293-0494-1888, Shiva Dhanush). The breadcrumb trail indicates the path: Billing and Cost Management > Budgets > Create budget. On the left, a sidebar menu lists various AWS services, with 'Billing and Cost Management' expanded. The main content area is titled 'Create budget' and features a progress indicator with five steps: 'Set your budget' (current), 'Configure alerts', 'Optional Attach actions', 'Review', and 'Review'. The 'Budget setup' section offers two options: 'Use a template (simplified)' and 'Customize (advanced)', with the latter selected. Below this, the 'Billing View' is shown as 'Based on selected billing view.' The 'Budget types' section lists four options: 'Cost budget - Recommended' (selected), 'Usage budget', 'Savings Plans budget', and 'Reservation budget'. The 'Cost budget - Recommended' option is highlighted with a blue border and a blue circle. The bottom of the page shows a footer with 'CloudShell', 'Feedback', and copyright information for Amazon Web Services, Inc. or its affiliates, along with links for 'Privacy', 'Terms', and 'Cookie preferences'.

aws [Search] [Alt+S] Global Account ID: 6293-0494-1888 Shiva Dhanush

Billing and Cost Management > Budgets > Create budget

Billing and Cost Management

Choose billing view [New](#)

Primary View

Home
Getting Started
Dashboards [New](#)

Billing and Payments

Bills
Payments
Credits
Purchase Orders

Cost and Usage Analysis

Cost Explorer
Cost Explorer Saved Reports
Cost Anomaly Detection
Free Tier
Data Exports
Customer Carbon Footprint Tool

Set your budget
Step 3
Configure alerts
Step 4 - Optional
Attach actions
Step 5
Review

Budget setup

☐ Use a template (simplified)
Use the recommended configurations. You can change some configuration options after the budget is created.

☒ **Customize (advanced)**
Customize a budget to set parameters specific to your use case. You can customize the time period, the start month, and specific accounts.

► **Billing View**
Based on selected billing view.

Budget types

☒ **Cost budget - Recommended**
Monitor your costs against a specified dollar amount and receive alerts when your user-defined thresholds are met. Using cost budgets, the budgeted amount you set represents your expected cloud spend. For example, you can set a cost budget for a business unit and then add additional parameters such as the associated member accounts.

☐ **Usage budget**
Monitor your usage of one or more specified usage types or usage type groups and receive alerts when your user-defined thresholds are met. Using usage budgets, the budgeted amount represents your expected usage. For example, you can use a usage budget to monitor the usage of certain services such as Amazon EC2 and Amazon S3.

☐ **Savings Plans budget**
Track the utilization or coverage associated with your Savings Plans and receive alerts when your percentage drops below a threshold you define. Setting a coverage target lets you see how much of your instance usage is covered by Savings Plans, while setting a utilization target lets you see if your Savings Plans are unused or underutilized.

☐ **Reservation budget**
Track the utilization or coverage associated with your reservations and receive alerts when your percentage drops below a

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 3: Set your budget page loads and fill in the details: MyBudget, Monthly, Recurring budget, set Month and Year, select Fixed – Budgeting method - Enter your budgeted amount – 20.00, select ‘All AWS services’ Scope options. Advanced options – leave it on default.

Click on **Next** button.

Billing and Cost Management

Choose billing view [New](#)

Primary View

Home

Getting Started

Dashboards [New](#)

Billing and Payments

Bills

Payments

Credits

Purchase Orders

Cost and Usage Analysis

Cost Explorer

Cost Explorer Saved Reports

Cost Anomaly Detection

Free Tier

Data Exports

Customer Carbon Footprint Tool

Details

Budget name

Provide a descriptive name for this budget.

MyBudget

Names must be between 1-100 characters.

Set budget amount

Period [Info](#)

Monthly

Budget renewal type

☒ **Recurring budget**
Recurring budgets renew on the first day of every monthly billing period.

☐ **Expiring budget**
Expiring monthly budgets stop renewing at the end of the selected expiration month.

Start month

Oct 2025

Budgeting method [Info](#)

Fixed
Create a budget that tracks against a single monthly budgeted amount.

Budget preview

Cost Data

Unable to load chart. An error has occurred with ce:GetCostAndUsage: DataUnavailableException - Cost Explorer was just enabled - it will take some time to prepare your cost and usage data. Please check back in 24 hours.

Alerts

No alerts configured.

© 2025, Amazon Web Services, Inc. or its affiliates. [Privacy](#) [Terms](#) [Cookie preferences](#)

Step 4: Configure alerts

Click on **Add an alert threshold**

Under Set alert threshold

Set Threshold: 80 and Trigger: Actual

Email recipients: enter your email id.

Click on **Next** button.

The screenshot displays the AWS Billing and Cost Management console interface. The left sidebar contains navigation links for Billing and Cost Management, including Home, Getting Started, Dashboards, Billing and Payments, and Cost and Usage Analysis. The main content area is titled 'Set alert threshold' and includes sections for Threshold, Trigger, Notification preferences, and Email recipients. The Threshold is set to 80 and the Trigger is set to Actual. A summary states: 'When your actual cost is greater than 80.00% (\$16.00) of your budgeted amount (\$20.00), the alert threshold will be exceeded.' The Email recipients section contains the email address dhanushshiva2003@gmail.com. A red error message is visible in the Budget preview section: 'Unable to load chart. An error has occurred with ce:GetCostAndUsage: DataUnavailableException - Cost Explorer was just enabled - it will take some time to prepare your cost and usage data. Please check back in 24 hours.' The bottom of the page shows the AWS footer with copyright information and links to Privacy, Terms, and Cookie preferences.

Billing and Cost Management

Choose billing view **New**

Primary View

Home

Getting Started

Dashboards **New**

Billing and Payments

Bills

Payments

Credits

Purchase Orders

Cost and Usage Analysis

Cost Explorer

Cost Explorer Saved Reports

Cost Anomaly Detection

Free Tier

Data Exports

Customer Carbon Footprint Tool

Set alert threshold

Threshold
When should this alert be triggered?
80 % of budgeted amount

Trigger
How should this alert be triggered?
Actual

Summary: When your actual cost is greater than 80.00% (\$16.00) of your budgeted amount (\$20.00), the alert threshold will be exceeded.

Notification preferences
Select one or more notification preferences to receive alerts.

Email recipients
Specify the email recipients you want to notify when the threshold has exceeded.
dhanushshiva2003@gmail.com
Maximum number of email recipients is 10.

► Amazon SNS Alerts - *Optional Info*

► AWS Chatbot Alerts

Budget preview

Cost Data

Unable to load chart. An error has occurred with ce:GetCostAndUsage: DataUnavailableException - Cost Explorer was just enabled - it will take some time to prepare your cost and usage data. Please check back in 24 hours.

Alerts

► Actual cost > 80% | No actions

© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Step 5: Under Attach actions – leave it on default

Click on **Next** button.

This screenshot shows the 'Attach actions' step (Step 4 of 5) in the AWS Budgets console. The left sidebar contains the 'Billing and Cost Management' menu. The main content area is titled 'Attach actions - Optional' and includes sections for 'Using budgets actions' and 'Alert #1 (0 actions attached)'. The 'Using budgets actions' section explains that a budget action allows you to define and trigger cost saving responses. The 'Alert #1' section shows a threshold of 80% measured against actual costs, with email recipients listed as dhanushshiva2003@gmail.com and Amazon SNS. The right sidebar shows the 'Budget preview' section, which displays 'Cost Data' and an error message: 'Unable to load chart. An error has occurred with ce:GetCostAndUsage: DataUnavailableException - Cost Explorer was just enabled - it will take some time to prepare your cost and usage data. Please check back in 24 hours.' The bottom of the page shows the AWS footer with copyright information and links to Privacy, Terms, and Cookie preferences.

Step 6: In review page – check and review all the options are set to what is desired.

Click on **Create budget** button.

Now, the budget has been created

This screenshot shows the 'Review' step (Step 5) in the AWS Budgets console. The left sidebar contains the 'Billing and Cost Management' menu. The main content area is titled 'Review' and includes sections for 'Step 1: Choose budget type' and 'Step 2: Set up your budget'. The 'Step 1: Choose budget type' section shows the 'Budget type' as 'Cost budget' with a description: 'Monitor your costs against a specified dollar amount and receive alerts when your user-defined thresholds are met.' The 'Step 2: Set up your budget' section shows the 'Budget details' table with the following data:

Budget details		
Name	Start date	Budget amount
MyBudget	Oct 2025	\$20.00
Period	End date	
Monthly	-	

The 'Additional budget parameters' section is empty. The 'Tags' section shows a table with columns 'Key' and 'Value'. The bottom of the page shows the AWS footer with copyright information and links to Privacy, Terms, and Cookie preferences.

Date: 9-10-2025

Exercise: AWS Identity Access Management (IAM) User and Group creation. Enable AWS IAM MFA. Create an AWS Account Alias (for Alternate Sign-in URL)

AWS Identity and Access Management (IAM) is a security service that helps you control who can access your AWS resources and what actions they can perform. It is a global AWS service. It allows you to securely manage users, groups, roles, and permissions in your AWS account.

Concept	Description
Root User	The account owner who created the AWS account. Has full access and should be used only for account setup.
Group	A collection of IAM users that share the same permissions. For example, a “Developers” group or “Students” group.
User	A person or application that interacts with AWS (e.g., student1, admin, developer). Each user has its own username, password, and access keys.
Policy	A JSON document that defines what actions are allowed or denied (e.g., “allow S3 read access”).
Role	A set of permissions that can be temporarily assumed by a user, service, or application — often used by EC2 instances or Lambda functions.

STEPS for AWS IAM User Group Creation


Step 1: Sign in to AWS Console

- Log in to your AWS Management Console using an administrator account.
- From the Services menu, search for IAM and open it.

The screenshot displays the AWS IAM Dashboard. At the top, there's a navigation bar with the AWS logo, a search bar, and account information (Account ID: 6293-0494-1888, Shiva Dhanush). The left sidebar shows the 'Identity and Access Management (IAM)' section with a search bar and a menu for 'Access management' (User groups, Users, Roles, Policies, Identity providers, Account settings, Root access management) and 'Access reports' (Access Analyzer, Resource analysis, Unused access, Analyzer settings, Credential report, Organization activity). The main content area features a blue banner for 'New access analyzers available'. Below this, the 'IAM Dashboard' section includes 'Security recommendations' (Root user has MFA, Root user has no active access keys), 'IAM resources' (User groups: 0, Users: 0, Roles: 2, Policies: 0, Identity providers: 0), and 'AWS Account' information (Account ID, Account Alias, Sign-in URL). A 'Quick Links' section at the bottom right provides links for security credentials and MFA. The footer shows 'CloudShell', 'Feedback', and copyright information for Amazon Web Services, Inc.

Step 2: Open Groups Section

- In the IAM dashboard, look at the left sidebar.
- Click on “User groups” → then click “Create group”.



Step 3: Name the Group

- Enter a Group name (example: Developers, Admins, Students, etc.).
- Group names must be unique within your account.

Step 4: Attach Permissions Policies


- You can attach IAM policies to define what members of the group can do.
Select the following:
 - AdministratorAccess → Full access to all services.
 - IAMUserChangePassword

Step 5: Review and Create

- Review all details (group name + permissions).
- Click “Create group” to finalize.

Group Successfully Created

- The new group now appears in the IAM dashboard.
- Any user added to this group automatically inherits all permissions attached to the group.



STEPS for AWS IAM User Creation

Step 1: Sign in to AWS Management Console

- Login to the AWS Management Console using your root user credentials.
- In the search bar, type IAM and open IAM service.

Step 2: Navigate to Users Section

- In the IAM dashboard's left panel (info panel), click on Users.
- Then click on the “Add users” button.

Step 3: Set User Details

- Enter a user name.
- Checkbox – ‘Provide user access to the AWS Management Console’.
- Select ‘I want to create an IAM user’ option

The screenshot shows the AWS IAM 'Create user' page. The 'User details' section is active, showing a 'User name' field with 'test-user-1'. Below it, there's a checkbox 'Provide user access to the AWS Management Console - optional' which is checked. A blue box highlights the 'User type' section with two options: 'Specify a user in Identity Center - Recommended' and 'I want to create an IAM user', with the latter being selected. Below this, the 'Console password' section shows 'Autogenerated password' selected, but 'Custom password' is also visible with a field containing 'test-user-1'.

Step 4: Set Permissions

You can grant permissions to the new user in three ways:

1. **Add user to group** – Assign predefined permission groups.
2. Copy permissions from an existing user.
3. Attach existing policies directly.

Set permissions

Add user to an existing group or create a new one. Using groups is a best-practice way to manage user's permissions by job functions. [Learn more](#)

The screenshot shows the 'Set permissions' section. Under 'Permissions options', 'Add user to group' is selected. Below this, the 'User groups (1/1)' section shows a table with one group: 'demo-group' with 0 users, attached policies 'AdministratorAccess and IAMU...', and created on '2025-10-13 (4 minutes ago)'. At the bottom, there's a 'Set permissions boundary - optional' section and navigation buttons: 'Cancel', 'Previous', and 'Next'.

Recommended: Use IAM **groups** to manage permissions easily.

Step 5: Set User Details and Tags (Optional)

- Add tags like Department: MCA or Role: Faculty/Student for easy identification.
- Click Next to review.

Tags - optional

Tags are key-value pairs you can add to AWS resources to help identify, organize, or search for resources. Choose any tags you want to associate with this user.

Key

Q Demo



Value - optional

Q Admin



Remove

Add new tag

You can add up to 49 more tags.

Step 6: Review and Create User

- Review all settings (user name, permissions, tags).
- Click Create user.

Step 7: Save Credentials


- After creation, AWS displays:
 - Console sign-in URL
 - Username
 - Password

NOTE: Download or copy these credentials immediately — they cannot be retrieved later.

The screenshot shows the AWS IAM console interface. At the top, there's a navigation bar with the AWS logo, a search bar, and account information (Account ID: 6293-0494-1888, Shiva Dhanush). Below the navigation bar, the breadcrumb trail shows 'IAM > Users > Create user'. A green success message banner at the top states 'User created successfully' and provides a 'View user' link. The main content area is titled 'Retrieve password' and includes a sub-header 'Console sign-in details'. It displays the 'Console sign-in URL' as 'https://629304941888.signin.aws.amazon.com/console', the 'User name' as 'test-user-1', and the 'Console password' as a masked string with a 'Show' link. A sidebar on the left lists the steps: 'Step 1: Specify user details', 'Step 2: Set permissions', 'Step 3: Review and create', and 'Step 4: Retrieve password' (which is currently selected). At the bottom of the main content area, there are three buttons: 'Cancel', 'Download .csv file', and 'Return to users list'. The footer of the console shows 'CloudShell', 'Feedback', and copyright information for 2025.

Step 8: Test the User Login

- Visit the IAM login URL provided (unique for your AWS account).
- Log in with the newly created username and password.
- Verify access and permissions.



Enable AWS IAM MFA


Step 1: Sign in to AWS Console as root user

Step 2: Open IAM Dashboard

- In the search bar, type IAM and select IAM (Identity and Access Management).
- From the left navigation pane, select Users.

Step 3: Select a User

- Click the user name for whom you want to enable MFA.
- This opens the User Summary page.



Step 4: Go to Security Credentials Tab

- Click the Security credentials tab.
- Scroll down to the section “Multi-factor authentication (MFA)”.

Step 5: Assign MFA Device

- Click “Assign MFA device”.
- Choose the MFA type:
 1. Virtual MFA device (e.g., Google Authenticator, Authy — most common)
 2. Security key (hardware-based, e.g., YubiKey)
 3. Authenticator app on phone

Step 6: Configure Virtual MFA

- If you select Virtual MFA device:
 1. Open your Google Authenticator or Authy app on your phone.
 2. Scan the QR code shown on the AWS screen.
 3. The app starts generating 6-digit codes.

Step 7: Verify MFA

- Enter two consecutive codes from your app in the verification fields.
- Click “Assign MFA”.

Step 8: Confirm Setup

- You will see a green checkmark confirming MFA is successfully assigned.
- The user now requires MFA each time they sign in.

The screenshot shows the AWS IAM console interface. At the top, there's a navigation bar with the AWS logo, a search bar, and user information. The main content area is titled "Identity and Access Management (IAM)" and shows the "Users" page for "test-user-1". A green notification banner at the top states "MFA device assigned". Below this, the "Summary" section displays the user's ARN, console access status, and creation date. The "Security credentials" tab is selected, showing the "Console sign-in" section with a link and password. At the bottom, the "Multi-factor authentication (MFA)" section shows the assigned MFA device and options to Remove, Resync, or Assign MFA device.

Create an AWS Account Alias (for Alternate Sign-in URL)

As an IAM user you can sign in using the default URL or create an account alias for it.

An Account Alias gives your AWS account a name instead of using the long numeric Account ID in your sign-in URL.

This makes it easier for IAM users to remember and log in.

Step 1:

- Sign in to the AWS Management Console using root user or an IAM user with administrative privileges.
- In the search bar, type IAM, and open the IAM service.

Step 2:

- In the left navigation pane, scroll down and select Dashboard.

Step 3:

- Under the “AWS Account” section, find “Account Alias”.
- Click on “Create” (or “Edit” if one already exists).

AWS Account

Account ID

 629304941888

Account Alias

demo-alias-1 [Edit](#) | [Delete](#)

Sign-in URL for IAM users in this account

 <https://demo-alias-1.signin.aws.amazon.com/console>

Step 4:

- In the pop-up box, enter your preferred alias name.
- Click “Create alias”.

Step 5:

- Once created, you'll see a new Sign-in URL displayed.

IAM user sign in ⓘ

Account ID or alias [\(Don't have?\)](#)

demo-alias-1

☐ Remember this account

IAM username

Password

☐ Show Password [Having trouble?](#)

Sign in

Sign in using root user email

Date: 14-10-2025

Exercise-3: Amazon S3 – Introduction, Bucket Creation and upload objects (files).

Amazon S3 (Simple Storage Service) is a fully managed, object-based storage service offered by AWS. It allows you to store and retrieve unlimited amounts of data from anywhere on the web.

Unlike traditional file systems that store data in folders, S3 stores data as objects inside buckets.

Each object has:

- Data (the actual file),
- Metadata (information about the file),
- Unique key (its name within the bucket).

You can store unlimited data in Amazon S3.

However, each AWS account can create up to 100 buckets by default (can be increased by request).

Each object (file) can be up to 5 TB (terabytes) in size.

Single upload limit: 5 GB (may vary), but larger files can be uploaded using Multipart Upload.

Amazon S3 is designed for:

- Durability: 99.999999999% (11 nines) – this means your data is extremely safe.
- Availability: 99.99% uptime – your data is almost always accessible.
- S3 automatically stores multiple copies of your data across multiple Availability Zones in a region.

Storage classes: S3 offers different storage classes depending on cost and frequency of access.

Storage Class	Description
S3 Standard	For frequently accessed data (default).
S3 Intelligent-Tiering	Automatically moves data between frequent - infrequent access tiers.
S3 Standard-IA	For infrequently accessed data but still quickly available.
S3 Glacier	For archival storage (low cost, slower retrieval).
S3 Glacier Deep Archive	For long-term archival (lowest cost).

Each object can be accessed through a **unique URL**.

S3 is commonly used for:


- Backup and archival
- Static website hosting
- Application data storage
- Media content delivery

Feature	Description
Buckets	Top-level containers for storing objects.
Objects	Actual files stored in S3 (e.g., images, HTML, PDFs).
Region	Each bucket is created in a specific AWS region.
Versioning	Maintains multiple versions of the same object for recovery.
Static Website Hosting	Allows you to host HTML pages directly from an S3 bucket.

Steps to Create an S3 Bucket and Upload an Image

Step 1: Open the S3 Service

- Sign in to your AWS Management Console.
- In the search bar, type S3 and click Amazon S3.



Step 2: Create a New Bucket

- Click “Create bucket”.
- Bucket type: General purpose
- Bucket name: Enter a name (Bucket names must be globally unique and lowercase.)
- AWS Region: Choose the region nearest to you.
- Scroll down and uncheck:
 - “Block all public access” → Uncheck this (for viewing file publicly).
- Confirm the warning checkbox.
- Keep all other settings as default.
- Click Create bucket.

Create bucket [Info](#)

Buckets are containers for data stored in S3.

General configuration

AWS Region
Asia Pacific (Mumbai) ap-south-1

Bucket type [Info](#)
☒ **General purpose**
Recommended for most use cases and access patterns. General purpose buckets are the original S3 bucket type. They allow a mix of storage classes that redundantly store objects across multiple Availability Zones.

☐ **Directory**
Recommended for low-latency use cases. These buckets use only the S3 Express One Zone storage class, which provides faster processing of data within a single Availability Zone.

Bucket name [Info](#)

Bucket names must be 3 to 63 characters and unique within the global namespace. Bucket names must also begin and end with a letter or number. Valid characters are a-z, 0-9, periods (.), and hyphens (-). [Learn More](#)

Copy settings from existing bucket - optional
Only the bucket settings in the following configuration are copied.

Format: s3://bucket/prefix

Object Ownership [Info](#)

Control ownership of objects written to this bucket from other AWS accounts and the use of access control lists (ACLs). Object ownership determines who can specify access to objects.

Object Ownership

☐ ACLs disabled (recommended)

All objects in this bucket are owned by this account. Access to this bucket and its objects is specified using only policies.

☒ ACLs enabled

Objects in this bucket can be owned by other AWS accounts. Access to this bucket and its objects can be specified using ACLs.

⚠ We recommend disabling ACLs, unless you need to control access for each object individually or to have the object writer own the data they upload. Using a bucket policy instead of ACLs to share data with users outside of your account simplifies permissions management and auditing.

Object Ownership

☒ Bucket owner preferred

If new objects written to this bucket specify the bucket-owner-full-control canned ACL, they are owned by the bucket owner. Otherwise, they are owned by the object writer.

☐ Object writer

The object writer remains the object owner.

💡 If you want to enforce object ownership for new objects only, your bucket policy must specify that the bucket-owner-full-control canned ACL is required for object uploads. [Learn more](#)

Block Public Access settings for this bucket

Public access is granted to buckets and objects through access control lists (ACLs), bucket policies, access point policies, or all. In order to ensure that public access to this bucket and its objects is blocked, turn on Block all public access. These settings apply only to this bucket and its access points. AWS recommends that you turn on Block all public access, but before applying any of these settings, ensure that your applications will work correctly without public access. If you require some level of public access to this bucket or objects within, you can customize the individual settings below to suit your specific storage use cases. [Learn more](#)

☐ Block all public access

Turning this setting on is the same as turning on all four settings below. Each of the following settings are independent of one another.

☐ Block public access to buckets and objects granted through new access control lists (ACLs)

S3 will block public access permissions applied to newly added buckets or objects, and prevent the creation of new public access ACLs for existing buckets and objects. This setting doesn't change any existing permissions that allow public access to S3 resources using ACLs.

☐ Block public access to buckets and objects granted through any access control lists (ACLs)

S3 will ignore all ACLs that grant public access to buckets and objects.

☐ Block public access to buckets and objects granted through new public bucket or access point policies

S3 will block new bucket and access point policies that grant public access to buckets and objects. This setting doesn't change any existing policies that allow public access to S3 resources.

☐ Block public and cross-account access to buckets and objects through any public bucket or access point policies

S3 will ignore public and cross-account access for buckets or access points with policies that grant public access to buckets and objects.

⚠ Turning off block all public access might result in this bucket and the objects within becoming public

AWS recommends that you turn on block all public access, unless public access is required for specific and verified use cases such as static website hosting.

☒ I acknowledge that the current settings might result in this bucket and the objects within becoming public.

Amazon S3 > Buckets

✔ Successfully created bucket "sd-demo-s3-bucket"
To upload files and folders, or to configure additional bucket settings, choose [View details](#).

General purpose buckets

All AWS Regions

Directory buckets

General purpose buckets (1) [Info](#)

Buckets are containers for data stored in S3.

Find buckets by name

Name	AWS Region	Creation date
☐ sd-demo-s3-bucket	Asia Pacific (Mumbai) ap-south-1	October 15, 2025, 11:49:48 (UTC+05:30)

Account snapshot [Info](#)

Updated daily

Storage Lens provides visibility into storage usage and activity trends. [View dashboard](#)

External access summary - new [Info](#)

Updated daily

External access findings help you identify bucket permissions that allow public access or access from other AWS accounts.

Step 3: Upload an Image File

1. Click on the newly created bucket name.
2. Click Upload → Add files.
3. Choose any image file from your computer.
4. Scroll down and click Upload.

Upload succeeded
For more information, see the Files and folders table.

Upload: status Close

After you navigate away from this page, the following information is no longer available.

Summary

Destination
s3://sd-demo-s3-bucket

Succeeded

1 file, 109.1 KB (100.00%)

Failed

0 files, 0 B (0%)

Files and folders

Configuration

Files and folders (1 total, 109.1 KB)

Find by name

< 1 >

Name	Folder	Type	Size	Status	Error
12_1600x.webp	-	image/webp	109.1 KB	Succeeded	-

Step 4: Make the Object Public

By default, S3 objects are private. To view them in a browser, make the file public.

1. After upload, go to the Objects tab inside your bucket.
2. Select the uploaded file.
3. Click Actions → Make public using ACL (or Object actions → Make public, depending on console version).
4. Confirm.

Make public Info

The make public action enables public read access in the object access control list (ACL) settings. [Learn more](#)

The object you've specified already has public read access enabled in the object access control list settings.

When public read access is enabled and not blocked by Block Public Access settings, anyone in the world can access the specified objects.

Specified objects

Find objects by name

Name	Type	Last modified	Size
12_1600x.webp	webp	October 15, 2025, 12:04:48 (UTC+05:30)	109.1 KB

Cancel Make public

Step 5: Copy the Object URL

1. Open the object again by clicking its name.
2. Scroll down to the **Object URL** section.
3. Copy this URL and open it in a new browser tab.

The image is displayed directly from the S3 bucket — meaning AWS S3 is serving that object over HTTP.

S3 URI

 s3://sd-demo-s3-bucket/12_1600x.webp

Amazon Resource Name (ARN)

 arn:aws:s3:::sd-demo-s3-bucket/12_1600x.webp

Entity tag (Etag)

 [e2dda3fb0ca810de72df165c27d6c449](#)

Object URL

 https://sd-demo-s3-bucket.s3.ap-south-1.amazonaws.com/12_1600x.webp

Date: 15-10-2025

Exercise-4: Amazon S3 – Static Website Hosting (Multi-Page website), Versioning, Cross-Region Replication rule.

Amazon S3 Static Website Hosting

Amazon S3 can host a static website – [a website consisting of only HTML, CSS, JavaScript, images, etc. – no server-side scripting like PHP or Python].

When you enable “Static Website Hosting,” your S3 bucket acts like a web server, and AWS provides a public website URL to access it.

You can create a multi-page static website (e.g., index.html, about.html, contact.html) and upload it to S3. Links within these pages allow users to navigate between them just like a normal website.

Steps to Create a Multi-Page Static Website on S3

Step 1: Create an S3 Bucket

- Open the AWS Management Console → Navigate to S3.
- Click Create bucket.
- Select Bucket type: General purpose
- Enter a unique bucket name (e.g., my-static-web-demo).
- Uncheck “Block all public access.”
- Click Create bucket.

Step 2: Prepare Website Files

Before uploading, organize your files in a folder structure as follows:

my-website/

```
|
|— index.html
|— about.html
|— contact.html
|— error.html
|— images/
    |— banner.jpg
```

Each HTML file should include navigation links.

Step 3: Upload Website Files

- Open your S3 bucket → Click Upload.
- Add all files and folders (HTML, CSS, JS, images).
- Click Upload to store them in S3.

Upload succeeded
For more information, see the Files and folders table.

Summary

Destination

s3://sd-demo-s3-bucket

Succeeded

5 files, 37.8 KB (100.00%)

Failed

0 files, 0 B (0%)

Files and folders

Configuration

Files and folders (5 total, 37.8 KB)

Find by name

Name	Folder	Type	Size	Status	Error
banner.jpg	images/	image/jpeg	32.7 KB	Succeeded	-
about.html	-	text/html	1.1 KB	Succeeded	-
contact.html	-	text/html	1.9 KB	Succeeded	-
error.html	-	text/html	976.0 B	Succeeded	-
index.html	-	text/html	1.2 KB	Succeeded	-

Step 4: Enable Static Website Hosting

- Go to the Properties tab of the bucket.
- Scroll down to Static website hosting → Click Edit.
- Choose Enable and select ‘Host a static website’.
- Set:
 - Index document: index.html
 - Error document: error.html
- Click Save changes.

Static website hosting

Use this bucket to host a website or redirect requests. [Learn more](#)

Static website hosting

- ☐ Disable
- ☒ Enable

Hosting type

- ☒ Host a static website
Use the bucket endpoint as the web address. [Learn more](#)
- ☐ Redirect requests for an object
Redirect requests to another bucket or domain. [Learn more](#)

For your customers to access content at the website endpoint, you must make all your content publicly readable. To do so, you can edit the S3 Block Public Access settings for the bucket. For more information, see [Using Amazon S3 Block Public Access](#)

Index document

Specify the home or default page of the website.

index.html

Error document - optional

This is returned when an error occurs.

error.html

Step 5: Make Files Public (Bucket Policy)

By default, your files are private. To make them public:

- Go to the Permissions tab → Bucket Policy → Edit.

- Paste the following policy (**replace bucket name**):

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "PublicReadGetObject",
      "Effect": "Allow",
      "Principal": "*",
      "Action": "s3:GetObject",
      "Resource": "arn:aws:s3:::my-static-web-demo/*"
    }
  ]
}
```

- Save the changes.

Bucket policy

The bucket policy, written in JSON, provides access to the objects stored in the bucket. Bucket policies don't apply to objects owned by other accounts. [Learn more](#)

[Edit](#)[Delete](#)

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "PublicReadGetObject",
      "Effect": "Allow",
      "Principal": "*",
      "Action": "s3:GetObject",
      "Resource": "arn:aws:s3:::sd-demo-s3-bucket/*"
    }
  ]
}
```

[Copy](#)

Step 6: Access Your Website

- Go to the Properties tab → Scroll to Static website hosting.
- Copy the Bucket Website Endpoint URL.
- Paste it into your browser — your homepage (index.html) should appear.
- Use the header links to navigate between pages (About, Contact, etc.).

Static website hosting

Use this bucket to host a website or redirect requests. [Learn more](#)

[Edit](#)

 We recommend using **AWS Amplify Hosting** for static website hosting

Deploy a fast, secure, and reliable website quickly with AWS Amplify Hosting. Learn more about [Amplify Hosting](#) or [View your existing Amplify apps](#)

[Create Amplify app](#)

S3 static website hosting

Enabled

Hosting type

Bucket hosting


Bucket website endpoint

When you configure your bucket as a static website, the website is available at the AWS Region-specific website endpoint of the bucket. [Learn more](#)

<http://sd-demo-s3-bucket.s3-website-ap-south-1.amazonaws.com>

When Will the Error Page Be Shown?

If a user enters a wrong URL or tries to access a file that doesn't exist (e.g., /abc.html), Amazon S3 automatically displays the file you set as the **Error document** (error.html).




Amazon S3 Versioning

Versioning allows you to keep multiple versions of an object in a bucket. If a file is accidentally deleted or overwritten, you can recover the previous version. Each version gets a unique version ID.

Steps to Enable Versioning

- Go to your S3 bucket
- Open the Properties tab
- Scroll to Bucket Versioning
- Click Edit → Enable
- Click Save changes

Now whenever you upload a file with the same name, S3 will keep both versions.



Find objects by prefix
 Show versions
< 1 >
⚙️

☐	Name	Type	Version ID	Last modified	Size	Storage class
☐	1MS24MC031- Shiva Dhanush [2024-26] Cloud Computing-Lab Manual (1).docx	docx	6yM574RG7fcQhpuKkNFk8Z_PEDNcvxwk	October 28, 2025, 12:00:47 (UTC+05:30)	3.4 MB	Standard
☐	1MS24MC031- Shiva Dhanush [2024-26] Cloud Computing-Lab Manual (1).docx	docx	LURhbd4m5nNc92a2qQnyjy_w6EmDYeHF	October 28, 2025, 11:59:18 (UTC+05:30)	3.4 MB	Standard

🔍 Successfully deleted objects
View details below.

Delete objects: status
Close

🔍 After you navigate away from this page, the following information is no longer available.

Summary

Source

s3://sd-demo-s3-bucket

Successfully deleted

🟢 2 objects, 6.7 MB

Failed to delete

0 objects

Failed to delete
Configuration

🔍 Failed to delete (0)

Find objects by name

Name	Folder	Version ID	Type	Last modified	Size	Error
No objects failed to delete.						

Note: You can view versions by clicking “List versions” in the bucket objects page.

To Restore or Delete a Specific Version

- Click the object name → Versions
- Select the desired version → Download / Delete
(Deleting only adds a delete marker — older versions are still stored.)

Cross-Region Replication (CRR) – is a rule created on S3

CRR automatically copies objects from one S3 bucket (source) to another (destination) in a different AWS Region.

Used for disaster recovery, compliance, or low-latency access in another region.

Requires Versioning to be enabled on both buckets.

Steps to Set Up CRR

NOTE: Enable Versioning on both:

Source bucket

Destination bucket

Step 1: Choose a different region before you create Destination bucket

Create Destination Bucket first.

The screenshot shows the AWS Management Console interface for the 'sd-paris-s3-bucket'. The top navigation bar includes the AWS logo, a search bar, and the user's account information (Shiva Dhanush, Account ID: 6293-0494-1888). The breadcrumb trail indicates the path: Amazon S3 > Buckets > sd-paris-s3-bucket. The main content area shows the bucket's details, including its name, region (Europe (Paris) eu-west-3), and creation time (October 28, 2025, 12:16:10 UTC+05:30). The 'Objects' section shows no objects in the bucket.

Name	Type	Last modified	Size	Storage class
No objects				

Step 2: Create Source Bucket

Give replication permission:

Source bucket → Management tab → Replication rules → Create rule

Step 3: Create Replication Rule page

Enter a Replication rule name

Status: Enabled

Source bucket section:

Choose a rule scope: select “Apply to all objects in the bucket”

Destination:

Select “Choose a bucket in this account”

Bucket name: Select the destination bucket

IAM role:

Select “Create new role”

[An IAM role gives Amazon S3 permission to replicate objects from your source bucket to your destination bucket.

S3 replication won't work unless you assign (or create) a role with the right permissions.

IAM Role: Either create a new role automatically or choose an existing one]

Replication rules (1)

View details

Edit rule

Delete

Actions

Create replication rule

Use replication rules to define options you want Amazon S3 to apply during replication such as server-side encryption, replica ownership, transitioning replicas to another storage class, and more. [Learn more](#)

< 1 >

Replication rule name	Status	Destination bucket	Destination Region	Priority	Scope	Storage class	Replica owner	Replication Time Control	KMS-encrypted objects (SSE-KMS or DSSE-KMS)	Replica modification sync
☐ myrule	Enabled	s3://sd-paris-s3-bucket	Europe (Paris) eu-west-3	0	Entire bucket	Same as source	Same as source	Disabled	Do not replicate	Disabled

Step 4: Save

Any new objects uploaded to the source bucket will automatically replicate to the destination region.

sd-mumbai-s3-bucket [Info](#)

Objects

Properties

Permissions

Metrics

Management

Access Points

Objects (1)

Copy S3 URI

Copy URL

Download

Open

Delete

Actions

Create folder

Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Show versions

< 1 >

☐	Name	Type	Last modified	Size	Storage class
☐	USN - NAME [2024-26] Cloud Computing-Lab Manual.docx	docx	October 28, 2025, 12:27:07 (UTC+05:30)	151.5 KB	Standard

sd-paris-s3-bucket [Info](#)

Objects

Properties

Permissions

Metrics

Management

Access Points

Objects (1)

Copy S3 URI

Copy URL

Download

Open

Delete

Actions

Create folder

Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Show versions

< 1 >

☐	Name	Type	Last modified	Size	Storage class
☐	USN - NAME [2024-26] Cloud Computing-Lab Manual.docx	docx	October 28, 2025, 12:27:07 (UTC+05:30)	151.5 KB	Standard

Note: Replication is not retroactive — only new uploads after enabling CRR are copied.

Reminder – Resource cleanup – to release/delete/terminate the resources created.

Date: 29-10-2025

Exercise-5: Overview of EC2, To Launch a Windows EC2 Instance and Connect via RDP Client

Amazon Elastic Compute Cloud (EC2) is a core AWS Compute service that lets you run virtual servers (instances) in the cloud.

Amazon EC2 is an Infrastructure as a Service (IaaS) offering from AWS.

It allows you to launch virtual machines to host applications and manage them remotely – wherever you are in the world.

Key concepts:

Instance - A virtual machine running in the AWS cloud.

AMI (Amazon Machine Image) - A pre-configured template that includes: OS (Linux, Windows, etc.), Application software, other configurations.

Instance Type - Defines hardware power:

Family	Example	Use Case
General Purpose	t2.micro	Basic web apps
Compute Optimized	c5.large	High-performance computing
Memory Optimized	r5.large	Databases, analytics
Storage Optimized	i3.large	Data warehousing
GPU Instances	g4dn.xlarge	ML/AI, graphics

EBS (Elastic Block Store) - Persistent storage for your EC2 instance. Acts like a hard drive — data remains even after instance stops. Types: SSD, HDD, etc.

Security Groups - Virtual firewalls controlling inbound and outbound traffic.

Example: Allow HTTP (port 80), SSH (port 22), HTTPS (port 443)

Key Pair - Used for secure login (SSH for Linux, RDP for Windows). Consists of a public key (stored in AWS) and private key (.pem) that you download.

Common Ways to Access EC2

- SSH (Linux instances)
- RDP (Windows instances) - Use Remote Desktop with Administrator password.
- User Data Script: Run automation commands during instance launch.

EC2 Use Cases

- Hosting static or dynamic websites
- Deploying web servers (Apache/Nginx)
- Running applications, APIs, or databases
- Machine Learning model hosting

- Batch processing jobs

Pricing Models

- On-Demand: Pay per hour/second; flexible.
- Reserved Instances: 1–3 year commitment; cheaper.
- Spot Instances: Unused capacity; up to 90% cheaper.
- Free Tier: t2.micro or t3.micro free for 6 months.

Lifecycle of an Instance

Step	Description
Launch	Choose AMI, type, key, security group
Running	Accessible and operational
Stop	Instance paused, EBS persists
Start	Boot again from same EBS
Terminate	Deleted permanently, data lost unless backed up

Two types of IPv4 addresses

When you launch an EC2 instance, AWS automatically assigns two types of IPv4 addresses depending on your network settings (VPC, subnet, etc.):

1. Private IPv4 Address

- Purpose: Used for internal communication within the same VPC (Virtual Private Cloud).
- Assigned Automatically: Yes, by AWS from the subnet's private IP range.
- Visibility: Not accessible from the Internet.
- Use Case:
 - Instance-to-instance communication inside AWS (e.g., web server ↔ database server).
 - Internal services that do not need public internet access.
- Persistence: The private IP remains attached to the instance until it is terminated.

2. Public IPv4 Address

- Purpose: Used for communication over the Internet.
- Assigned Automatically:
 - Yes, if your subnet is public (i.e., auto-assign public IP enabled).
 - No, if it's a private subnet.
- Visibility: Accessible from the Internet.
- Use Case:
 - Accessing the instance via SSH or HTTP from your local system.
 - Hosting web applications publicly.
- Persistence:
 - The public IP changes each time you stop/start the instance.
 - To make it permanent, you can assign an Elastic IP (static public IP).

Remote Desktop Protocol [RDP], is a secure communication protocol developed by Microsoft that allows a user to connect to and control another computer remotely. It establishes an encrypted channel to transmit keyboard and mouse inputs from the client to the remote computer and send screen information back to the client, providing a graphical user interface for remote access.

An **RDP client** is the software or app that you use to make this remote connection.

It connects to a remote Windows server or Windows EC2 instance that is running an RDP server (which listens on port **3389**). It is pre-installed in Windows.

Steps to Launch a Windows EC2 Instance and Connect via RDP Client

Step 1: Sign in to AWS Management Console

Select the **Region** closest to you.

Step 2: Open the EC2 Dashboard

In the AWS Console, search for **EC2** in the search bar.

Click on EC2 → Instances → Launch Instance.


Under Name and Tags, give your instance a name.

Step 3: Choose an Amazon Machine Image (AMI)

Under Application and OS Images (Amazon Machine Image) → click Browse more AMIs or choose: Microsoft Windows Server 2022 Base (Free tier eligible).

▼ Application and OS Images (Amazon Machine Image) [Info](#)

An AMI contains the operating system, application server, and applications for your instance. If you don't see a suitable AMI below, use the search field or choose **Browse more AMIs**.



Amazon Machine Image (AMI)

Step 4: Choose Instance Type

Choose t3.micro (Free-tier eligible).

Step 5: Configure Key Pair

Under Key pair (login), choose an existing key pair or create a new one.

If creating a new key pair:

- Choose type: RSA
- Format: .pem
- Download and save it safely — it's required to decrypt your Windows password later.

▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - *required*

demo-instance-key



[Create new key pair](#)

For Windows instances, you use a key pair to decrypt the administrator password. You then use the decrypted password to connect to your instance.

Step 6: Configure Network Settings

Leave default VPC and Subnet settings.

Under Firewall (security group) →

- Select Create security group.
- Allow RDP (port 3389) access from My IP (for better security) or anywhere (0.0.0.0/0).

Firewall (security groups) [Info](#)

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☒ Create security group

☐ Select existing security group

We'll create a new security group called 'launch-wizard-1' with the following rules:

☒ Allow RDP traffic from
Helps you connect to your instance

Anywhere
0.0.0.0/0

☐ Allow HTTPS traffic from the internet
To set up an endpoint, for example when creating a web server

☐ Allow HTTP traffic from the internet
To set up an endpoint, for example when creating a web server

Step 7: Launch the Instance

Review all configurations.

Click Launch Instance.

Wait until the Instance state changes to Running and Status check = 3/3 passed.

Note:

Wait approximately 5 minutes after instance launch to allow AWS to fully initialize the instance and make the Administrator password available.

When a Windows EC2 instance is first launched, AWS needs a few minutes to:

Initialize the instance, Attach the root volume, Generate the Windows Administrator password (encrypted using your key pair).

Step 8: Get the Administrator Password


Select your running instance → click Connect → choose RDP Client tab.
Click Get Password (you must wait about 4 minutes after launch).
Upload your .pem key file and click Decrypt Password.
Copy the Public IPv4 address and Administrator Password shown.

Passwprd: UHaDcm)R@@owE8BpPlSW4kbi7wletu8-

Step 9: Connect Using RDP Client

On Windows system:

1. Open Remote Desktop Connection (from Start Menu).
2. Enter your instance's Public IPv4 address.
3. Click Connect → Enter:
 - Username: Administrator
 - Password: *(the decrypted password from AWS)*
4. Click OK → accept the certificate → the remote Windows desktop opens!



Step 10: Verify Connection

- You should now see a Windows Server desktop running inside your local window.
- You can open File Explorer, browse settings, or install software (within the free-tier limits).

Note:

- Always **Stop** (not Terminate) the instance when not in use to avoid charges.
- RDP uses port 3389, so ensure it's open in the security group.
- Avoid sharing your decrypted password or key pair.

DATE: 30-10-2025

Exercise 6: Launch a Linux EC2 Instance and Connect using SSH through PowerShell/Linux Terminal and PuTTY on Windows.

Note: Choosing the Correct Key Pair Format

While creating a Key Pair, you are asked to select the Private Key File Format — either .pem or .ppk. The correct choice depends on the operating system and the method you will use to connect to your EC2 instance.

Scenario	Key File Format	Explanation
Using PuTTY on Windows	.ppk	The .ppk file is specific to the GUI based PuTTY application, a popular SSH client for Windows system.
Using PowerShell on Windows, Linux terminal on Linux	.pem	The .pem file is the default AWS key format used by the OpenSSH client available in PowerShell (Windows), Linux, and macOS terminals. Used for CLI.

Steps to Launch a Linux EC2 Instance and Connect using SSH through PowerShell/Linux

Step 1: Sign in to AWS Management Console

select the nearest AWS Region.

Step 2: Open EC2 Service

In the search bar, type EC2 and click EC2 Dashboard.


Select Instances → Launch Instance.

Step 3: Configure Instance Details

Under Name and Tags, give your instance a name, e.g., *Linux-SSH-Demo*.

Under Application and OS Images (AMI) → select Amazon Linux 2 AMI (Free tier eligible).

Under Instance type, choose t3.micro (Free-tier eligible).



Step 4: Create or Select a Key Pair

Under Key pair (login) → choose Create new key pair.

Choose:

- Key pair type: RSA
- Private key file format: .pem (for SSH via PowerShell/Linux)

Click Create key pair → the .pem file will automatically download.

Save it securely on your local machine (you'll need it for SSH login).

Step 5: Configure Network Settings

Under Network settings, leave the defaults.

In Firewall (security group) → select Create security group.

Ensure SSH (port 22) is allowed:

- Type: SSH
- Protocol: TCP
- Port Range: 22
- Source: My IP (recommended for security) or Anywhere (0.0.0.0/0).

Step 6: Launch the Instance

Review the configuration → click Launch Instance.


Wait for the Instance State to show Running and Status Checks: 3/3 passed.

Step 7: Get the Public IP Address

Select your instance → In the summary section →

Copy the Public IPv4 address — you'll use this to connect.

Step 8: Connect using SSH



(A) On Windows using PowerShell


- Open **PowerShell** (search “PowerShell” from the Start menu).
- Navigate to the folder where your .pem file is saved:
`cd "C:\Users\<YourName>\Downloads"`
- Connect using the SSH command:
`ssh -i "keyfile.pem" ec2-user@<Public-IP-address>`
- When prompted, type **yes** to continue connecting.
- You’ll now be logged into your EC2 Linux terminal.

(B) On Linux Terminal (Ubuntu / macOS)

- Open **Terminal**.
- Navigate to the directory where your .pem file is stored.
- Set proper permission for the key file:
`chmod 400 keyfile.pem`
- Connect to the instance:
`ssh -i keyfile.pem ec2-user@<Public-IP-address>`
- Type **yes** when prompted.
- You will be connected to your EC2 instance remotely.

Step 9: Verify Connection

- Once connected, the command prompt will appear as:
`[ec2-user@ip-172-31-xx-xx ~]$`
- Try a few commands:
uname -a and sudo yum update -y to confirm access.



```
C:\Users\WCA-01\Downloads>ssh -i "cc_class.pem" ec2-user@ec2-57-183-17-120.ap-northeast-1.compute.amazonaws.com
ec2-user@ec2-57-183-17-120:~$
Amazon Linux 2023
https://aws.amazon.com/linux/amazon-linux-2023

Last login: Tue Nov 4 06:10:10 2025 from 49.206.203.162
[ec2-user@ip-172-31-26-189 ~]$
```

Step 10: Stop Instance after Use

Go to the EC2 console.

Select your instance → Instance State → Stop Instance (to avoid charges).

Steps to Install PuTTY on Windows


PuTTY is a client program for the SSH, Telnet and Rlogin network protocols. These protocols allow you to interact with a remote server as if you were sitting right in front of it.

It is primarily used in the Windows platform.

In an era where cloud computing and remote servers are the norm, being able to securely connect and interact with these servers is vital. It provides a secure and reliable way to connect to these remote systems. It supports a range of network protocols including the secure ones like SSH.

Use the correct, safe download link: Official download page: This is the only genuine PuTTY source.

<https://www.chiark.greenend.org.uk/~sgtatham/putty/latest.html>



Under the **Package files** section, look for:

MSI ('Windows Installer')

64-bit x86: putty-64bit-0.83-installer.msi

Click the link: **putty-64bit-0.83-installer.msi**



Once it downloads, open the file and follow: Next → Next → Install → Finish

After installation, you will have:

- PuTTY – to connect to your EC2 instance via SSH
- PuTTYgen – to convert .pem to .ppk (if needed)
- Pageant – optional key manager

You can open PuTTY by typing **PuTTY** in the Windows search bar/start menu.

Steps to Launch a Linux EC2 Instance and Connect using PuTTY on Windows

Step 1: Sign in to AWS Management Console

Select your nearest AWS Region.

Step 2: Open the EC2 Service

In the AWS Console search bar, type EC2 and select EC2 Dashboard.
Click Instances → Launch Instance.

Step 3: Configure Instance Details

Under Name and Tags, enter an instance name, e.g., *Linux-PuTTY-Demo*.

Under Application and OS Images (AMI) → choose Amazon Linux 2 AMI (Free tier eligible).

Under Instance Type, select t2.micro (Free tier eligible).

Step 4: Create or Select a Key Pair


Under Key pair (login) → select Create new key pair.

Choose the following:

- Key pair type: RSA
- Private key file format: .ppk (for PuTTY on Windows)

Click Create key pair → a .ppk file will download automatically.

Save it securely — this file is required to connect later.



Step 5: Configure Network Settings

Under Network settings, leave default VPC/Subnet settings.

Under Firewall (security group):

- Select Create security group.
- Ensure SSH (port 22) is allowed:
 - Type: SSH
 - Protocol: TCP
 - Port Range: 22
 - Source: anywhere.

Step 6: Launch the Instance

Review all configurations.

Click Launch Instance.

Wait until the Instance State changes to Running and Status Check = 3/3 passed.

Step 7: Obtain the Public IP

Select your instance → In the summary section →

Copy the Public IPv4 address or Public DNS (IPv4) — you'll use this to connect from PuTTY.

Step 8: Connect using PuTTY

Open **PuTTY** on your Windows system.

In the **Host Name (or IP address)** field, enter: `ec2-user@<Public-IP-address>`


In the **Category** list on the left, expand: Connection → SSH → Auth → Credentials

Click **Browse** → locate and select your .ppk key file downloaded earlier.

Click **Open**.

When prompted, click **Accept** to trust the host.

A terminal window opens — you're now connected to your EC2 Linux instance!



Step 9: Verify Connection

Once connected, your prompt should look like: `[ec2-user@ip-172-31-xx-xx ~]$`

Try verifying: `uname -a` or update packages: `sudo yum update -y`

DATE: 5-11-2025

Exercise 7: Hosting a static website on EC2 instance.

- Manual Installation of Apache (httpd) Web Server on EC2 for Static Website Hosting.
- Launching an EC2 Instance with User Data Script to Automatically Install Apache and Host a Static Website.

When you create an EC2 instance, it's just a **bare machine** — It does not have the software to host a website yet. To host a website:

- You need a web server software → Apache (httpd)
- You put your website files (HTML, CSS, etc.) in a special folder (usually /var/www/html)
- Apache listens on port 80 (HTTP) or port 443 (HTTPS) for incoming connections

Apache HTTP Server (often called Apache or httpd) is a web server software.

It runs on a computer (like your EC2 instance) and delivers web pages (HTML, images, CSS, etc.) to users over the HTTP/HTTPS protocol.

So, whenever someone types your website's URL (like http://your-ec2-ip/), Apache receives that request and serves your webpage files from your server to the browser.

httpd stands for HTTP Daemon.

A *daemon* in Linux is a background service that keeps running to handle requests.

So, httpd is the background process that runs the Apache web server.

When you install Apache, you're really installing the **httpd service**.

File Locations (default) – Website files go into: /var/www/html

Steps for Manual Installation of Apache (httpd) Web Server on EC2 for Static Website Hosting.

Part 1 – Creating an EC2 Instance

Step 1: Sign in to AWS Management Console

In the search bar, type EC2 and open the EC2 Dashboard.

Step 2: Launch a New Instance

Click “Launch Instance.”

Give a name. (e.g., MyApacheServer)

Step 3: Choose an Amazon Machine Image (AMI)

Select Amazon Linux 2 AMI (Free tier eligible)

Step 4: Choose Instance Type

Choose t3.micro (Free tier eligible).

Step 5: Create / Select Key Pair

Under Key pair (login), choose:

Create new key pair → Give a name → File format = .pem

Download the key file → Keep it safe.

Step 6: Configure Network Settings

Click → Allow SSH traffic

→ Allow HTTP traffic from the internet. (This automatically opens port 80 for web access).

Step 7: Launch Instance

Review → Click **Launch Instance**.

Wait for the instance to start.

Step 8: Connect to the Instance

Using powershell type the following commands:

- Open PowerShell (search “PowerShell” from the Start menu).
- Navigate to the folder where your .pem file is saved:
cd "C:\Users\<YourName>\Downloads"
- Connect using the SSH command:
ssh -i "keyfile.pem" ec2-user@<Public-IP-address>
- When prompted, type yes to continue connecting.
- You’ll now be logged into your EC2 Linux terminal.

Part 2 – Manual Installation of Apache (httpd) Web Server

Step 1: Update the Packages

```
sudo yum update -y
```

Step 2: Install Apache (httpd)

```
sudo yum install httpd -y
```

(This installs the Apache Web Server package.)

Step 3: Start the Apache Service

```
sudo systemctl start httpd
```

Step 4: Enable Apache to Start on Boot

```
sudo systemctl enable httpd
```

Step 5: Check Apache Status

sudo systemctl status httpd [It should show active (running). Press q to exit status view.]

```
ec2-user@ip-172-31-19-112:~$ sudo apt update -y
Warning: Permanently added 'ec2-35-78-248-84.ap-northeast-1.compute.amazonaws.com' (ED25519) to the list of known hosts.
Amazon Linux 2023
https://aws.amazon.com/linux/amazon-linux-2023


[ec2-user@ip-172-31-19-112 ~]$ sudo apt update -y
sudo: apt: command not found
[ec2-user@ip-172-31-19-112 ~]$ sudo yum update -y
Amazon Linux 2023 Kernel Livepatch repository
Dependencies resolved.
Nothing to do.
Complete!
[ec2-user@ip-172-31-19-112 ~]$ sudo yum install httpd -y
Last metadata expiration check: 0:00:14 ago on Wed Nov 5 05:53:03 2025.
Dependencies resolved.
=====
Package                Architecture      Version           Repository        Size
=====
Installing:
httpd                  x86_64            2.4.65-1.amzn2023.0.2  amazonlinux      47 k
Installing dependencies:
apr                    x86_64            1.7.5-1.amzn2023.0.4  amazonlinux      129 k
apr-util               x86_64            1.6.3-1.amzn2023.0.2  amazonlinux      97 k
apr-util-ldb           x86_64            1.6.3-1.amzn2023.0.2  amazonlinux      32 k
generic-logos-httpd    noarch            18.0.0-12.amzn2023.0.3  amazonlinux      19 k
httpd-core             x86_64            2.4.65-1.amzn2023.0.2  amazonlinux      1.4 M
httpdfilesystem        noarch            2.4.65-1.amzn2023.0.2  amazonlinux      15 k
httpd-tools            x86_64            2.4.65-1.amzn2023.0.2  amazonlinux      81 k
libserf1               x86_64            1.0.0-4.amzn2023.0.2  amazonlinux      215 k
mailcap                noarch            2.1.49-3.amzn2023.0.3  amazonlinux      33 k
Installing weak dependencies:
apr-util-openssl       x86_64            1.6.3-1.amzn2023.0.3  amazonlinux      15 k
mod_httpd               x86_64            2.0.27-1.amzn2023.0.3  amazonlinux      166 k
mod_lua                 x86_64            2.4.65-1.amzn2023.0.2  amazonlinux      60 k
=====
```

Step 7: Test Apache Server

Copy your Public IPv4 address from the EC2 dashboard.

Paste it into a browser: http://

You should see the Apache Test Page



It works!

Part 3 – Host a Static Website on Apache

Step 1: Move to Web Root Folder

cd /var/www/html

Step 2: Create your HTML file

sudo nano index.html

Step 3: Add HTML content

Paste the following code:

```
<!DOCTYPE html>
```

```
<html>
<head>
<title>Welcome to My Website</title>
</head>
<body style="text-align:center; background-color: #f0f0f0;">
<h1>Hello from Apache on AWS EC2!</h1>
<p>This is a static website hosted on Apache web server.</p>
</body>
</html>
```

Press Ctrl + O, Enter, then Ctrl + X to save and exit.

After Pressing Ctrl + O You'll see:

File Name to Write: index.html [Just press Enter to confirm saving with that name].


To Exit Nano: press Ctrl + X [This closes the Nano editor and will be back to the Linux prompt]

Step 4: Restart Apache

`sudo systemctl restart httpd`

Step 5: View Website

- In browser: `http://<your-ec2-public-ip>`
- You'll see your custom HTML page.



Optional: Add Multiple Pages

- Upload other pages like `about.html`, `contact.html` in the same directory.

Steps for launching an EC2 Instance with User Data Script to Automatically Install Apache and Host a Static Website.

In this method, you don't manually install Apache or upload files after connecting. Instead, you write a **shell script** in the **User Data** section (during instance creation). When the EC2 instance starts for the first time, it automatically:

1. Installs Apache (httpd)
2. Starts the service
3. Creates an index.html webpage
4. Hosts your website immediately after launch

Step 1: Sign in to AWS Management Console

Go to EC2 Dashboard → Click Launch Instance

Step 2: Name and OS

Name: AutoApacheWebServer

AMI: Choose Amazon Linux 2 AMI (Free tier eligible)

Instance Type: t3.micro

Step 3: Key Pair

Select existing key pair (or create new) → Format = .pem

Step 4: Network Settings

Allow HTTP traffic from the Internet (port 80)

Allow SSH traffic (port 22)

Step 5: Add User Data Script

Scroll down to Advanced Details → User data box.

This script:

- Updates all packages
- Installs and starts Apache
- Creates a sample index.html webpage in /var/www/html

Paste the following shell script:


```
#!/bin/bash
# Update packages
yum update -y

# Install Apache Web Server
yum install -y httpd

# Start Apache
systemctl start httpd
systemctl enable httpd

# Create website content
echo "<html>"
```

```
<head> <title>Welcome to My Auto Web Server</title> </head>
<body style='text-align:center; background-color:#e9f5ff;'>
<h1>Welcome to EC2 Instance</h1>
<h2>Apache Installed Automatically via User Data Script</h2>
<p>This is a static website hosted on EC2 using User Data automation.</p>
</body>
</html>" > /var/www/html/index.html
```



Step 6: Launch the Instance


Click Launch Instance

Wait for the status → Running

Step 7: Test Your Web Server

Copy the Public IPv4 address of your instance.

Paste it in your browser: <http://<your-public-ip>>



You should immediately see your webpage without logging into EC2!

DATE: 7-11-2025

Exercise 8: Create a Custom AMI from a Working EC2 Instance.

Launch an EC2 Instance using a Custom AMI

Delete the custom AMI

AMI (Amazon Machine Image) is a pre-configured template that contains only Operating System (e.g., Amazon Linux, Ubuntu, Windows) needed to launch an EC2 instance.

When you launch a new EC2 instance, you select an AMI as the base image.

A **Custom AMI** is created from your own *running EC2 instance* after you've installed software, uploaded website files, or applied settings.

Benefits:

1. **Reusability** – Launch multiple identical servers quickly.
2. **Backup** – Acts as a snapshot of your configured instance.
3. **Auto Scaling** – Used by Auto Scaling Groups to create identical instances automatically.
4. **Disaster Recovery** – You can recreate your setup if the original instance fails.
5. **Time-saving** – No need to reinstall Apache or re-upload files each time.

A Custom AMI is like a master copy of your EC2 setup. It ensures your website or app can be duplicated instantly and consistently.

Steps to Create a Custom AMI from a Working EC2 Instance

To capture the current configuration — installed packages, website files, and settings — into a reusable Amazon Machine Image (AMI).

Step 1: Select the running instance

Go to EC2 → Instances.

Select the instance that already hosts your website.

Or create an instance (add user data for web hosting)

Step 2: Create Image

From the Actions → Image and templates → Create image.

St

Storage type	Device	Snapshot	Size	Volume type	IOPS	Throughput	Delete on termination	Encrypted
EBS	/dev/sd...	Create new snapshot	8	EBS General Purpose	3000		☒ Enable	☐ Enable

Description: AMI created from configured Apache website instance

Leave “No reboot” unchecked (so the filesystem is consistent).

Step 4: Storage volumes

The root volume will appear automatically; keep defaults unless you need more space.


Step 5: Create image

Click **Create image**.

A confirmation message appears; note the **Image ID**.

Step 6: Verify creation

In left panel → **AMIs** → refresh until **Status = Available**.



Your custom AMI is now saved in that region and can be used to launch identical webserver instances.

Steps to Launch an EC2 Instance using a Custom AMI

To launch a new EC2 instance from a previously created Custom Amazon Machine Image (AMI) — containing your configured website and software.

Step 1: Go to AMIs

Open the EC2 Service dashboard.

In the left navigation pane, click AMIs (under “Images”).


Step 2: Select Your Custom AMI

Locate the AMI you created earlier (e.g., MyWebsiteAMI).

Ensure Status = Available.

Step 3: Launch Instance from AMI

Select the AMI → click Launch instance from image.



Step 4: Configure Instance Details

Name: WebServer-from-Custom-AMI

Instance type: t3.micro (Free Tier eligible)

Key pair: Choose an existing .pem key for SSH access.

Network settings:

VPC: Default

Subnet: Public

Security Group: Allow HTTP (80) and SSH (22)



Step 5: Storage (EBS Volume)

Keep default root volume (e.g., 8 GB gp3).

Step 6: Review and Launch

Click Launch instance.

Wait until the instance state becomes Running.



Step 7: Access the Website

Copy the Public IPv4 address from the instance details.

Open in a browser → `http://<Public-IP>`

You should see your same website, confirming the custom AMI works.

A new EC2 instance is successfully launched using the Custom AMI, automatically containing the OS, Apache, configurations, and website files — no manual setup required.

Steps to Delete the custom AMI

Deleting a **custom AMI** involves **deregistering** the image and then **deleting its associated EBS snapshot**.

When you deregister an AMI, it is removed from your account and can no longer be used to launch new instances.

However, the **snapshot** that was created along with the AMI still remains in your storage and **continues to incur charges, so you must delete it separately** to free up space and stop costs.

This ensures your AWS environment remains clean and cost-efficient.

Step 1 – Open EC2 Dashboard

Sign in to the AWS Management Console.

Navigate to EC2 service.

In the left navigation pane, scroll down to Images → AMIs.

Step 2 – Locate Your Custom AMI

In the Owned by me tab, you will see all AMIs you created (custom AMIs).

Select the AMI you want to delete.


You can identify it by Name or AMI ID.

Step 3 – Deregister the AMI

Select the AMI → Click on Actions dropdown.

Choose Deregister AMI.

Confirm by clicking Deregister in the pop-up.



- The AMI is now deregistered (unavailable for future instance launches).
- The snapshot is also deleted, freeing up EBS storage and costs.

DATE: 11-11-2025

Exercise 9: Mini Project: Hosting a Multi-Page Website using EC2 and S3

Website Structure:

We'll make a **3-page static website**:

1. **Home Page (index.html)**

- Welcome message
- Navigation bar linking to other pages
- Short intro about the site

2. Gallery Page (gallery.html)

- Dynamically loads and displays images hosted in an S3 bucket
- Uses public S3 URLs or signed URLs

3. About Page (about.html)

- Information about the project and AWS services used

Step 1: Create a S3 Bucket and upload the images

1. Go to AWS Console → S3 → Create bucket

- Name: <bucket-name>
- Region: same as your EC2 instance
- Uncheck “Block all public access” (for demo)

2. Upload a few sample images.

3. Under each image → Permissions → Make public (or set a bucket policy like this):

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "PublicReadGetObject",
      "Effect": "Allow",
      "Principal": "*",
      "Action": "s3:GetObject",
      "Resource": "arn:aws:s3:::<bucket-name>/*"
    }
  ]
}
```

Amazon S3 > Buckets > ec2-s3-karthik

ec2-s3-karthik

Objects Metadata Properties Permissions Metrics Management Access Points

Objects (3)

Find objects by prefix

Name	Type	Last modified	Size	Storage class
logo.png	png	November 1, 2023, 14:38:08 (UTC+05:30)	2.5 KB	Standard
spring_boot.png	png	November 2, 2023, 14:38:09 (UTC+05:30)	3.8 KB	Standard
spring.png	png	November 2, 2023, 14:38:09 (UTC+05:30)	4.2 KB	Standard

Amazon S3 > Buckets > ec2-s3-karthik > Edit bucket policy

Edit bucket policy

Bucket policy

The bucket policy, written in JSON, provides access to the objects stored in the bucket. Bucket policies don't apply to objects owned by other accounts. [Learn more](#)

Bucket ARN
arn:aws:s3:::ec2-s3-karthik

Policy

```
1 {
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Sid": "PublicReadGetObject",
6       "Effect": "Allow",
7       "Principal": "*",
8       "Action": "s3:GetObject",
9       "Resource": "arn:aws:s3:::ec2-s3-karthik/*"
10    }
11  ]
12 }
```

Edit statement

Select a statement

Select an existing statement in the policy or add a new statement.

+ Add new statement

4. Note down the public URLs:

https://ec2-s3-karthik.s3.ap-northeast-1.amazonaws.com/Spring_Boot.png


<https://ec2-s3-karthik.s3.ap-northeast-1.amazonaws.com/Java.png>

<https://ec2-s3-karthik.s3.ap-northeast-1.amazonaws.com/Spring.png>

Step 2: Launch an EC2 Instance

1. Launch an EC2 Instance

- Amazon Linux 2 or Ubuntu
- Security group → open ports 22 (SSH) and 80 (HTTP)



2. Connect via SSH

```
ssh -i your-key.pem ec2-user@<EC2-Public-IP>
```

3. Install Apache

- `sudo yum update`
- `sudo yum install httpd -y`
- `sudo systemctl start httpd`
- `sudo systemctl enable httpd`

4. Check in browser:

<http://<EC2-Public-IP>> → should show Apache test page.

Step 3: Deploy Website Files

1. Change pwd to **/var/www/html** and create html files

- **cd /var/www/html/**

➤ **sudo nano index.html**

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <title>CloudGallery - Home</title>
  <link rel="stylesheet" href="style.css">
</head>
<body>
  <header>
    <h1>Welcome to CloudGallery</h1>
    <nav>
      <a href="index.html">Home</a> |
      <a href="gallery.html">Gallery</a> |
      <a href="about.html">About</a>
    </nav>
  </header>
  <section>
    <h2>Your photos, served from the cloud ☁</h2>
    <p>This website is hosted on an AWS EC2 instance and fetches images from an AWS S3 bucket.</p>
  </section>

  <footer>
    <p>© 2025 CloudGallery | Powered by AWS</p>
  </footer>
</body>
</html>
```

➤ **sudo nano gallery.html**

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <title>CloudGallery - Gallery</title>
  <link rel="stylesheet" href="style.css">
</head>
<body>
  <header>

</header>
  <section>
    <h2>Images Fetched from S3</h2>
    <div class="gallery">
      
      
      
    </div>
  </section>
  <footer>
    <p>Images served from AWS S3</p>
  </footer>
</body>
</html>
```

➤ sudo nano about.html

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <title>About - CloudGallery</title>
  <link rel="stylesheet" href="style.css">
</head>
<body>
  <header>
    <h1>About</h1>
  <nav>
    <a href="index.html">Home</a> |
    <a href="gallery.html">Gallery</a> |
    <a href="about.html">About</a>
  </nav>
</header>

  <section>
    <h2>About This Project</h2>
    <p>
      CloudGallery is a static website hosted on an AWS EC2 instance. It demonstrates integration between
      EC2 and S3 by fetching image assets directly from a public S3 bucket.
    </p>
    <p>
      This project highlights the simplicity of using AWS services together for scalable and cost-effective hosting.
    </p>
  </section>

  <footer>
    <p>© 2025 CloudGallery</p>
  </footer>
</body>
</html>
```

➤ sudo nano style.css

```
body {
  font-family: Arial, sans-serif;
  background: #f7f9fb;
  color: #333;
}

header {
  background: #0073bb;
```

```

color: white;
padding: 1rem;
text-align: center;
}

nav a {
color: white;
margin: 0 10px;
}

section {
padding: 2rem;
text-align: center;
}

.gallery {
display: flex;
justify-content: center;
flex-wrap: wrap;
gap: 15px;
}

.gallery img {
width: 300px;
height: 200px;
border-radius: 10px;
box-shadow: 0 0 10px rgba(0,0,0,0.2);
}

footer {
background: #222;
color: white;
text-align: center;
padding: 1rem;
}

box-shadow: 0 0 10px rgba(0,0,0,0.2);
}

footer {
background: #222;
color: white;
text-align: center;
padding: 1rem;
}

```

Step 4: Adjust Permissions & Test



➤ `sudo chmod -R 755 /var/www/html`

Open your browser and visit:

➤ **http://<EC2-Public-IP>/index.html**

You should see:

- Home page
- Gallery fetching S3 images
- About page



DATE: 12-11-2025

Exercise 10: Creation and Configuration of a Custom AWS Virtual Private Cloud (VPC)

[Including Public and Private Subnets, Internet Gateway, NAT Gateway, Route Tables]

Virtual Private Cloud [VPC]

It is a logically isolated section of the AWS Cloud where you can launch your AWS resources (like EC2 instances, databases, etc.) in a customized, secure network environment — similar to having your own private data center inside AWS.

It is a virtual network dedicated to your AWS account.

It gives you complete control over your networking environment, including IP address ranges, subnets, route tables, and network gateways.

Components of a VPC

Component	Description
CIDR Block (IP Range)	The range of IP addresses for your VPC (e.g., 10.0.0.0/16).
Subnets	Smaller divisions inside your VPC; can be Public (accessible from internet) or Private (internal only).
Internet Gateway (IGW)	Allows internet access for resources in public subnets.
Route Tables	Define how traffic is directed between subnets and gateways.
NAT Gateway / NAT Instance	Enables instances in private subnets to connect to the internet without being exposed.
Security Groups	Virtual firewalls that control inbound/outbound traffic at the instance level.
Network ACLs (Access Control Lists)	Additional firewall at the subnet level.
VPC Peering	Connects two VPCs so they can communicate privately.

Default VPC and Custom VPC

Characteristics of Default VPC

When you first create an AWS account, AWS automatically creates a default VPC for you in each region.

Feature	Description
----------------	--------------------

Best For	Beginners, quick testing, learning environments, or temporary setups.
Created Automatically	One Default VPC per AWS Region (created by AWS).
Ready to Use	You can launch EC2 instances immediately — no setup needed.
CIDR Block	Always uses 172.31.0.0/16.
Subnets	One default subnet in each Availability Zone within the region.
Internet Connectivity	Each default subnet is a public subnet (has a route to Internet Gateway).
Route Table	Already configured to connect to the Internet Gateway.
Security Groups & NACLs	Default ones are automatically created and allow basic communication.

Characteristics of Custom VPC

A Custom VPC is created manually by the user to have complete control over the network configuration.

Feature	Description
Best For	Production environments, enterprise setups, or multi-tier architectures.
Created Manually	You define the VPC and its settings yourself.
CIDR Block	You can choose your own IP range (e.g., 10.0.0.0/16).
Subnets	You decide how many subnets, and whether they are public or private.
Internet Connectivity	You attach your own Internet Gateway.
Route Tables	Must be created and configured manually.
Security	You can create custom Security Groups and NACLs as needed.

Subnets

- Subdivisions inside a VPC, used to organize resources.
- Each subnet belongs to one Availability Zone.
- Two types:
 - **Public Subnet** → Connected to Internet Gateway; for web servers.
 - **Private Subnet** → No direct internet access; for databases or internal apps.
- CIDR examples:
 - Public: 10.0.1.0/24
 - Private: 10.0.2.0/24

Internet Gateway (IGW)

- A gateway that connects your VPC to the Internet.

- Required for instances in a public subnet to receive internet traffic.
- Must be attached to the VPC and referenced in the route table.
- Supports bi-directional communication (inbound and outbound).

NAT Gateway

- Enables outbound internet access for private subnet instances.
- Allows downloads and updates (e.g., OS patches) without exposing private IPs.
- Deployed inside a public subnet.
- Needs an Elastic IP for internet access.
- Traffic is one-way: private → internet only (not vice versa).

Route Tables

- Define rules (routes) that determine where network traffic goes.
- Each subnet must be associated with one route table.
- Common routes:
 - For public subnet → 0.0.0.0/0 → Internet Gateway
 - For private subnet → 0.0.0.0/0 → NAT Gateway
- Ensures proper separation between public and private networks.

Security Groups

- **Instance-level firewalls** controlling inbound and outbound traffic.
- Stateful: if inbound traffic is allowed, corresponding outbound is automatically allowed.
- Rules are based on protocol, port number, and source/destination.
- Example:
 - WebServer-SG: allows HTTP(80), HTTPS(443), SSH(22)
 - Database-SG: allows MySQL(3306) from WebServer-SG only

Network ACL (Access Control List)

- **Subnet-level firewall**, acts as an additional layer of security.
- **Stateless**: inbound and outbound rules must be defined separately.
- Default NACL allows all traffic; custom NACLs can be restrictive.
- Used for fine-grained control or compliance environments.

EC2 Instances

- Virtual machines running inside your subnets.
- Public subnet → hosts Web Server (Apache).

- Private subnet → hosts Database Server (MySQL).
- Public EC2s get a public IP; private ones use private IPs only.
- Controlled by their Security Groups and Route Tables.

Elastic IP (EIP)


- A **static, public IPv4 address** that can be attached to an EC2 or NAT Gateway.
- Remains constant even if the instance is stopped or restarted.
- Useful for **NAT Gateways**, load balancers, or fixed server access.

Example: Two-Tier Architecture

Tier	Subnet	Function	Internet Access
Web Tier	Public Subnet	Hosts front-end web server	Yes (via IGW)
Database Tier	Private Subnet	Hosts back-end database	Outbound only (via NATGW)

Create a new VPC for a web application:

- One **Public Subnet** for web servers (via Internet Gateway)
- One **Private Subnet** for databases (via NAT Gateway)
- Custom route tables and tighter security rules.



Objective


To design and configure a **Virtual Private Cloud (VPC)** in AWS with:

- One **Public Subnet** for web servers (via **Internet Gateway**)
- One **Private Subnet** for databases (via **NAT Gateway**)
- Separate **Route Tables** for public and private subnets
- Custom **Security Groups** to enforce network isolation

Step 1 – Create a New VPC

- Open AWS Management Console → VPC.
- Click Create VPC.
- Choose VPC only option.
- Enter:
 - Name → MyWebAppVPC
 - IPv4 CIDR block → 10.0.0.0/16
 - Tenancy → Default
- Click Create VPC.

This allocates a private IP range for your virtual network.



Step 2 – Create Subnets

We'll create two subnets inside the VPC.


(a) Public Subnet

- Go to Subnets → Create subnet.
- Select VPC: MyWebAppVPC.
- Choose Availability Zone A.
- Enter:
 - Name → PublicSubnet
 - IPv4 CIDR block → 10.0.1.0/24
- Click Create subnet.

(b) Private Subnet

- Click Create subnet.

- Select same VPC.
- Choose Availability Zone B.
- Enter:
 -
 - Name → PrivateSubnet
 - IPv4 CIDR block → 10.0.2.0/24
- Click Create subnet.



Step 3 – Create and Attach an Internet Gateway

- In the left navigation pane, scroll down and click on “Internet Gateways”.
- Internet Gateways → Create Internet Gateway.
 - Name: MyWebApp-IGW
- Click Create Internet Gateway

At this point, the IGW exists but is not yet connected to your VPC.

- **Attach the Internet Gateway to Your VPC**
 - Select the IGW you just created (checkbox).
 - Click on the “Actions” drop-down.
 - Choose “Attach to VPC.”
 - From the list, select your **VPC name**, e.g. MyWebAppVPC.
 - Click “Attach Internet Gateway.”

Now the IGW is linked to your VPC and the public subnet can reach the internet.


Step 4 – Configure Route Table for IGW

Public Route Table

- Go to Route Tables → Create route table.
 - Name → PublicRT
 - VPC → MyWebAppVPC

- Under Routes → Edit routes → Add route
 - Destination: 0.0.0.0/0
 - Target: Internet Gateway (MyWebApp-IGW)
- Under Subnet Associations → Edit subnet associations → Select PublicSubnet → Save.

Now the PublicSubnet has internet access.




Step 5 – Create a NAT Gateway

NOTE: NAT Gateway must always be created in a public subnet

- Go to NAT Gateways → Create NAT Gateway.
- Name → NAT-GW
- Choose:
 - Subnet → PublicSubnet
 - Elastic IP Allocation ID → Click Allocate Elastic IP
- Click Create NAT Gateway.

This allows instances in private subnet to access the internet (e.g., for updates) without being exposed.

Note: NAT Gateway is a paid resource.




Step 6 – Configure Route Table for IGW

Private Route Table

- Create another Route Table → Name PrivateRT.
- Under Routes → Edit routes → Add route
 - Destination: 0.0.0.0/0
 - Target: NAT Gateway (MyWebApp-NATGW)
- Under Subnet Associations → Select PrivateSubnet → Save.

PrivateSubnet traffic goes through the NAT Gateway.



Note:

Each subnet in a VPC can be associated with **only one route table**.


If a subnet is not explicitly associated with a custom table, it will automatically use the **Main Route Table** created by default with the VPC.

DATE: 14-11-2025

Exercise 11: Launching and Configuring EC2 Instances for Web Server in Public Subnet and Database Server in Private Subnet Using NAT Gateway for Outbound Access

Step 1 — Launch Windows instance in Public Subnet


- Go to EC2 → Launch Instance.
- Choose:
 - AMI → Windows Server (Free tier eligible)
 - Instance type → t2.micro / t3.micro
- Select:
 - VPC → MyWebAppVPC
 - Subnet → PublicSubnet
 - Auto assign Public IP → Enable
- Security Group:
 - Create WebInstance-SG
 - Allow:
 - RDP (3389) → Anywhere (for practice)
- Launch instance using key pair (.pem).



Step 2 — Launch Windows instance in Private Subnet

- Click Launch Instance.
- Choose:
 - Windows Server AMI
- Select:
 - VPC → MyWebAppVPC
 - Subnet → PrivateSubnet
 - Auto assign Public IP → Disable
- Security Group:
 - Create DBInstance-SG
 - Allow RDP only from WebInstance-SG
- Launch instance.

The DBInstance is now **fully isolated** and cannot be accessed directly from the internet.



Step 3 — Connect to Public Windows Instance

1. Select WebInstance → Click Connect.

2. Choose RDP Client tab.
3. Click Get Password.
4. Upload your .pem key pair.
5. It shows the decrypted Windows Administrator password.
6. Copy the Public IP into a document for further use.

Connect using RDP

- Open Remote Desktop Connection (mstsc).
- In the dialog:
 - Computer → Public IP of WebInstance
- Click Connect.
- Credentials dialog:
 - Username → Administrator
 - Password → Paste decrypted password
- Click OK.

You are now inside the public Windows instance.

Step 4 — connect to private windows instance


We cannot connect directly from your laptop → No Public IP.

We must connect **from within WebInstance** (jump server / bastion host).

Remote Login from WebInstance → DBInstance

- While logged into WebInstance, open Remote Desktop Connection. **[mstsc]**
- Enter:
 - Computer → Private IP of DBInstance
- Click Connect.
- Credentials:
 - Username → Administrator
 - Password → Paste the decrypted password (same key pair)
- Click OK.

You are now inside the private Windows instance.



Step 5 — testing NAT gateway connectivity (Verify Internet Access)

From WebInstance (Public Subnet)

- Open a browser → Internet should work (through IGW).


From DBInstance (Private Subnet)

- Open browser → Internet should also work (through NAT Gateway).
- Internet icon will show "Internet Access".

Security Check

- DBInstance cannot receive inbound traffic from internet.
- Only outbound is allowed via NAT (safe for DB servers).

Test: From DBInstance (Private Subnet) → ping Google (Outbound Allowed)



Steps

Connect to DBInstance (using RDP from WebInstance).

Open Command Prompt inside DBInstance.

Type: ping google.com

Expected Result

- It will successfully ping google.com
- You will see replies like:
Reply from 142.250.xxx.xxx: bytes=32 time=20ms TTL=115

Why does this work?

- Outbound traffic is allowed from Private Subnet → NAT Gateway → Internet.
- NAT Gateway acts as a proxy for the private instance.
- So the private instance can reach internet, but internet cannot reach the private instance.

What will NOT work?

From the Private DBInstance: ping <Your Own Public IP>
or anyone trying to ping: ping <DBInstance Private IP>

Both will fail because:

- Inbound traffic is blocked
- DBInstance has no Public IP
- SG allows inbound only from WebInstance-SG

Private instance → internet = YES (via NAT Gateway)

Internet → private instance = NO (fully blocked)

Because NAT Gateway is **one-way** only.

Elastic IP address

- A static IP address is a permanent address that doesn't change. You can manually configure a device to have a static IP address.
- An Elastic IP address is static and has to be used in a specific Region, it cannot be moved to a different Region.
- An Elastic IP address comes from Amazon's pool of IPv4 addresses.
- To use an Elastic IP address, you first allocate one to your account, and then associate it with your instance or a network interface.
- When you associate an Elastic IP address with an instance or its primary network interface, if the instance already has a public IPv4 address associated with it, that public IPv4 address is released back into Amazon's pool of public IPv4 addresses and the Elastic IP address is associated with the instance instead.
- You can disassociate an Elastic IP address from a resource, and then associate it with a different resource.
- A disassociated Elastic IP address remains allocated to your account until you explicitly release it.
- You are charged for all Elastic IP addresses in your account, regardless of whether they are associated or disassociated with an instance.
- Static IP addresses are especially important in cases where a device has to be quickly found over the internet on a permanent basis.
- Web Servers: A website must have one or more static IP addresses to be assigned to the domain always point to the correct server.

DATE: 19-11-25

Exercise 12: Deploying and Testing a Load-Balanced Web Application

By Creating a Web Server, building a Custom AMI, configuring a Target Group, Launching an Application Load Balancer (ALB), Registering the Instance in the Target Group, and Testing the ALB DNS

Objective

To deploy a scalable and highly available web application on AWS by configuring:

- EC2 web server
- Custom AMI
- Target Group
- Application Load Balancer (ALB)
- Testing the ALB DNS

Description for Each Component

EC2 Web Server

What: A virtual Linux machine that hosts your web application files.

Why: It serves the actual website content that users access through the ALB.

Custom AMI

What: A snapshot/template of your configured EC2 instance (with Apache + website files).

Why: Ensures every new instance launched by the Auto Scaling Group has the same setup automatically.

Target Group

What: A collection of EC2 instances that the ALB sends traffic to.

Why: It allows the load balancer to forward requests only to healthy instances in the Auto Scaling Group.

Application Load Balancer (ALB)

What: A managed service that distributes incoming HTTP traffic across multiple EC2 instances.

Why: Ensures high availability and balanced traffic distribution for the web application.

Security Group

What: A virtual firewall controlling inbound/outbound traffic to EC2 instances and ALB.

Why: Ensures only required traffic (HTTP/SSH) is allowed for the application components.

VPC + Subnets

What: A private network environment where all resources run.

Why: Provides isolation, routing, and a structured network setup for public & private components.

Launch Template

What: A reusable blueprint containing AMI, instance type, key pair, and security group settings.

Why: The ASG uses this template to launch identical EC2 instances.

STEP 1 — Launch Base EC2 Instance

- Open AWS Console → EC2 → Launch Instance
- Name: Base-WebServer
- AMI: Amazon Linux 2
- Instance type: t3.micro
- Key Pair: your .pem file
- Security Group:
 - HTTP (80) → Anywhere
 - SSH (22) → My IP
- Launch the instance

STEP 2 — Install Apache and Create Web Page

- SSH into the instance.
- Install Apache:
 - `sudo yum install httpd -y`
 - `sudo systemctl start httpd`
 - `sudo systemctl enable httpd`
- Move to Web Root Folder - `cd /var/www/html`
- Create a webpage with hostname:
 - `echo "<h1>Hello from Instance 1 — $(hostname)</h1>" | sudo tee /var/www/html/index.html`

Test in browser using the instance's public IP.

STEP 3 — Create Custom AMI

- Go to EC2 → Instances
- Select Base-WebServer
- Actions → Image and Templates → Create Image
- Name: WebServer-AMI
- Create
- Wait until AMI status = Available



This AMI now contains Apache + your index.html.

STEP 4 — Create Target Group

- EC2 → Target Groups → Create Target Group

- Target type: Instances
- Name: WebApp-TG
- Protocol: HTTP
- Port: 80
- Health Check Path: /
- Create (do not register instances manually)


Purpose: Target Group holds the list of EC2 instances behind the Load Balancer.



STEP 5 — Create Application Load Balancer (ALB)


- EC2 → Load Balancers → Create Load Balancer
- Choose Application Load Balancer
- Name: WebApp-ALB
- Scheme: Internet-facing

- Listeners: HTTP (80)
- Select two public subnets
- Security Group: allow HTTP (port 80)
- Forward to Target Group → WebApp-TG
- Create



STEP 6 — Test the Load Balancer Using ALB DNS

- Go to AWS Console → EC2 → Load Balancers
 - Select your ALB: ALB-WebServer
- Copy the ALB DNS Name
 - You will see something like:
 - WebApp-ALB-123456789.ap-south-1.elb.amazonaws.com
- Open a Browser and Paste the DNS Name
 - Enter: `http://<your-alb-dns-name>`
 - Example: `http://WebApp-ALB-123456789.ap-south-1.elb.amazonaws.com`
- View the Output
 - You should see the webpage you created earlier:
 - Hello from Instance 1 — ip-xx-xx-xx-xx
 - This confirms that: The ALB is working
 - The Target Group is routing traffic
 - The instance created from your custom AMI is serving the web page
- Optional – Refresh Multiple Times
 - If you later use Auto Scaling with multiple instances:
 - Clicking Refresh will show different hostnames
 - This confirms load balancing across multiple servers



Hello from ip-172-31-35-81.ap-northeast-1.compute.internal

DATE: 20-11-25

Exercise 13: Stress Testing a Linux EC2 Instance (CPU Load Test)

Objective:

To generate high CPU load on an Amazon EC2 Linux instance using the stress-ng tool and observe CPU utilization in CloudWatch.

Step 1 — Launch a Linux EC2 Instance

1. Log in to the AWS Console.
2. Go to **EC2 → Instances → Launch Instance**.
3. Name: **StressTest-EC2**
4. AMI: **Amazon Linux 2023 (Free Tier eligible)**
5. Instance type: **t2.micro / t3.micro**
6. Key pair: Select or create a .pem key.
7. Security Group:
 - Allow **SSH (22)** from My IP.
 - Allow **HTTP (80)** from anywhere (optional).
8. Launch the instance.

Step 2 — Connect to the EC2 Instance

Use Windows PowerShell

Navigate to the folder where pem file is located and type the following command:

```
ssh -i "your-key.pem" ec2-user@<public-ip>
```

Step 3 — Install Apache (Optional, to verify the instance is working)

```
sudo yum install httpd -y
```

```
sudo systemctl start httpd
```

```
sudo systemctl enable httpd
```

Create a test web page:

```
echo "<h1>EC2 Stress Test Demo — $(hostname)</h1>" | sudo tee /var/www/html/index.html
```

STEP 4 — Install and Run the Stress Testing Tool (stress-ng)

(For Amazon Linux 2023 EC2 Instances)

Move to the ec2-user home directory

(It is recommended to run stress-ng from the home folder)

`cd ~`

Install stress-ng

Amazon Linux 2023 includes stress-ng directly through dnf, so install it using:

`sudo dnf install stress-ng -y`

Run a CPU Stress Test

Generate high CPU load using 4 CPU workers for 2 minutes:

`stress-ng --cpu 4 --timeout 120`

Expected Result

- Terminal becomes busy during the 120-second load test
- After completion, you will see:

successful run completed in 120.02s

- CPU usage will spike to 90–100% (you can check using top)

```
ec2-user@ip-172-31-17-191: ~
Installing:
stress-ng          x86_64          0.15.05-1.amzn2023          amazonlinux          2.3 M
Installing dependencies:
x86_64          1.0.5-25.amzn2023.0.3          amazonlinux          153 k
Judy            x86_64          0.10.0-7.amzn2023.0.2          amazonlinux          109 k
libbsd         x86_64          1.0.10-9.amzn2023.0.3          amazonlinux          92 k
lkctp-tools    x86_64          1.0.18-9.amzn2023.0.3          amazonlinux          92 k

Transaction Summary
=====
Install 4 Packages
Total download size: 2.7 M
Installed size: 9.7 M
Downloading Packages:
(1/4): lkctp-tools-1.0.18-9.amzn2023.0.3.x86_64.rpm          1.1 MB/s | 92 kB  00:00
(2/4): libbsd-0.10.0-7.amzn2023.0.2.x86_64.rpm             1.1 MB/s | 109 kB 00:00
(3/4): Judy-1.0.5-25.amzn2023.0.3.x86_64.rpm               2.4 MB/s | 153 kB 00:00
(4/4): stress-ng-0.15.05-1.amzn2023.x86_64.rpm             2.4 MB/s | 2.3 MB 00:00
-----
Total: 12 MB/s | 2.7 MB  00:00

Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
Preparing : lkctp-tools-1.0.18-9.amzn2023.0.3.x86_64          1/1
Installing : libbsd-0.10.0-7.amzn2023.0.2.x86_64          2/4
Installing : Judy-1.0.5-25.amzn2023.0.3.x86_64            3/4
Installing : stress-ng-0.15.05-1.amzn2023.x86_64          4/4
Running scriptlet: stress-ng-0.15.05-1.amzn2023.x86_64    4/4
Verifying : Judy-1.0.5-25.amzn2023.0.3.x86_64            1/4
Verifying : libbsd-0.10.0-7.amzn2023.0.2.x86_64          2/4
Verifying : lkctp-tools-1.0.18-9.amzn2023.0.3.x86_64      3/4
Verifying : stress-ng-0.15.05-1.amzn2023.x86_64          4/4

Installed:
Judy-1.0.5-25.amzn2023.0.3.x86_64 libbsd-0.10.0-7.amzn2023.0.2.x86_64 lkctp-tools-1.0.18-9.amzn2023.0.3.x86_64 stress-ng-0.15.05-1.amzn2023.x86_64

Complete!
[ec2-user@ip-172-31-17-191 ~]$ stress-ng --cpu 4 --timeout 120
stress-ng: info: [25013] setting to a 120 second (2 mins, 0.00 secs) run per stressor
stress-ng: info: [25013] dispatching hogs: 4 cpu
```

Step 5 — Run Stress Test (CPU Load Generation)

Move to home directory to avoid permission issues:

`cd ~`

Run CPU stress for 2 minutes using 4 CPU workers:

`stress-ng --cpu 4 --timeout 120`

You should see messages like:

stress-ng: info: dispatching hogs: 4 cpu

stress-ng: info: successful run completed in 120.02s

This will increase CPU usage to ~100% for the duration.

Step 6 — Observe CPU Utilization in CloudWatch

Go to CloudWatch → Metrics → EC2 → Per-Instance Metrics

Select:

CPUUtilization → InstanceId → your instance

View graph in 1-minute or 5-minute intervals.

You should see a sharp spike during the 2-minute stress period.

Step 7 — Stop or Terminate the Instance


To avoid charges:

Option A — Stop the instance (safe):

Actions → Instance State → Stop

Option B — Terminate (permanently delete):

Actions → Instance State → Terminate



Expected Output

- CPU Utilization in CloudWatch should reach 90–100%.
- Stress test completes without errors.
- Students understand how CPU load affects EC2 metrics.

DATE: 22-11-25

Exercise–14: Mini Project

Deploying a Load-Balanced Web Application using Application Load Balancer (ALB), Auto Scaling Group (ASG), Custom AMI, and Target Group on AWS with CloudWatch Alarms & Stress Testing for Auto Scaling Validation.

Auto Scaling Group (ASG)

What: A service that automatically launches or terminates EC2 instances based on demand.

Why: Provides scalability so your application can handle variable traffic.

CPU-Based Scaling Policies

What: Rules that add or remove EC2 instances based on CPU usage thresholds.

Why: Ensures your application automatically scales up during heavy load and scales down to save cost.

Stress Testing (using stress/stress-ng tool)

What: A method to artificially increase CPU load on instances.

Why: Used to validate whether the Auto Scaling Group is scaling up/down correctly.

STEP 7A — Create Launch Template

Go to EC2 → Launch Templates → Create launch template

Template Name

- Launch template name: LT- WebServer

Choose AMI (Custom AMI)

Under Application and OS Images (AMI):

- Click My AMIs
- Select your custom AMI: AMI-WebServer (the AMI you created earlier)

Instance Type

Under Instance type: Select t3.micro

Key Pair

Under Key pair (login): Select your existing key pair (pemfile1.pem)

Security Group

Under Network settings → Firewall (security groups):

- Choose Select existing security group
- Select a security group that allows:
 - SSH (Port 22)
 - HTTP (Port 80)

Example: launch-wizard-1 (already has both rules)

Storage

Leave default: 8 GiB gp3 (root volume)

Create Template

Click Create launch template

Launch Template is now ready to be used by the Auto Scaling Group.

STEP 7B — Create Auto Scaling Group (ASG)

(Using the Launch Template you created in Step 7A)

Open ASG Wizard

Go to EC2 → Auto Scaling Groups → Create Auto Scaling Group

Name and Choose Launch Template

Auto Scaling group name: ASG-WebServer

Launch template: Choose LT-WebServer

Click Next.

Select Network (VPC + Subnets)

Since we are using default VPC, select:

VPC: vpc-0c952a6cabfbe594b (Default VPC)

Subnets (select ANY 2)

Example: ap-south-1a and ap-south-1b

(These are public subnets in default VPC — good for web servers)

Click Next.

Attach Load Balancer

Under Load balancing options:

Select: Attach to an existing load balancer

Under Select load balancers to attach:

Choose: Choose from your load balancer target groups

Under dropdown:

Select your Target Group: TG-WebServer (or whatever name you created)

AWS will automatically show:

- Load balancer: ALB-WebServer
- Type: Application/HTTP

Click Next.

Configure Group Size

Set:

Desired capacity = 1

Minimum capacity = 1

Maximum capacity = 3

We are not adding scaling policies now (stress test comes later), so:

Select No scaling policies

Click Next.

Skip Notifications and Tags

No changes.

Click Next.

Review and Create

Check:

- Launch template = LT-WebServer
- Load balancer = ALB-WebServer
- Target group = TG-WebServer
- Subnets = 2 selected
- Desired = 1 instance

Then click: Create Auto Scaling Group

ASG Creation Done

Check if ASG launched an instance.

Verify ASG Instance Creation

Go to EC2 → Auto Scaling Groups → ASG-WebServer

Open it and check: Under Activity tab

You should see: Successful — Launching a new EC2 instance

Under Instances tab: You should see an instance like: i-xxxxxxxxxxxxx (Running)

This instance was launched automatically by the ASG using custom AMI through the Launch Template.

Now that ASG is created and one instance is running, we will add Auto Scaling Policies so the group can:

- Scale OUT (add more instances when CPU is high)
- Scale IN (remove instances when CPU is low)

STEP 8 — Configure Auto Scaling Policies (Scale Out & Scale In)

PART 1 – Create High-CPU Alarm (for Scale-OUT)

Open CloudWatch and start alarm

Go to CloudWatch → left menu Alarms → All alarms.

Click Create alarm.

Choose metric (CPU of ASG)

Click Select metric.

Navigate:

EC2 → By Auto Scaling Group (or By AutoScalingGroupName).

Click your group ASG-WebServer.

Select CPUUtilization.

Click Select metric.

Set condition (CPU > 60)

Statistic: Average

Period: 1 minute

Under Conditions:

Threshold type: Static

“Whenever CPUUtilization is...” → choose Greater

“than...” → 60

In Additional configuration, datapoints to alarm: 1 out of 1 (default is fine).

Click Next.

Skip actions (no SNS, no ASG here)

On Configure actions page, do NOT add anything.

Just click Next.

If a popup says “*No actions configured*”, click Proceed without actions.

Name and create

Alarm name: ASG-CPU-High-60

Description (optional): CPUUtilization > 60% for 1 minute (ASG-WebServer)

Click Next, then Create alarm.

You now have one metric alarm with no actions.

PART 2 – Attach High-CPU alarm to ASG as Scale-OUT Policy

Open ASG automatic scaling

Go to EC2 → left menu Auto Scaling groups.

Click ASG-WebServer.

Go to the Automatic scaling tab.

Create Scale-OUT policy

Click Create dynamic scaling policy.

Policy type: Simple scaling

Scaling policy name: ScaleOut-CPU60

CloudWatch alarm:

Click the dropdown and select ASG-CPU-High-60.

Take the action:

Action: Add

Value: 1

Unit: capacity units
And then wait: 300 seconds (default is fine).
Click Create.

Now your ASG knows:
When alarm ASG-CPU-High-60 goes to ALARM → Add 1 instance.

PART 3 – Create Low-CPU Alarm (for Scale-IN)

Create another alarm in CloudWatch

Go again to CloudWatch → Alarms → Create alarm.

Select metric exactly like before:

EC2 → By Auto Scaling Group → ASG-WebServer → CPUUtilization → Select metric.

Set condition (CPU < 20)

Statistic: Average

Period: 1 minute

Conditions:

Threshold type: Static

“Whenever CPUUtilization is...” → Lower

“than...” → 20

Datapoints to alarm: 1 out of 1 (default).

Click Next.

Skip actions again

Do not add SNS / Auto Scaling / EC2 actions.

Click Next → if popup appears, Proceed without actions.

Name and create

Alarm name: ASG-CPU-Low-20

Description (optional): CPUUtilization < 20% for 1 minute (ASG-WebServer)

Next → Create alarm.

PART 4 – Attach Low-CPU alarm to ASG as Scale-IN Policy

Open ASG automatic scaling again

Back to EC2 → Auto Scaling groups → ASG-WebServer.

Go to Automatic scaling tab.

Create Scale-IN policy

Click Create dynamic scaling policy.

Policy type: Simple scaling

Scaling policy name: scalein-cpu20

CloudWatch alarm:

Choose ASG-CPU-Low-20 from the dropdown.

Take the action:

Action: Remove

Value: 1

Unit: capacity units

And then wait: 300 seconds.

Click Create.

CHECK the resources created in STEP 8-

In EC2 → Auto Scaling groups → ASG-WebServer → Automatic scaling

- You should see:
- ScaleOut-CPU60 – Simple scaling – Add 1 capacity units – Alarm: ASG-CPU-High-60
- scalein-cpu20 – Simple scaling – Remove 1 capacity units – Alarm: ASG-CPU-Low-20

In CloudWatch → Alarms

- ASG-CPU-High-60
- ASG-CPU-Low-20
- Both will show Actions: No actions – this is OK because ASG is using them.

STEP 9 — Perform Stress Testing for Auto Scaling Validation

To artificially increase CPU load on EC2 instances and observe automatic scaling behavior triggered by CloudWatch alarms.

Part A — Connect to the ASG-launched EC2 Instance

Go to AWS Console → EC2 → Instances

Select the instance launched by ASG

(NOT the original base instance you created manually)

Click Connect

Choose EC2 Instance Connect → Connect

Part B — Install Stress Testing Tool

For Amazon Linux 2023:

```
sudo dnf install stress-ng -y
```

Confirm installation:

```
stress-ng --version
```

Part C — Apply CPU Load (Trigger Scale Out)

Run stress tool for 2 minutes with 4 CPU workers:

stress-ng --cpu 4 --timeout 120

- ✓ This will spike CPU usage above **60%**
- ✓ Your Scale-Out Policy should activate
- ✓ CloudWatch alarm → ALARM state → ASG launches 1 more instance

Part D — Monitor Auto Scaling Activity

Check scaling progress in:

Service	What to Monitor
EC2 → Auto Scaling Groups → Activity	Shows launching of new instance
EC2 → Instances	New instance appears with different hostname
CloudWatch → Alarms	High CPU alarm turns ALARM

Scaling may take **2–4 minutes**

Refresh page to see updates

Part E — Test Load Balancing

Copy ALB DNS Name → paste in browser → refresh multiple times:

You should see:

Hello from Instance 1 — ip-xx-xx-xx-xx

Hello from Instance 2 — ip-aa-aa-aa-aa

- ✓ Confirms ALB is distributing traffic across multiple instances

Part F — Scale In (CPU comes down)

After 2 minutes, stress test stops automatically → CPU drops below 20%

- ✓ Low CPU alarm triggers
- ✓ ASG removes the extra instance
- ✓ Only 1 instance remains (desired capacity)

Scaling-in may take 3–5 minutes

Extra points

This exercise imitates a production-grade architecture used in companies.

What is a Target Group?

A Target Group is a collection of EC2 instances that receive traffic from the Application Load Balancer (ALB).

It defines:

- Which instances to send traffic to
- On which port (example: 80)
- Health check rules (path /, interval, threshold)

In simple words:

Target Group = list of servers behind the load balancer.

Why do we need two steps: Create Launch Template and Create Auto Scaling Group?

Because both have different roles:

Launch Template = WHAT to launch

It contains the configuration of one EC2 instance, such as:

- AMI
- Instance type
- Security Group
- Key pair
- Storage
- User data

This is just a blueprint.

Auto Scaling Group = WHEN & HOW MANY to launch

The ASG uses the Launch Template to automatically:

- Launch new instances
- Remove instances
- Maintain desired capacity
- Scale based on CPU/memory demand

Where to find `http://<alb-dns-name>?`

You can find the ALB DNS Name in the AWS Console:

Path: EC2 → Load Balancers → Select your ALB → Description tab

There you will see: DNS name: mywebapp-alb-12345678.ap-south-1.elb.amazonaws.com

Use it in your browser as: `http://mywebapp-alb-12345678.ap-south-1.elb.amazonaws.com`

This is the public URL of your Load Balancer.

What is ALB DNS?

ALB DNS is the publicly accessible DNS name automatically created by AWS for your Application Load Balancer.

Example: myapp-alb-123456789.ap-south-1.elb.amazonaws.com

When a user types this URL:

- Traffic goes to the ALB
- ALB forwards to Target Group
- Target Group sends request to one of the EC2 instances

ALB DNS = The website URL of your Load Balancer.

DNS means Domain Name System.

DNS is a system that converts domain names to IP addresses.

Example: `www.amazon.com` → `52.95.120.1`

A DNS Server is only one part of this system.

DNS = System that translates names to IPs.

DNS Server = Machine that performs the translation.

Why Do we NEED CloudWatch alarms for ASG scaling exercise?

We are using “Simple Scaling”

- Based on CPU > 60 and CPU < 20
- Which requires manually created CloudWatch alarms

Without the alarms, ASG has NO trigger, so it cannot scale.

Delete these resources in this exact order

To avoid dependency errors:

1. Delete Auto Scaling Group (ASG)

- EC2 → Auto Scaling Groups
- Select ASG-WebServer
- Actions → Delete

This will automatically terminate any instances managed by ASG.

2. Delete Launch Template

- EC2 → Launch Templates
- Select your template
- Actions → Delete template

3. Delete Application Load Balancer

- EC2 → Load Balancers
- Select your ALB
- Actions → Delete

4. Delete Target Group

- EC2 → Load Balancing → Target Groups
- Select your target group
- Actions → Delete

****5. Delete the Custom AMI**

- EC2 → AMIs
- Select your AMI → **Deregister**
- Check Snapshots → delete the related snapshot also.

6. Delete any Manual EC2 Instances and other resources

If you still have any instances running:

- EC2 → Instances → Select → **Terminate**

Resource clean-up NOTE:

- **ALB** charges per hour
- **ASG instances** can launch unexpectedly
- **Custom AMI snapshot** costs storage
- **Running EC2** will cost hourly

DATE: 22-11-25

Exercise–15: Hosting a WordPress Content Management Website on Amazon LightSail.

Amazon LightSail is a beginner-friendly cloud platform in AWS that provides:

- Virtual servers (simplified EC2)
- Fixed monthly pricing (predictable cost)
- Simple UI and easy setup
- One-click application launch (WordPress, LAMP, Node.js, etc.)
- Built-in networking (Static IP, DNS, Firewall)
- Automatic SSH terminal
- Snapshots (backups)

It is ideal for:

- Personal websites
- Student projects
- Small web apps
- WordPress blogs
- Quick prototyping

Why LightSail Instead of EC2?

LightSail = “EC2 with all hard things simplified.”

Feature	EC2	LightSail
Pricing	Pay per hour	Fixed monthly
Setup	Complex (VPC, SG, AMI)	Simple

SSH	Need .pem or PuTTY	Built-in SSH
DNS	Basic	Easy DNS panel
Apps	Install manually	One-click WordPress, LAMP

3. What is WordPress in LightSail?

WordPress is a full-stack content management system (CMS) pre-installed on LightSail.

It includes:

Frontend (UI)

- HTML, CSS, JS
- Themes
- Templates
- Page layouts

Backend

- PHP (server-side logic)
- WordPress core engine

Database

- MySQL/MariaDB
Stores pages, media, lessons, posts, settings.

Server

- Apache web server

Hosting

- LightSail VM + static IP + firewall

Thus, WordPress is a full-stack application (LAMP stack).

4. LightSail Free-Tier Eligibility

- 3 months free
- Only for the 512 MB RAM plan
- Suitable for WordPress, but slightly slow
- After 3 months → ~₹350–₹400/month
- recommended: 1 GB RAM plan.

PART A — Create Server on LightSail

Step 1 — Open LightSail Console

AWS Console → Search “LightSail” → Open it.

LightSail dashboard appears.

Step 2 — Create Instance

Click Create Instance.

Step 3 — Select Platform:

Choose Linux/Unix. (WordPress works best on Linux)

Step 4 — Select Blueprint → App + OS → WordPress

[LightSail will install automatically:

- Apache
- PHP
- MySQL/MariaDB
- WordPress application

No manual setup required.]

Step 5 — Choose Instance Plan

512 MB RAM (free-tier for 3 months but slower)

Step 6 — Name the Instance

Give a name (example MyCMS)

Click Create Instance.

Step 7 — Wait Until Status = Running

Takes 1–2 minutes.

PART B — Access and Configure WordPress

Step 8 — Connect Through Browser-Based SSH

Click **Connect** → **Connect using SSH**

(LightSail gives built-in SSH. No need for .pem file.)

Step 9 — Retrieve WordPress Admin Password

In SSH terminal, run:

```
cat bitnami_application_password
```

This outputs a long randomly generated password.

Copy it — you need it for login. [example: INBV03eMcWX:]

Step 10 — Open Your WordPress Website

Copy the **Public IP** from the instance details.

Visit: `http://<your-lightsail-ip>`

You will see a WordPress homepage.

Step 11 — Login to WordPress Admin

Visit: `http://<your-lightsail-ip>/wp-admin`

Login:

- Username: **user**

- Password: *paste the password from SSH*
- You now have full access to your site's backend.

PART C — Build ‘CMS Website’ [Experiential learning]

Step 12 — Install a Clean Education Theme

Go to:

Dashboard → Appearance → Themes → Add New

Recommended theme: Astra (best)

Activate the theme. [install and activate]

Step 13 — Create Main Website Pages

Dashboard → Pages → Add New

Create:

1. Home
2. Courses & Subjects
3. Lesson Notes
4. Assignments
5. Study Material (PDFs/PPTs)
6. Contact

Step 14 — Create Lesson-Wise Pages

Step 15 — Upload PDFs, PPTs, Notes

Go to:

Dashboard → Media → Add New

Upload:

- PDFs
- DOCX files
- PPTs
- Images
- Notes

Then insert into pages using “Add Media”.

Step 16 — Create Menus

Dashboard → Appearance → Menus

Add:

Home | Courses | Notes | Assignments | Contact

Add submenus under Courses and Notes.

PART D — Secure & Finalize

Step 18 — Attach Static IP

LightSail → Networking → Create Static IP → Attach to instance.

Reason:

Without this, your IP changes → site breaks.

Step 19 — Take Snapshot

LightSail → Snapshots → Create snapshot.

Acts as a backup of your entire site.

DATE: 22-11-25

Exercise–16: Building a Basic Python Flask Web Application — Fundamentals of Web Requests & Form Handling (Pre-Requisite for AWS RDS Connectivity Lab)

[Prerequisite for AWS RDS exercise. This particular exercise is a local Python Flask exercise. No AWS service is used here.]

To design and implement a simple Web Application using Python Flask that:

- Displays an HTML form to accept Name and Password.
- Reads the form data on the server side when the form is submitted.
- Displays the submitted values on a new result page.

After completing this exercise, you will be able to:

- Understand how an HTML form sends data to the server.
- Explain how Flask receives form data using the request object.
- Read request parameters (request.form) in Flask.
- Understand basic frontend–backend communication.
- Handle a simple HTTP POST request in a Flask application.

Folder Structure

Create the following structure manually:

```
FlaskFormApp/  
├── app.py  
├── templates/  
│   ├── form.html  
│   └── result.html
```

Requirements

- Python installed (3.x)
- Flask library
- Any code editor (VS Code / PyCharm / Notepad etc.)
- Web browser (Chrome / Edge / Firefox)

STEP 1: Install Flask

Open **Command Prompt / Terminal** and run:
pip install flask

STEP 2: Create Project Folder

Create a folder, for example:
C:\Users\MCA\FlaskFormApp

STEP 3: Create templates Folder

Inside C:\Users\MCA\FlaskFormApp, create a sub-folder named templates
FlaskFormApp/templates/
All HTML files will be placed in this templates folder.

STEP 4: Create HTML Form — templates/form.html

Create a new file form.html inside templates folder and type the following code:

```
<!DOCTYPE html>
<html>
<head>
  <title>HTML Form</title>
</head>
<body>
  <h2>Enter Your Details</h2>
  <form action="/submit" method="post">
    Name: <input type="text" name="uname"><br><br>
    Password: <input type="password" name="pwd"><br><br>
    <button type="submit">Submit</button>
  </form>
</body>
</html>
```

Note:

- action="/submit" – sends the form data to the /submit route.
- method="post" – form data is sent using HTTP POST method.
- name="uname" and name="pwd" – these names are used in Flask to read values.

STEP 5: Create Result Page — templates/result.html

Create another file result.html inside templates folder:

```
<!DOCTYPE html>
<html>
<head>
  <title>Result</title>
</head>
<body>
  <h2>Form Submission Result</h2>
  <p><strong>Name:</strong> {{ name }}</p>
  <p><strong>Password:</strong> {{ password }}</p>
</body>
</html>
```

Note: {{ name }} and {{ password }} are **placeholders** (Jinja2 template syntax) that Flask will fill with the values sent from app.py.

STEP 6: Create Flask Backend — app.py

In the FlaskFormApp folder, create app.py and type:

```
from flask import Flask, request, render_template

app = Flask(__name__)

@app.route('/')
def home():
    return render_template('form.html')

@app.route('/submit', methods=['POST'])
def submit():
    name = request.form['uname']
    password = request.form['pwd']
    return render_template('result.html', name=name, password=password)

if __name__ == "__main__":
    app.run(debug=True)
```

Explanation:

- `@app.route('/') →` Home URL, shows the form using form.html.
- `@app.route('/submit', methods=['POST']) →` Handles form submission.
- `request.form['uname'] →` Reads the value of input with name="uname".
- `render_template('result.html', name=name, password=password) →` Sends values to result page.

STEP 7: Run the Application

- Open **Command Prompt** / **Terminal** inside the FlaskFormApp folder.
- Run:
 - `python app.py`
- You will see something like:
 - * Running on `http://127.0.0.1:5000`
- Open a browser and type the URL: <http://127.0.0.1:5000>
- To stop the Flask server, go to the terminal where it is running and press **Ctrl + C**.

DATE: 26-11-25

EXERCISE 17: Flask Web App with Registration & Login Using AWS RDS (MySQL)
(Full Deployment on AWS EC2 + AWS RDS)

To design and deploy a Python Flask web application on an AWS EC2 instance that allows:

1. User Registration – store Name and Password in AWS RDS MySQL database.
2. User Login – validate user credentials from the RDS database.
3. Welcome Page – display a message on successful login.

This exercise demonstrates a complete cloud-based web application workflow using Flask + MySQL (RDS).

System Architecture

Flow: Browser → Flask App (on EC2) → MySQL Database (RDS)

- Frontend: HTML templates (registration form, login form, welcome page)
- Backend: Python Flask application running on EC2
- Database: AWS RDS – MySQL engine
- Hosting: Linux EC2 instance
- Web Server: Flask's built-in development server (Apache/NGINX not used in this exercise)

Folder Structure (on EC2 instance)

Create the following structure inside the EC2 instance:

```
FlaskLoginApp/  
├── app.py  
├── db_config.py  
└── templates/  
    ├── register.html  
    ├── login.html  
    ├── success.html  
    ├── error.html  
    └── welcome.html
```

From EC2 Create Database and Table in RDS MySQL

Where:

The SQL commands are executed from the EC2 instance, by connecting to the RDS MySQL database using the MySQL client.

When:

After the RDS MySQL instance status becomes “Available” and after allowing EC2 security group access in RDS inbound rules.

Create:

- A database: flaskdb
- A table: users

SQL: Create Database and Table

```
CREATE DATABASE flaskdb;
```

```
USE flaskdb;
```

```
CREATE TABLE users (  
    id INT AUTO_INCREMENT PRIMARY KEY,  
    name VARCHAR(50) NOT NULL,  
    password VARCHAR(50) NOT NULL  
);
```

- id – unique identifier for each user (auto-increment primary key).
- name – username (cannot be NULL).
- password – password (cannot be NULL).

Note: For learning purposes, we store plain text. In real systems, passwords must be **hashed**.

Flask Application Code (Backend)

6.1 db_config.py

This file stores the **RDS MySQL connection details**.

db_config.py

```
db_config = {  
    'host': 'your-rds-endpoint.amazonaws.com',  
    # e.g., flaskdb-instance.xxxxxx.ap-south-1.rds.amazonaws.com  
    'user': 'admin',           # master username created in RDS  
    'password': 'YourPasswordHere', # master password created in RDS  
    'database': 'flaskdb'      # database name created in MySQL  
}
```

Replace host, user, and password with the actual values from the RDS instance.

6.2 app.py

This is the **main Flask application**.

app.py


```
from flask import Flask, request, render_template
import mysql.connector
from db_config import db_config

app = Flask(__name__)

def get_connection():
    return mysql.connector.connect(**db_config)

@app.route('/')
def home():
    return render_template('register.html')

@app.route('/register', methods=['POST'])
def register():
    name = request.form['uname']
    password = request.form['pwd']
    conn = get_connection()
    cur = conn.cursor()
    sql = "INSERT INTO users (name, password) VALUES (%s, %s)"
    cur.execute(sql, (name, password))
    conn.commit()

    cur.close()
    conn.close()

    return render_template('success.html')

@app.route('/login')
def login():
    return render_template('login.html')

@app.route('/validate', methods=['POST'])
def validate():
    name = request.form['uname']
    password = request.form['pwd']

    conn = get_connection()
    cur = conn.cursor()

    sql = "SELECT * FROM users WHERE name = %s AND password = %s"
    cur.execute(sql, (name, password))
    user = cur.fetchone()
    cur.close()
    conn.close()
```

```
if user:
    return render_template('welcome.html', username=name)
else:
    return render_template('error.html')

if __name__ == "__main__":
    app.run(host='0.0.0.0', port=5000, debug=True)
```

7. HTML Templates (Frontend)

All HTML files go inside the templates/ folder.

7.1 templates/register.html

```
<!DOCTYPE html>
<html>
<head>
    <title>User Registration</title>
</head>
<body>
    <h2>User Registration</h2>
    <form action="/register" method="post">
        <label>Name:</label>
        <input type="text" name="uname" required><br><br>
        <label>Password:</label>
        <input type="password" name="pwd" required><br><br>
        <button type="submit">Register</button>
    </form>
    <br>
    <a href="/login">Already registered? Click here to Login</a>
</body>
</html>
```

7.2 templates/login.html

```
<!DOCTYPE html>
<html>
<head>
    <title>User Login</title>
</head>
<body>
    <h2>User Login</h2>
```

```
<form action="/validate" method="post">
  <label>Name:</label>
  <input type="text" name="uname" required><br><br>
  <label>Password:</label>
  <input type="password" name="pwd" required><br><br>
  <button type="submit">Login</button>
</form>
<br>
<a href="/">New user? Click here to Register</a>
</body>
</html>
```

7.3 templates/welcome.html

```
<!DOCTYPE html>
<html>
<head>
  <title>Welcome</title>
</head>
<body>
  <h2>Welcome, {{ username }}!</h2>
  <p>You have successfully logged in.</p>
</body>
</html>
```

7.4 templates/success.html

```
<!DOCTYPE html>
<html>
<head>
  <title>Registration Successful</title>
</head>
<body>
  <h3>Registration Successful!</h3>
  <a href="/login">Click here to Login</a>
</body>
</html>
```

7.5 templates/error.html

```
<!DOCTYPE html>
<html>
<head>
  <title>Login Error</title>
```

```
</head>
<body>
  <h3>Invalid Username or Password</h3>
  <a href="/login">Try Again</a>
</body>
</html>
```

Deployment Steps –


PHASE 1 – Launch EC2 Linux Instance

- Open AWS Management Console → EC2 → Launch instance.
- Configuration:
 - Name: FlaskLoginServer
 - AMI: Amazon Linux 2023 (Free Tier eligible)
 - Instance type: t3.micro (Free Tier)
 - Key pair: create or select existing
 - Network: Default VPC and any public subnet
 - Auto-assign Public IP: Enable
- Security Group for EC2 – Inbound rules:

Type	Port	Source	Purpose
SSH	22	My IP	To connect via SSH/EC2 Connect
Custom TCP	5000	0.0.0.0/0	To test Flask app in browser

- Outbound: Allow all traffic.
- Click Launch instance.

The screenshot displays the AWS Management Console interface for an EC2 instance. The breadcrumb navigation at the top shows 'EC2 > Instances > i-04e767f357d6f6e36'. The instance name is 'Flask_App'. The summary is divided into several sections: 'Instance ID' (i-04e767f357d6f6e36), 'IPv6 address' (none), 'Hostname type' (IP name: ip-172-31-20-82.ap-northeast-1.compute.internal), 'Answer private resource DNS name' (IPv4 (A)), 'Auto-assigned IP address' (55.76.30.220 [Public IP]), 'IAM Role' (none), 'IMDSv2' (Required), 'Operator' (none), 'Public IPv4 address' (35.76.30.220), 'Instance state' (Running), 'Private IP DNS name (IPv4 only)' (ip-172-31-20-82.ap-northeast-1.compute.internal), 'Instance type' (t3.micro), 'VPC ID' (vpc-0d57af8226f8027a), 'Subnet ID' (subnet-07e5bc67b305511d7), 'Instance ARN' (arn:aws:ec2:ap-northeast-1:499497677555:instance/i-04e767f357d6f6e36), 'Private IPv4 addresses' (172.31.20.82), 'Public DNS' (ec2-35-76-30-220.ap-northeast-1.compute.amazonaws.com), 'Elastic IP addresses' (none), 'AWS Compute Optimizer finding' (Opt-in to AWS Compute Optimizer for recommendations), 'Auto Scaling Group name' (none), and 'Managed' (false). The bottom navigation bar includes tabs for 'Details', 'Status and alarms', 'Monitoring', 'Security', 'Networking', 'Storage', and 'Tags'. The footer shows 'CloudShell', 'Feedback', 'Console Mobile App', and copyright information for Amazon Web Services, Inc.




PHASE 2 – Connect to EC2 and Install Dependencies

- Go to EC2 → Instances → Select instance → Connect.
- Choose EC2 Instance Connect (browser-based SSH) → Connect.
Or open PowerShell on windows and connect using ssh command
- Update packages:
`sudo yum update -y`
- Check Python version (Amazon Linux 2023 has Python 3):
`python3 --version`
- Install pip (if not already installed):
`sudo dnf install python3-pip -y`
- Install Flask and MySQL Connector for Python:
`pip3 install flask`
`pip3 install mysql-connector-python`

PHASE 3 – Create Flask Project Structure on EC2

- Create project folder:
`mkdir FlaskLoginApp`
`cd FlaskLoginApp`
- Create templates folder:
`mkdir templates`
- Create the Python and HTML files using nano:
`nano db_config.py` → paste code for db_config.py
`nano app.py` → paste code for app.py
`nano templates/register.html` → paste register page
`nano templates/login.html` → paste login page
`nano templates/welcome.html` → paste welcome page

Save each file using: CTRL + O → Enter → CTRL + X



PHASE 5 – Verify RDS Security Group (Auto-configured)

- Go to AWS Console → RDS → Databases → flaskdb-instance.
- Open Connectivity & security section.
- Under VPC security groups, click the linked security group name.
- In Inbound rules, verify that the following rule exists:

Type	Port	Source
MySQL / Aurora	3306	EC2 Security Group

5. If the rule exists, EC2 to RDS connectivity is correctly configured.

PHASE 6 – Create Database and Table in RDS from EC2

- From the EC2 terminal, install MySQL client (MariaDB client):

```
sudo dnf install mariadb105 -y
```

 # if this fails, try: `sudo dnf install mariadb -y`
- Connect to the RDS MySQL instance from EC2:

```
mysql -h <RDS-endpoint> -u admin -p
```

 Example: `mysql -h flaskdb-instance.xxxxxx.ap-south-1.rds.amazonaws.com -u admin -p`
- Inside the MySQL prompt, run following commands:

```
CREATE DATABASE flaskdb;
USE flaskdb;

CREATE TABLE users (
  id INT AUTO_INCREMENT PRIMARY KEY,
  name VARCHAR(50) NOT NULL,
  password VARCHAR(50) NOT NULL
);
```
- Exit MySQL: `exit;`

PHASE 7 – Configure Flask to Use RDS and Run Application

- Open db_config.py on EC2 and update with **correct RDS details**:

```
nano db_config.py
Ensure:
db_config = {
    'host': 'your-rds-endpoint.amazonaws.com',
    'user': 'admin',
    'password': 'YourPasswordHere',
    'database': 'flaskdb'
}
```

Save and exit.


- Start the Flask app:
cd ~/FlaskLoginApp
python3 app.py
You should see: * Running on http://0.0.0.0:5000

PHASE 8 – Test the Application from Browser

- In your local system browser, open: http://<EC2-public-IP>:5000
Example: http://13.232.122.81:5000
You should see: Registration Page (default route /)
- After registration → success message with link to /login
- Login Page → if correct credentials → Welcome Page
Check that data is stored in flaskdb.users table in RDS. - SELECT * FROM users;

Expected Output

1. Registration Page
 - User enters Name and Password.
 - On submit, data is inserted into users table in RDS MySQL.




User Registration

Name:

Password:

[Already registered? Click here to Login](#)




Registration Successful!

[Click here to Login](#)

2. Login Page
 - User enters same Name and Password.
 - Flask executes SELECT query on users table to check match.
3. Welcome Page


- For valid credentials:

“Welcome, <username>! You have successfully logged in.”



Welcome, Karthik!

You have successfully logged in.



Invalid Username or Password

[Try Again](#)

- For invalid credentials:

“Invalid Username or Password” message with link to try again.

Result

A **fully working cloud-based registration and login system** was successfully deployed using:

- **Flask** as the backend web framework
- **HTML** templates for frontend pages
- **MySQL** database hosted on **AWS RDS**
- **AWS EC2** instance as the application server

This exercise demonstrates a **complete end-to-end mini project** on AWS using **Flask + MySQL (RDS)**.

DATE: 28-11-25

Exercise 18: Creating and operating a NoSQL Key-Value Database using Amazon DynamoDB

To create an Amazon DynamoDB table and perform CRUD operations (Create, Read, Update, Delete) using the AWS Management Console, and understand Partition Key, Sort Key, Query vs Scan, and table cleanup.

Phase 1: Create DynamoDB Table

Create a DynamoDB table to store Student Records.

Store items like: USN, Name, Semester, Course, Attendance, IA1Marks

Step 1: Open DynamoDB

- Login to AWS Console

- Search → type DynamoDB → open Amazon DynamoDB

Step 2: Create a Table


- Left menu → Tables
- Click Create table
- Fill details:

Table details

- Table name: MCA_StudentLabInternals
- Partition key (PK): USN (Type: String)
- Sort key (SK): CourseCode (Type: String)
(Enable sort key option by setting sort key)

Why this design?

- One student can have multiple courses → Sort key separates records per course.
- This creates a composite primary key (PK + SK).



Step 3: Table settings

- Under Table settings, choose:
Default settings (recommended for lab)
- Ensure Capacity mode is: On-demand (default already selected)

Step 4: Create

- Click Create table
- Wait until table status becomes Active

Phase 2: Insert Items (Create Data)

Step 5: Open Table and Add Items

- Click the table: MCA_StudentLabInternals
- Click: Explore table items
- Click Create item

Step 6: Add Item 1 (Student 1 - Course 1)

- You will see attributes:
 - USN (String)
 - CourseCode (String)

Enter:

- USN = 1MS24MCA001
- CourseCode = CCL301

Now add additional attributes (click **Add new attribute**):

- Name (String) = Arun
- Semester (Number) = 3
- Attendance (Number) = 86
- IA1Marks (Number) = 18

Click **Create item**

Create item

You can add, remove, or edit the attributes of an item. You can nest attributes inside other attributes up to 32 levels deep. [Learn more](#)

Attributes

Attribute name	Value	Type	Remove
USN - Primary key	1MS24MCA001	String	
CourseCode - Sort key	CCL301	String	
Attendance	90	Number	Remove
Name	Arun	String	Remove
Semester	3	Number	Remove
IA1Marks	30	Number	Remove

[Add new attribute](#) [Cancel](#) [Create item](#)

Repeat Create item with:

- USN = 1MS24MCA001
- CourseCode = DBS301
- Name = Arun
- Semester = 3
- Attendance = 88
- IA1Marks = 20

Step 8: Add Item 3 (Another student)

Create item:

- USN = 1MS24MCA002
- CourseCode = CCL301
- Name = Chitra
- Semester = 3
- Attendance = 74
- IA1Marks = 14

You should now see 3 items in the table list.

Phase 3: Read Data (Get / Query / Scan)

Step 9: Get a single item (exact PK + SK)

- In “Explore table items”, use search/filter:
- Use USN = 1MS24MC001 and CourseCode = CCL301
- Open the item → verify all attributes

Key concept: To uniquely identify an item in a PK+SK table, you must provide **both**.

The left screenshot shows the 'Scan or query items' interface for the 'StudentLabInternals' table. The 'Scan' tab is selected. The 'Select a table or index' dropdown is set to 'Table - StudentLabInternals'. The 'Select attribute projection' dropdown is set to 'All attributes'. The 'Filters - optional' section shows two filters: 'USN' with a value of '35mc046' and 'CourseCode' with a value of '24mcy03'. The 'Run' button is highlighted.

The right screenshot shows the results of the scan. The title is 'Table: StudentLabInternals - Items returned (2)'. The scan started on December 19, 2025, 09:47:19. The results are displayed in a table with columns: USN (String), CourseCode (String), Attendance, Name, and Semester. The table contains two items:

USN (String)	CourseCode (String)	Attendance	Name	Semester
35mc046	24mcy03	90	Arun	3
35mc046	25mcy02	90	Arun	3

Step 10: Query – fetch all courses for one student

- Click Run query
- Set Partition key condition: USN equals 1MS24MCA001
- Run

Expected output:

- Items for Arun in both CCL301 and DBS301

Key concept: Query works only with Partition Key (and optional Sort Key conditions). It is efficient.

Step 11: Scan (Not for large tables)

- Click Scan
- Run scan without filters

Expected output:

- All items (Arun + Chitra)

Key concept: Scan reads the entire table → slow/costly for big tables.

Phase 4: Update Data

Step 12: Update IA1 marks of Chitra for CCL301


- Open item where:
USN = 1MS24MCA002
CourseCode = CCL301
- Click Edit
- Change: IA1Marks from 14 to 16
- Click Save changes

Verify updated value appears.

Phase 5: Delete Data

Step 13: Delete one item (Arun's DBS301 record)

- Select item:
 - USN = 1MS24MCA001
 - CourseCode = DBS301
- Click **Delete**
- Confirm delete



Now total items should reduce by 1.

Phase 6: Cleanup (Must Do to avoid charges)

Step 15: Delete the Table

- DynamoDB → Tables
- Select MCA_StudentLabInternals
- Click Delete
- Type confirmation (if asked) and delete

Ensure table disappears from list.

DATE: 03-12-25

Exercise 19: ElastiCache (Redis) as an In-Memory Cache

To implement Amazon ElastiCache (Redis) as an in-memory caching service by deploying a Redis cluster and performing basic cache operations from an EC2 instance.

Application-Level Data Flow (Logical)

Step 1: User requests data


Step 2: Application checks ElastiCache (Redis)

Step 3: If data exists → return from cache (FAST)

Step 4: If data does not exist → fetch from RDS

Step 5: Store fetched data in ElastiCache

Step 6: Return data to user



This exercise provides **exposure** to an industry-used in-memory data store.

Redis improves performance by serving frequently accessed data from memory and is typically used **in front of databases** in real-world applications.

Feature	RDS	ElastiCache
Storage	Disk	RAM
Data lifetime	Permanent	Temporary
Access speed	Slower	Very fast
TTL support	No	Yes

Phase-1 → EC2 + valkey-cli ready

Phase-2 → Redis cache created

Phase-3 → Connect EC2 → Redis and run commands

Phase-1: Launch EC2 Client using Amazon Linux 2023 (for Valkey / Redis Client)

Purpose of Phase-1

To create an EC2 instance using **Amazon Linux 2023** that will act as a **client machine** to connect to Amazon ElastiCache (Redis OSS) using **valkey-cli**.

Step 1: Select AWS Region

- Choose ONE region and stick to it for the entire exercise
ap-south-1 (Mumbai) (*recommended*)
- Ensure ElastiCache and EC2 will be in the SAME region

Step 2: Launch EC2 Instance

1. Go to **AWS Console** → **EC2**
2. Click **Launch instance**

Step 3: Configure Instance Basics

- Name: RedisClient-AL2023
- AMI: Select Amazon Linux 2023 AMI
- Instance Type: Select t3.micro (Free Tier eligible)
- Key Pair: Select an existing key pair OR create a new one (RSA, .pem)
- Network & Security
 - Network
 - VPC: Default VPC
 - Subnet: No preference
 - Auto-assign public IP: Enabled
 - Security Group (VERY IMPORTANT)
Create a new security group:
 - Security Group Name: SG-RedisClient
 - Description: Security group for Redis client EC2

Inbound Rules

Type	Port	Source
SSH	22	My IP (or Anywhere for lab)

Do NOT add Redis (6379) here (The client does not listen on 6379)

Outbound rules: Allow all (default)

Step 5: Storage

- Keep default

Step 6: Launch Instance

- Click Launch instance
- Wait until Instance State = Running

Step 7: Connect to EC2

1. EC2 → Instances → select RedisClient-AL2023
2. Click Connect
3. Choose EC2 Instance Connect
4. Click Connect

You are now inside the EC2 terminal.

Step 8: Install Valkey Client

Run the following commands:

```
sudo dnf update -y
```

```
sudo dnf install -y valkey
```

Verify installation:

```
valkey-cli --version
```

Expected Output (example)

```
valkey-cli 8.x.x
```

This confirms the EC2 is ready as a Redis-compatible client.

```

ec2-user@ip-172-31-7-242: ~
$ sudo dnf install -y valkey
Waiting for process with pid 1641 to finish.
Amazon Linux 2023 Kernel Livepatch repository
Dependencies resolved.
Nothing to do.
Complete!
Last metadata expiration check: 0:00:02 ago on Wed Dec 24 06:31:09 2025.
Dependencies resolved.
=====
Package                Architecture      Version           Repository        Size
-----
Installing:
valkey                  x86_64            8.0.6-3.amzn2023.0.2  amazonlinux      1.6 M
=====
Transaction Summary
-----
Install 1 Package
Total download size: 1.6 M
Installed size: 5.8 M
Downloading Packages:
valkey-8.0.6-3.amzn2023.0.2.x86_64.rpm                                14 MB/s | 1.6 MB  00:00
-----
Total:
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
Preparing
Running scriptlet: valkey-8.0.6-3.amzn2023.0.2.x86_64                                1/1
Installing
Running scriptlet: valkey-8.0.6-3.amzn2023.0.2.x86_64                                1/1
Running scriptlet: valkey-8.0.6-3.amzn2023.0.2.x86_64                                1/1
Verifying
Running scriptlet: valkey-8.0.6-3.amzn2023.0.2.x86_64                                1/1
Installed:
valkey-8.0.6-3.amzn2023.0.2.x86_64
Complete!
ec2-user@ip-172-31-7-242: ~
$ valkey --version
-bash: valkey: command not found
ec2-user@ip-172-31-7-242: ~
$ valkey-cli --version
valkey-cli 8.0.6
ec2-user@ip-172-31-7-242: ~
$

```

Note: why install valkey and not redis??

Amazon Linux 2023 does not provide redis-cli directly.

AWS now provides **Valkey**, which is fully Redis-protocol compatible.

Hence, we use **valkey-cli** to connect to ElastiCache Redis.

Phase-2: Create Amazon ElastiCache (Redis OSS) Cluster

Purpose of Phase-2

To create an Amazon ElastiCache Redis OSS cluster that will act as an in-memory cache, accessible from the EC2 client created in Phase-1.

Step 1: Confirm AWS Region

- Ensure you are in the **same region** used in Phase-1
EC2 and ElastiCache must be in the SAME region and VPC.

Step 2: Open ElastiCache Console

1. AWS Console → **Services**
2. Select ElastiCache
3. Click Create cache

Step 3: Select Cache Engine

- **Engine: Redis OSS**

Step 4: Choose Deployment Settings

- Deployment option: Node-based cluster
- Creation method: Easy create

Easy create is used to avoid advanced production settings.

Step 5: Select Configuration

- **Configuration: Demo**

This automatically selects:

- A small, low-cost node (e.g., cache.t4g.micro)
- Suitable for labs and practice

Step 6: Provide Cluster Information

- Cache name: lab-redis
- Description: Optional (lab-redis)

Step 7: Configure Network and Subnet Group

Network

- Network type: IPv4

Subnet Group

- Select: Create a new subnet group
 - Subnet group name: redis-subnet-group
 - VPC: Default VPC
 - Subnets: Leave AWS auto-selected subnets unchanged

No manual subnet selection required.

Step 8: Configure Security

Security Group

- Create or select a security group:
 - Name: SG-RedisCache

Inbound Rule for Redis

Add **one inbound rule** to SG-RedisCache:

Type	Port	Source
Custom TCP	6379	SG-RedisClient

This allows only the EC2 client to access Redis.


Step 9: Authentication

- Authentication: Disabled

Step 10: Create Cache

1. Review all settings
2. Click Create
3. Wait until Status = Available

This may take a few minutes.



Step 11: Note the Redis Endpoint

1. Click on the cache name lab-redis
2. Copy the **Primary Endpoint**
 - Example: lab-redis.xxxxxx.cache.amazonaws.com

This endpoint will be used in **Phase-3** to connect from EC2.

Note:

We have created an in-memory Redis cache using Amazon ElastiCache. It runs inside the AWS VPC and does not have a public IP. Only our EC2 client is allowed to access it using port 6379.

Phase-3: Connect EC2 to ElastiCache Redis and Execute Cache Commands

Purpose of Phase-3

To connect the EC2 client (Amazon Linux 2023) to the ElastiCache Redis OSS cluster using valkey-cli and demonstrate basic in-memory cache operations.

Pre-checks - Before connecting, confirm:

- EC2 and ElastiCache are in the same AWS region
- EC2 security group = SG-RedisClient
- Redis security group = SG-RedisCache
- Redis security group allows port 6379 from SG-RedisClient
- Redis Status = Available
- You have copied the Primary Endpoint

Step 1: Connect to EC2

1. AWS Console → **EC2**
2. Select instance RedisClient-AL2023
3. Click **Connect**
4. Choose **EC2 Instance Connect**
5. Click **Connect**

You are now inside the EC2 terminal.

Step 2: Verify Valkey Client

Run: valkey-cli --version

Expected Output (example)

valkey-cli 8.x.x

This confirms the EC2 is ready to act as a Redis-compatible client.

Step 3: Connect to ElastiCache Redis

Use the configuration Endpoint from Phase-2.

valkey-cli -h <CONFIGURATION_ENDPOINT> -p 6379 --tls -c

Example: valkey-cli -h lab-redis.xxxxxx.cache.amazonaws.com -p 6379

Expected Result

You should see a prompt like:

lab-redis.xxxxxx.cache.amazonaws.com:6379>

This means the connection is successful.

Step 4: Execute Redis / Valkey Commands

These commands simulate what an application does internally.

Test Connectivity

PING

Expected Output

PONG

Store and Retrieve Data (SET / GET)

SET course "Cloud Computing"

GET course

Expected Output

"Cloud Computing"

Demonstrates **key-value storage in memory**.

Counter Example (INCR)

INCR visits

INCR visits

GET visits

Expected Output

"2"

NOTE: Common real-world use - page views, hit counters.

4.4 Set Data with Expiry (TTL)

SET notice "Results Published" EX 60

TTL notice

GET notice

Expected Output

- TTL shows a value ≤ 60
- GET returns:

"Results Published"

Shows **temporary cache data**.

4.5 Verify Automatic Expiry

Wait ~60 seconds, then run:

GET notice

Expected Output

(nil)

Confirms data is **automatically removed from memory**.

4.6 Delete Data Manually


DEL course

GET course

Expected Output: (nil)

Step 5: Exit Redis Client

EXIT



```
ec2-user@ip-172-31-7-242:~$ redis-cli
clustercfg.lab-redis.bqxcqw.ap1.cache.amazonaws.com:6379> ping
PONG
clustercfg.lab-redis.bqxcqw.ap1.cache.amazonaws.com:6379> set course "cloud computing"
OK
clustercfg.lab-redis.bqxcqw.ap1.cache.amazonaws.com:6379> get course
"cloud computing"
clustercfg.lab-redis.bqxcqw.ap1.cache.amazonaws.com:6379> INCR visits
(integer) 1
clustercfg.lab-redis.bqxcqw.ap1.cache.amazonaws.com:6379> INCR visits
(integer) 2
clustercfg.lab-redis.bqxcqw.ap1.cache.amazonaws.com:6379> get visits
"2"
clustercfg.lab-redis.bqxcqw.ap1.cache.amazonaws.com:6379> SET notice "Results Published" EX 60
OK
clustercfg.lab-redis.bqxcqw.ap1.cache.amazonaws.com:6379> TTL notice
(integer) 55
clustercfg.lab-redis.bqxcqw.ap1.cache.amazonaws.com:6379> get notice
"Results Published"
clustercfg.lab-redis.bqxcqw.ap1.cache.amazonaws.com:6379> get notice
(nil)
clustercfg.lab-redis.bqxcqw.ap1.cache.amazonaws.com:6379> DEL course
(integer) 1
clustercfg.lab-redis.bqxcqw.ap1.cache.amazonaws.com:6379> GET course
(nil)
clustercfg.lab-redis.bqxcqw.ap1.cache.amazonaws.com:6379> EXIT
[ec2-user@ip-172-31-7-242 ~]$
```

Note:

“The application first checks Redis.

If data is present, it is returned immediately from memory.

If not present, the application fetches data from the database and stores it in Redis with a TTL.

Redis automatically removes the data after expiry.”

Phase-3 Outcome

- EC2 successfully connected to ElastiCache Redis
- In-memory key–value operations verified
- Temporary storage and TTL behavior observed
- Redis used as a **cache**, not as a primary database

Common Errors & Quick Fix

Problem	Reason	Fix
Connection timeout	Wrong SG or region	Check SG-RedisCache inbound rule
Command hangs	Redis not Available	Wait & retry
(nil) output	Key expired	Expected behavior

Clean-Up (Mandatory)

Step 1: Delete Redis Cluster

- ElastiCache → Select lab-redis → Delete
- Disable snapshots

Step 2: Terminate EC2 Instance

- EC2 → Instances → Terminate RedisClient-EC2

Step 3: Optional

- Delete unused security groups

Redis Commands Explanation

- **SET** course "Cloud Computing"
- **GET** course
- **EXPIRE** course 30
- **TTL** course
- The above commands simulate application caching behavior.
The SET command represents storing frequently accessed data in cache.
The EXPIRE command shows that cached data is temporary and stored in memory.
After expiry, the application would fetch fresh data from the database again.

DATE: 05-12-25

Exercise–20: Deploy a simple Flask application on AWS Elastic Beanstalk and verify it using the public URL.

Phase A — Create the Flask App

Step A1: On your system, create a folder

Create a folder: eb-flask-lab

Step A2: Create application.py

Create a file named application.py:

```
from flask import Flask

app = Flask(__name__)

@app.route("/")
def home():
    return "Hello from Flask on Elastic Beanstalk!"

@app.route("/health")
def health():
    return "OK"

if __name__ == "__main__":
    app.run()
```

Step A3: Create requirements.txt

Create a file named **requirements.txt** and add:

```
Flask==3.0.3
gunicorn==22.0.0
```

Step A4: Create Procfile (MOST IMPORTANT)

Create a file named **exactly**:

Procfile (no extension)

Contents:

```
web: gunicorn application:app
```

[Windows Notepad method:

File name: Procfile, Save as type: **All Files**, Encoding: UTF-8]

Step A5: Verify folder contents

Your folder must contain exactly these 3 files:

eb-flask-lab

- |— application.py
- |— requirements.txt
- |— Procfile (Type should show as “File”, NOT Procfile.txt)

Phase B — Create the ZIP correctly (root-level ZIP)

Step B1: Select the 3 files

Select:

- application.py
- requirements.txt
- Procfile

Step B2: Create ZIP

Right click → Compress to → ZIP file

Rename the ZIP as:

eb-flask-lab.zip

Step B3: Confirm ZIP contents (one-time check)

Open the ZIP and confirm it shows (at top level):

application.py

requirements.txt

Procfile

If you see a folder inside the ZIP, that is wrong.

Phase C — Deploy on Elastic Beanstalk (AWS Console)

Step C1: Open Elastic Beanstalk

AWS Console → Search Elastic Beanstalk → Open

Ensure region is Asia Pacific (Mumbai) (ap-south-1)

Step C2: Create Environment

Click Create environment

Environment tier - Select: Web server environment

Application information

- Application name: EB-Flask-Lab (or any name)

Environment information

- Environment name: EB-Flask-Lab-env (or any name)
- Domain: leave blank (auto)

Platform

- Platform: Python
- Platform branch: latest Python on Amazon Linux
- Platform version: recommended

Step C3: Upload your code (IMPORTANT SETTINGS)

In Application code:

1. Select: Upload your code
2. Version label: type v1
3. Source code origin: Local file
(Do NOT choose “Public S3 URL”)
4. Click Choose file and select:
5. eb-flask-lab.zip

Presets

Select: Single instance (free tier eligible)

Click Next

Phase D — Configure Service Access (IAM Roles)

On Configure service access page:

Step D1: Service role

If dropdown shows “No options”:

- Click Create role (opens IAM)
- Use case: Elastic Beanstalk
- Keep default permissions (AWS auto selects)
- Create role name:
- aws-elasticbeanstalk-service-role

Return to EB tab → click Refresh → select this role.

Step D2: EC2 instance profile

If dropdown shows “No options”:

- Click Create role
- Use case: EC2
- Attach policy: AWSElasticBeanstalkWebTier
- Role name: aws-elasticbeanstalk-ec2-role


Return to EB tab → refresh → select this role.

Step D3: EC2 Key pair

Leave blank (optional)

Click Next

Leave remaining pages as default → Create environment



Part E — Wait for Deployment

Step E1: Monitor

Open Events tab.

Final success condition:

- Green message: Environment successfully launched
- Health: OK
- Domain URL appears

Phase F — Test the Application

Step F1: Open the domain URL

Click the Domain link shown in EB.


Expected output: Hello from Flask on Elastic Beanstalk!

Step F2: Test /health

In the browser address bar, append /health

Example: `http://<your-domain>.elasticbeanstalk.com/health`

Expected output: OK



Hello from Flask on Elastic Beanstalk!

Phase G — Cleanup (must do after lab)

Step G1: Terminate Environment

Elastic Beanstalk → Environment

Actions → Terminate environment

(Optional) After termination: Delete application if required.

Common Mistakes (quick checklist)

Profile must be Profile (no .txt)

ZIP must contain files at root level (not folder)

Upload must be **Local file** (not Public S3 URL)

IAM roles must be created/selected if “No options”

DATE: 10-12-25

Exercise-21: Deploy a Flask application integrated with DynamoDB on AWS Elastic Beanstalk and verify the CRUD operations using CURL.

A) LOCALLY — Prepare Folder + Code (Before AWS)

1) Create folder

Create a folder named: flask-ddb-lab

2) Inside it, create 4 files

Your folder must look like:

```
flask-ddb-lab/  
  application.py  
  requirements.txt  
  Procfile  
  runtime.txt
```

3) Paste code into application.py

```
from flask import Flask, request, jsonify  
import boto3  
import os  
  
app = Flask(__name__)  
  
REGION = os.getenv("AWS_REGION", "ap-south-1")  
TABLE_NAME = os.getenv("TABLE_NAME", "student")  
  
dynamodb = boto3.resource("dynamodb", region_name=REGION)  
table = dynamodb.Table(TABLE_NAME)  
  
@app.route("/")  
def home():  
    return "Flask + DynamoDB on Elastic Beanstalk is working!"  
  
@app.route("/health")  
def health():  
    return jsonify(status="ok")  
  
# CREATE  
@app.route("/student", methods=["POST"])  
def create_student():
```

```

data = request.get_json()
table.put_item(Item={
    "StudentID": data["StudentID"],
    "Name": data["Name"],
    "Dept": data.get("Dept", "")
})
return jsonify(message="Student created"), 201

# READ
@app.route("/student/<student_id>", methods=["GET"])
def get_student(student_id):
    resp = table.get_item(Key={"StudentID": student_id})
    item = resp.get("Item")
    if not item:
        return jsonify(error="Student not found"), 404
    return jsonify(item)

# UPDATE
@app.route("/student/<student_id>", methods=["PUT"])
def update_student(student_id):
    data = request.get_json()
    resp = table.update_item(
        Key={"StudentID": student_id},
        UpdateExpression="SET #n = :n, Dept = :d",
        ExpressionAttributeNames={"#n": "Name"},
        ExpressionAttributeValues={" :n": data["Name"], ":d": data.get("Dept", "")},
        ReturnValues="UPDATED_NEW"
    )
    return jsonify(message="Student updated", updated=resp["Attributes"])

# DELETE
@app.route("/student/<student_id>", methods=["DELETE"])
def delete_student(student_id):
    table.delete_item(Key={"StudentID": student_id})
    return jsonify(message="Student deleted")

```

4) requirements.txt

```

Flask==3.0.0
boto3==1.34.0
gunicorn==21.2.0

```

5) Procfile

web: gunicorn application:app

6) runtime.txt

python-3.11

7) Zip it

Select these 4 files **together** → right click → **Send to** → **Compressed (zipped) folder**.

Rename the zip: flask-ddb-lab.zip

When you open the zip, you must directly see:

- application.py
- requirements.txt
- Procfile
- runtime.txt

(No extra folder inside.)

B) AWS — DynamoDB Table

Create table

AWS Console → **DynamoDB** → Tables → **Create table**

- Table name: student
 - Partition key: StudentID (String)
- Create → wait till **ACTIVE**

C) AWS — IAM Role (for EB EC2 to access DynamoDB)

Create role

AWS Console → **IAM** → Roles → **Create role**

- Trusted entity: **AWS service**
 - Use case: **EC2**
 - Attach policy: **AmazonDynamoDBFullAccess**
 - Role name: EB-EC2-DynamoDB-Role
- Create role

[once role is created, for next time it will be automatically selected. No need to create again]

D) AWS — Elastic Beanstalk Environment (Sample app first)

Create EB app/environment

AWS Console → **Elastic Beanstalk** → **Create application**


- Environment tier: **Web server environment**
- Application name: flask-ddb-lab
- Platform: **Python**
- Application code: **Sample application**
- Presets: **Single instance (free tier eligible)**

Configure service access


On **Configure service access** page:

- Service role: leave default
- **EC2 instance profile: select EB-EC2-DynamoDB-Role**
Next → Next → Review → **Create environment**

Wait until you can open the EB domain and see sample page.



Click on Domain: [Flask-ddb-lab-env.eba-jtqiug36.ap-south-1.elasticbeanstalk.com](https://flask-ddb-lab-env.eba-jtqiug36.ap-south-1.elasticbeanstalk.com)



E) AWS — Deploy Your ZIP

Upload and deploy


EB Environment → **Upload and deploy**

- Upload: flask-ddb-lab.zip
- Version label: v1

Click **Deploy**

After deploy, opening the EB domain should show:

“Flask + DynamoDB on Elastic Beanstalk is working!”



Flask + DynamoDB on Elastic Beanstalk is working!

F) Add Environment Variables in Elastic Beanstalk (MANDATORY)

Step 1

Go to: AWS Console → Elastic Beanstalk → Environments → flask-ddb-lab-env

Step 2

Click: Configuration

Step 3

Under **Software**, click: Edit

Step 4

Scroll to **Environment properties**

Add these two entries:

Name	Value
TABLE_NAME	student
AWS_REGION	ap-south-1

Step 5

Click: Apply

Step 6

Wait for **Environment update** to complete
(Status may show **Updating** / **Pending** — wait until it finishes)

G) Verify Application After Env Update

Open the **EB Environment URL** in browser:

`http://<your-eb-domain>/`

Expected output: Flask + DynamoDB on Elastic Beanstalk is working!

H) Test Health Endpoint (GET)


Step 1

Open **AWS CloudShell** (or Command Prompt / PowerShell)

Step 2

Run: `curl http://<your-eb-domain>/health`

Expected output: `{"status":"ok"}`

A screenshot of the AWS CloudShell interface. The title bar shows 'CloudShell' with a close button. Below the title bar, there's a tab labeled 'ap-northeast-1' with a plus sign to its right. The terminal window shows a command prompt '~ \$' followed by the command 'curl http://Flask-ddb-lab-env.eba-z3ftuabu.ap-northeast-1.elasticbeanstalk.com/health'. The output of the command is displayed on the next line: '{"status":"ok"}'. The prompt '~ \$' is shown again on the following line, indicating the command has completed.

```
CloudShell
ap-northeast-1 | +
~ $ curl http://Flask-ddb-lab-env.eba-z3ftuabu.ap-northeast-1.elasticbeanstalk.com/health
{"status":"ok"}
~ $
```

AWS CloudShell - Steps

1. In AWS Console (top right)
2. Search **CloudShell** in the search bar
3. Wait until terminal opens (you'll see \$ prompt)
4. Paste the following command (replace domain):

I) CREATE Item in DynamoDB (POST)

Step 1

Run:

```
curl -X POST http:// Flask-ddb-lab-env-1.eba-k9hmvizi.ap-south-1.elasticbeanstalk.com /Student \
-H "Content-Type: application/json" \
-d '{"StudentID":"101","Name":"Anita","Dept":"MCA"}'
```

Expected output:

```
{"message":"Student created"}
```

J) READ Item from DynamoDB (GET)

Run:

```
curl http://<your-eb-domain>/student/101
```

Expected output (JSON with student data).

K) UPDATE Item (PUT)

Run:

```
curl -X PUT http://<your-eb-domain>/student/101 \
-H "Content-Type: application/json" \
-d '{"Name":"Anita S","Dept":"MCA-III"}'
```

Expected output:

```
{"message":"Student updated"}
```


L) DELETE Item (DELETE)

Run:

```
curl -X DELETE http://<your-eb-domain>/student/101
```

Expected output:

```
{"message":"Student deleted"}
```



The screenshot shows a CloudShell terminal window with the following commands and output:

```
~$ clear
~$ curl -X POST http://Flask-ddb-lab-env-e3ftuabu.ap-northeast-1.elasticbeanstalk.com/student -H "Content-Type: application/json" -d '{"StudentID":"101","Name":"Anita","Dept":"MCA"}'
{"message":"Student created"}
~$ curl http://Flask-ddb-lab-env-e3ftuabu.ap-northeast-1.elasticbeanstalk.com/student/101
{"Dept":"MCA","Name":"Anita","StudentID":"101"}
~$ curl -X PUT http://Flask-ddb-lab-env-e3ftuabu.ap-northeast-1.elasticbeanstalk.com/student/101 -H "Content-Type: application/json" \
-d '{"Name":"Anita S","Dept":"MCA-III"}'
{"message":"Student updated","updated":{"Dept":"MCA-III","Name":"Anita S"}}
~$ curl -X DELETE http://Flask-ddb-lab-env-e3ftuabu.ap-northeast-1.elasticbeanstalk.com/student/101
{"message":"Student deleted"}
```

M) Verify in DynamoDB Console

Go to:

AWS Console → DynamoDB → Tables → Students → Explore table items

Confirm record creation / update / deletion.

DATE: 12-12-25

Exercise 22: CloudFormation – Launch EC2 (Amazon Linux 2023) + Install Apache using UserData.

Create an EC2 instance using **AWS CloudFormation (YAML template)** and automatically install **Apache web server (httpd)** using **UserData**, then verify the website from a browser.

Part A — Create Key Pair (One-time setup)

Step A1: Open EC2 Key Pairs

1. AWS Console → Search **EC2**
2. Left menu → **Key Pairs** (under “Network & Security”)
3. Click **Create key pair**

Step A2: Create key pair

- Name: pemkeypair (any name is fine)
- Key pair type: **RSA**
- Private key file format:
 - **.pem** (recommended)

Click **Create key pair**

A file will download like: pemkeypair.pem

Note: CloudFormation uses **key pair NAME** (pemkeypair), not the file name.

Part B — Create CloudFormation Template File (Amazon Linux 2023 + Apache)

Step B1: Create YAML file on your computer

1. Open **Notepad**
2. Paste the full YAML template given below
3. Save as: **ec2-apache-al2023.yaml**
 - Save type: **All files**
 - Encoding: UTF-8 (if asked)

Full CloudFormation Template (Amazon Linux 2023)

Important:

- Do **not** add .pem anywhere.

- You will select Key Pair from dropdown during stack creation.

AWS::Template::FormatVersion: "2010-09-09"

Description: Launch EC2 (Amazon Linux 2023) and install Apache

Parameters:

KeyName:

Type: AWS::EC2::KeyPair::KeyName

Description: Select an existing EC2 Key Pair

AmazonLinux2023AMI:

Type: AWS::SSM::Parameter::Value<AWS::EC2::Image::Id>

Default: /aws/service/ami-amazon-linux-latest/al2023-ami-kernel-default-x86_64

Resources:

WebServerSG:

Type: AWS::EC2::SecurityGroup

Properties:

GroupDescription: Allow SSH and HTTP

SecurityGroupIngress:

- IpProtocol: tcp

FromPort: 22

ToPort: 22

CidrIp: 0.0.0.0/0

- IpProtocol: tcp

FromPort: 80

ToPort: 80

CidrIp: 0.0.0.0/0

WebServerInstance:

Type: AWS::EC2::Instance

Properties:

InstanceType: t3.micro

KeyName: !Ref KeyName

ImageId: !Ref AmazonLinux2023AMI

SecurityGroupIds:

- !Ref WebServerSG

UserData:

Fn::Base64: |

#!/bin/bash

dnf update -y

dnf install -y httpd

systemctl enable httpd

systemctl start httpd

echo "<h1>Apache Installed via CloudFormation (Amazon Linux 2023)</h1>" >

/var/www/html/index.html

Outputs:

WebsiteURL:

Description: Apache Website URL

Value: !Sub "http://\${WebServerInstance.PublicDnsName}"

Part C — Create CloudFormation Stack (Console Steps)

Step C1: Open CloudFormation

1. AWS Console → Search **CloudFormation**
2. Click **Stacks**
3. Click **Create stack** → **With new resources (standard)**

Step C2: Prepare template (select correct options)

1. Under **Prepare template**:
Select **Choose an existing template**
2. Under **Template source**:
Select **Upload a template file**
3. Click **Choose file** → select ec2-apache-al2023.yaml
4. Click **Next**

Part D — Specify Stack Details

Step D1: Stack name

- Stack name: EC2-Apache-AL2023

Click **Next**

Step D2: Parameters

- Under **KeyName**, select your key pair name from dropdown
Example: pemkeypair

Click **Next**

Part E — Configure Stack Options (keep default)

1. Leave everything as default
2. Click **Next**

Part F — Review and Create


1. Scroll down
2. Click **Create stack**

Part G — Monitor Stack Creation

1. Wait for Stack status to become:
CREATE_COMPLETE
2. Open the stack → click **Outputs** tab
3. Copy **WebsiteURL**
4. Paste URL in browser

Expected Output in browser:

Apache Installed via CloudFormation UserData (Amazon Linux 2023)!



Operation ID	Timestamp	Logical ID	Status	Detailed status	Status reason
6f121ef2-7094-4632-9231-44a5a6b0b5a9	-	-	-	-	-
6f121ef2-7094-4632-9231-44a5a6b0b5a9	2026-01-02 09:50:30 UTC+0530	EC2-Apache-AL2023	CREATE_COMPLETE	-	-
6f121ef2-7094-4632-9231-44a5a6b0b5a9	2026-01-02 09:50:28 UTC+0530	WebServerInstance	CREATE_COMPLETE	-	-
6f121ef2-7094-4632-9231-44a5a6b0b5a9	2026-01-02 09:50:16 UTC+0530	WebServerInstance	CREATE_IN_PROGRESS	-	Resource creation Initiated
6f121ef2-7094-4632-9231-44a5a6b0b5a9	2026-01-02 09:50:14 UTC+0530	WebServerInstance	CREATE_IN_PROGRESS	-	-
6f121ef2-7094-4632-9231-44a5a6b0b5a9	2026-01-02 09:50:14 UTC+0530	WebServerSG	CREATE_COMPLETE	-	-

Part H — Verify in EC2 (Optional)

1. Go to **EC2** → **Instances**
2. Instance should be running
3. Security group should allow:
 - SSH 22
 - HTTP 80

Troubleshooting (Common Errors)

1) Website not opening

Check:

- EC2 instance status checks are **2/2 passed**
- Security group has **port 80 open**
- Wait 2–3 minutes (UserData takes time)

2) Stack went to ROLLBACK

Go to stack → **Events** tab → check the failure reason

Most common reason:

- Wrong Key Pair selected / key pair does not exist

The screenshot displays the AWS CloudFormation console. On the left, the 'Stacks' list shows 'EC2-Apache-AL2023' with a status of 'CREATE_COMPLETE'. The main panel shows the 'Outputs' tab for this stack. The output table has the following data:

Key	Value	Description	Export name
WebsiteURL	http://ec2-18-176-178-214.ap-northeast-1.compute.amazonaws.com	Apache Website URL	-



Apache Installed via CloudFormation (Amazon Linux 2023)

Part I — Clean-up

To avoid charges:

1. CloudFormation → Stacks
2. Select EC2-Apache-AL2023
3. Click **Delete**
4. Confirm

This deletes EC2 + Security Group automatically.

DATE: 17-12-25

Exercise–23: EC2 + S3 (Static Content Pulled from S3 using CloudFormation)

Create an **S3 bucket** using CloudFormation.

Upload a static HTML file (index.html) to S3.

Launch **EC2 (Amazon Linux 2023) + Apache** using CloudFormation.

EC2 will **pull index.html from S3** and host it via Apache.

Pre-requisites

- Region: **Mumbai (ap-south-1)**
- A Key Pair exists (example: pemkeypair)

PART 0 — One-time: Create Key Pair

EC2 → Key Pairs → Create key pair

- Name: pemkeypair
- Format: .pem
Download it and keep safe.

PART 1 — Stack-1: Create S3 Bucket (CloudFormation)

Step 1.1: Create S3 template file

Create a file: **stack1-s3-bucket.yaml** and paste:

AWSTemplateFormatVersion: "2010-09-09"

Description: Create an S3 bucket to store website content (index.html)

Resources:

WebsiteBucket:

Type: AWS::S3::Bucket

Outputs:

BucketName:

Description: S3 Bucket Name (use this to upload index.html)

Value: !Ref WebsiteBucket


Step 1.2: Create stack

CloudFormation → Stacks → Create stack → With new resources

- Choose an existing template
- Upload a template file → stack1-s3-bucket.yaml
- Stack name: S3-Website-Bucket
- Create stack

Step 1.3: Copy bucket name

Open stack → **Outputs** → copy **BucketName**



PART 2 — Upload Static Content to S3 (Manual)

Step 2.1: Create index.html on your PC

Create a file named **index.html** with this content:

```
<!DOCTYPE html>
```

```
<html>
```

```
<head>
```

```
<title>EC2 + S3 Demo</title>
```

</head>

<body>

<h1>Hello! This page was pulled from S3 to EC2 automatically.</h1>

<p>Deployed using CloudFormation + UserData</p>

</body>


</html>

Step 2.2: Upload to S3 bucket

S3 → Buckets → open your bucket (from Output) → Upload

- Upload **index.html**
- Keep it at **root** (no folder)
- Upload

Now S3 has: s3://<your-bucket-name>/index.html



PART 3 — Stack-2: Launch EC2 + Apache + Pull from S3

This stack will:

- Create IAM Role (permission to read the bucket)
- Create Security Group (22, 80)
- Launch EC2 (Amazon Linux 2023)
- Install Apache
- Copy index.html from S3 to /var/www/html/index.html

Step 3.1: Create EC2 template file

Create a file: **stack2-ec2-pull-from-s3.yaml** and paste:

AWSTemplateFormatVersion: "2010-09-09"

Description: Launch EC2 (Amazon Linux 2023), install Apache, and pull index.html from S3

Parameters:

KeyName:

Type: AWS::EC2::KeyPair::KeyName

Description: Select an existing EC2 Key Pair to enable SSH access

BucketName:

Type: String

Description: Enter the S3 bucket name created in Stack-1 (Output)

SubnetId:

Type: AWS::EC2::Subnet::Id

Description: Select a PUBLIC subnet (default VPC public subnet recommended)

VpcId:

Type: AWS::EC2::VPC::Id

Description: Select the VPC (default VPC recommended)

Resources:

WebServerRole:

Type: AWS::IAM::Role

Properties:

AssumeRolePolicyDocument:

Version: "2012-10-17"

Statement:

- Effect: Allow

Principal:

Service: ec2.amazonaws.com

Action: sts:AssumeRole

Policies:

- PolicyName: S3ReadOnlyForWebsiteBucket

PolicyDocument:

Version: "2012-10-17"

Statement:

- Effect: Allow

Action:

- s3:GetObject

- s3:ListBucket

Resource:

- !Sub "arn:aws:s3:::\${BucketName}"
- !Sub "arn:aws:s3:::\${BucketName}/*"

WebServerInstanceProfile:

Type: AWS::IAM::InstanceProfile

Properties:

Roles:

- !Ref WebServerRole

WebServerSG:

Type: AWS::EC2::SecurityGroup

Properties:

GroupDescription: Allow SSH (22) and HTTP (80)

VpcId: !Ref VpcId

SecurityGroupIngress:

- IpProtocol: tcp

FromPort: 22

ToPort: 22

CidrIp: 0.0.0.0/0

- IpProtocol: tcp

FromPort: 80

ToPort: 80

CidrIp: 0.0.0.0/0

WebServerInstance:

Type: AWS::EC2::Instance

Properties:

InstanceType: t2.micro

KeyName: !Ref KeyName

SubnetId: !Ref SubnetId

SecurityGroupIds:

- !Ref WebServerSG

IamInstanceProfile: !Ref WebServerInstanceProfile

Amazon Linux 2023 AMI for Mumbai (ap-south-1)

ImageId: ami-02b49a24cfb95941c

UserData:

Fn::Base64: !Sub |

#!/bin/bash

dnf update -y

dnf install -y httpd

```
systemctl enable httpd
systemctl start httpd
```

```
# pull index.html from S3 to Apache web root
aws s3 cp s3://${BucketName}/index.html /var/www/html/index.html
```

```
systemctl restart httpd
```

Outputs:

WebsiteURL:

Description: Open this URL in browser

Value: !Sub "http://\${WebServerInstance.PublicDnsName}"

Step 3.2: Create stack

CloudFormation → Stacks → Create stack

- Upload template → stack2-ec2-pull-from-s3.yaml
- Stack name: EC2-Pull-S3-Website
- Parameters:
 - **KeyName:** select your key pair (e.g., pemkeypair)
 - **BucketName:** paste bucket name from Stack-1 Output
 - **VpcId:** select **default VPC**
 - **SubnetId:** select a **PUBLIC subnet** in default VPC
(usually the subnet name shows “public” or you can pick any default subnet that gives public IP; default subnets generally work)

Create stack → wait for **CREATE_COMPLETE**

CloudFormation → Stacks → EC2-Pull-S3-Website → Operation: 3f24cc37-cbeb-44a0-8baf-1a956856b124

Operation type: CREATE_STACK


Operation events (15)

Search events ☐ Display failures only

Timestamp	Logical ID	Status	Detailed status	Status reason	Hook invocations
2026-01-14 15:21:35 UTC+0530	WebServerInstance	CREATE_COMPLETE	-	-	-
2026-01-14 15:21:24 UTC+0530	EC2-Pull-S3-Website	CREATE_IN_PROGRESS	CONFIGURATION_COMPLETE	Eventual consistency check initiated	-
2026-01-14 15:21:24 UTC+0530	WebServerInstance	CREATE_IN_PROGRESS	CONFIGURATION_COMPLETE	Eventual consistency check initiated	-
2026-01-14 15:21:21 UTC+0530	WebServerInstance	CREATE_IN_PROGRESS	-	Resource creation initiated	-
2026-01-14 15:20:19 UTC+0530	WebServerInstance	CREATE_IN_PROGRESS	-	-	-
2026-01-14 15:20:18 UTC+0530	WebServerInstanceProfile	CREATE_IN_PROGRESS	CONFIGURATION_COMPLETE	Eventual consistency check initiated	-
2026-01-14 15:20:17 UTC+0530	WebServerInstanceProfile	CREATE_IN_PROGRESS	-	Resource creation initiated	-
2026-01-14 15:20:16 UTC+0530	WebServerInstanceProfile	CREATE_IN_PROGRESS	-	-	-
2026-01-14 15:20:16 UTC+0530	WebServerRole	CREATE_COMPLETE	-	-	-
2026-01-14 15:20:07 UTC+0530	WebServerSG	CREATE_COMPLETE	-	-	-
2026-01-14 15:20:00 UTC+0530	WebServerSG	CREATE_IN_PROGRESS	-	Resource creation initiated	-
2026-01-14 15:19:59 UTC+0530	WebServerSG	CREATE_IN_PROGRESS	-	Resource creation initiated	-

PART 4 — Verify Output

Stack-2 → Outputs → copy WebsiteURL → open in browser.



Hello! This page was pulled from S3 to EC2 automatically.

Deployed using CloudFormation + UserData

Troubleshooting (common)

1) Website not opening

- Wait 1–2 minutes (UserData takes time)
- Ensure Security Group has **HTTP 80 open**
- Ensure you selected a **public subnet** (needs internet to reach S3)

2) Stack-2 fails with S3 access error

- BucketName typed wrong
- index.html not uploaded at root of bucket

Clean-up

Delete Stack-2 first

CloudFormation → select EC2-Pull-S3-Website → Delete

Empty the S3 bucket

S3 → bucket → delete index.html (empty bucket)

Delete Stack-1

CloudFormation → select S3-Website-Bucket → Delete

DATE: 19-12-25

Exercise–24: AWS Lambda: Input Processing, Business Logic Execution, and CloudWatch Logging

Create one Lambda function that:

1. prints a welcome message
2. reads JSON input event
3. performs business logic (grade calculation)
4. writes logs → view them in CloudWatch Logs

Yes — this is the **first combined Lambda lab exercise** (basics + input + logic + logging). Below is a **complete step-by-step procedure** using **AWS Console**.

Create the Lambda Function

Step 1: Open Lambda

AWS Console → Search Lambda → Open AWS Lambda

Step 2: Create function

Click Create function

- Select - Author from scratch
- Function name: StudentGradeLogger
- Runtime: Python 3.11
- Architecture: x86_64 (default)


Step 3: Permissions (Execution Role)

Under “Permissions”

- Choose: Create a new role with basic Lambda permissions

This automatically gives permission to write logs to CloudWatch.

Click **Create function**



Add Code (Business Logic + Logs)

Step 4: Paste this code

In **Code** tab → lambda_function.py → paste and **Deploy**

```
import json

def lambda_handler(event, context):
    # Welcome message
    print("Lambda invoked successfully")

    # Read input from event
    student_name = event["StudentName"]
    marks = event["Marks"]

    # Business logic: grade calculation
    if marks >= 75:
        result = "Pass"
        grade = "A"
    else:
        result = "Fail"
        grade = "F"

    # Log output
    print("Student:", student_name)
    print("Marks:", marks)
    print("Grade:", grade)
    print("Result:", result)

    # Return response
    return {
        "statusCode": 200,
        "body": json.dumps({
            "StudentName": student_name,
            "Marks": marks,
            "Grade": grade,
            "Result": result
        })
    }
```

Click **Deploy**.

Test with Input Events (JSON)

Step 5: Create a test event (Valid case)

Go to Test (top right) → Configure test event

- Event name: ValidInput
- Paste this JSON:

```
{
  "StudentName": "Anita",
  "Marks": 86
}
```


Click **Save**

Step 6: Run test

Click **Test**

Expected response (sample):

- statusCode: 200
- body will include Grade A and Result Pass



Step 7: Test invalid input (Missing Marks)

Create another test event:


- Event name: MissingMarks

```
{
  "StudentName": "Rahul"
}
```

Run Test

Expected:

- statusCode: 400
- message: Missing 'Marks'



Step 8: Test invalid marks range

Event name: InvalidMarks

```
{
  "StudentName": "Priya",
  "Marks": 150
}
```

Expected:

- statusCode: 400
- message: Marks must be 0–100

View Logs in CloudWatch

Step 9: Open logs from Lambda directly

On the Lambda function page:

- Go to **Monitor** tab
- Click **View CloudWatch logs**

Log events

You can use the filter bar below to search for and match terms, phrases, or values in your log events. [Learn more about filter patterns](#)

Filter events - press enter to search

Clear 1m 30m 1h 12h Custom UTC timezone Display

Timestamp	Message
2026-01-09T04:55:59.153Z	Grade: A
2026-01-09T04:55:59.153Z	Result: Pass
2026-01-09T04:55:59.155Z	END RequestId: c05aacdd-0d14-4a8e-a08c-b17eefa71ed8
2026-01-09T04:55:59.155Z	REPORT RequestId: c05aacdd-0d14-4a8e-a08c-b17eefa71ed8 Duration: 1.22 ms Billed Duration: 2 ms Memory Size: 128 MB Max Memory Used: 34 MB
2026-01-09T04:56:42.020Z	START RequestId: 0fc8ade8-6fed-42ed-a548-e1e06d106d97 Version: \$LATEST
2026-01-09T04:56:42.021Z	[ERROR] KeyError: "Marks" Traceback (most recent call last): File "/var/task/lambda_function.py", line 9, in lambda_handler marks = event["Marks"]
2026-01-09T04:56:42.025Z	Lambda invoked successfully
2026-01-09T04:56:42.047Z	END RequestId: 0fc8ade8-6fed-42ed-a548-e1e06d106d97
2026-01-09T04:56:42.047Z	REPORT RequestId: 0fc8ade8-6fed-42ed-a548-e1e06d106d97 Duration: 25.49 ms Billed Duration: 26 ms Memory Size: 128 MB Max Memory Used: 34 MB
2026-01-09T04:57:10.071Z	START RequestId: b58ab002-09d1-4531-9c73-4e11dca37332 Version: \$LATEST
2026-01-09T04:57:10.071Z	Lambda invoked successfully
2026-01-09T04:57:10.071Z	Student: Priya
2026-01-09T04:57:10.071Z	Marks: 150
2026-01-09T04:57:10.071Z	Grade: A
2026-01-09T04:57:10.071Z	Result: Pass
2026-01-09T04:57:10.073Z	END RequestId: b58ab002-09d1-4531-9c73-4e11dca37332
2026-01-09T04:57:10.073Z	REPORT RequestId: b58ab002-09d1-4531-9c73-4e11dca37332 Duration: 1.26 ms Billed Duration: 2 ms Memory Size: 128 MB Max Memory Used: 34 MB

No newer events at this moment. Auto retry paused. [Resume](#)

Back to top

CloudShell Feedback Console Mobile App © 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

You will see:

- Log group: /aws/lambda/StudentGradeLogger
- Open latest **log stream**
- You should see all print() outputs:
 - “Lambda invoked successfully!”
 - “Received event...”
 - grade/result logs
 - error logs for invalid tests

Cleanup (Avoid charges, keep account clean)

Lambda itself usually costs nothing in small use, but cleanup is good practice:

1. Lambda → Functions → select StudentGradeLogger → **Delete**
2. CloudWatch → Log groups → find /aws/lambda/StudentGradeLogger → **Delete**
3. IAM role created (optional):
 - IAM → Roles → search role name created for Lambda → delete (only if not used elsewhere)

DATE: 24-12-2025

Exercise-25 Mini Project

Event-Driven Notification System using AWS Lambda and Amazon SNS

Simulating real-time alerts from events using serverless computing

- An event in JSON format is given as input to an AWS Lambda function.
- The Lambda function processes the event and generates a notification message.
- The message is published to an Amazon SNS topic.
- SNS sends the notification to the subscribed email address.
- The execution details are verified using CloudWatch Logs.

Create SNS Topic + Email Subscription

Step 1: Open SNS

AWS Console → Search SNS → Open Simple Notification Service

Step 2: Create a Topic

SNS → Topics → Create topic

- Type: Standard
 - Name: NotificationTopic
- Click Create topic

Step 3: Copy Topic ARN

Open the created topic → copy Topic ARN


(You will paste it into Lambda code later)

Step 4: Create an Email Subscription

Inside the same topic:

Click Create subscription

- Protocol: Email
 - Endpoint: (your email id)
- Click Create subscription



Step 5: Confirm Subscription

Open your email inbox → find AWS SNS confirmation mail → click Confirm subscription
→ Status in SNS should become Confirmed.

Create Lambda Function


Step 6: Open Lambda

AWS Console → Search Lambda → Open AWS Lambda

Step 7: Create function

Click Create function

- Select: Author from scratch
 - Function name: NotificationSimulator
 - Runtime: Python 3.11
 - Permissions: Create a new role with basic Lambda permissions
- Click Create function



Add Lambda Code (with SNS Publish)

Step 8: Paste code in lambda_function.py

Lambda → **Code** tab → open lambda_function.py → remove existing code → paste:

Replace <SNS_TOPIC_ARN> with your copied Topic ARN.

```
import json
import boto3

sns = boto3.client('sns')

TOPIC_ARN = "<SNS_TOPIC_ARN>"
```



```
def lambda_handler(event, context):
    print("Notification Simulator invoked")

    event_type = event["eventType"]
    user = event["user"]

    if event_type == "ORDER_PLACED":
        message = f"Hi {user}, your order is placed successfully."
        priority = "NORMAL"

    elif event_type == "PAYMENT_FAILED":
        message = f"Hi {user}, your payment has failed. Please retry."
        priority = "HIGH"

    elif event_type == "LOW_ATTENDANCE":
        message = f"Hi {user}, your attendance is low. Please take action."
        priority = "HIGH"

    else:
        message = f"Hi {user}, unknown event received."
        priority = "LOW"

    print("Event Type:", event_type)
    print("Message:", message)
    print("Priority:", priority)

    # Publish to SNS (Actual notification delivery)
    sns.publish(TopicArn=TOPIC_ARN, Message=message, Subject=f"{event_type} [{priority}]")

    return {
        "statusCode": 200,
        "body": json.dumps({
            "EventType": event_type,
            "Message": message,
            "Priority": priority
        })
    }
```

Click **Deploy**

Wait for “Successfully updated function...”

Give Lambda Permission to Publish to SNS

Step 10: Open Lambda Execution Role

Lambda → Configuration → Permissions

Click the Role name (execution role link)

Step 11: Attach SNS Publish Policy

IAM Role page → Add permissions → Attach policies

Attach: AmazonSNSFullAccess

Click Add permissions

Now Lambda can publish to SNS.

Test the Mini Project

Step 12: Create a test event


Lambda → Test tab → Create new test event

Event name: PaymentFailed

Paste:

```
{
  "eventType": "PAYMENT_FAILED",
  "user": "Anita"
}
```

Click **Save**



Step 13: Run Test

Click **Test**

Expected:

- Execution: **Succeeded**

- Email should arrive within a few seconds:
Subject like: PAYMENT_FAILED [HIGH]
Message: “Hi Anita, your payment has failed. Please retry.”

Verify CloudWatch Logs

Step 14: View logs

Lambda → Monitor → View CloudWatch logs

Open latest log stream.

Verify these prints exist:

- Notification Simulator invoked
- Event Type:
- Message:
- Priority:

▶	Timestamp	Message
		No older events at this moment. Retry
▶	2026-01-14T06:23:48.804Z	INIT_START Runtime Version: python:3.11.v109 Runtime Version ARN: arn:aws:lambda:ap-northeast-1::runtime:49f733259c7ce7e0deee75ff91c6...
▶	2026-01-14T06:23:49.250Z	START RequestId: 4810bbfd-a583-4aab-a86f-357e148ac3dc Version: \$LATEST
▶	2026-01-14T06:23:49.250Z	Notification Simulator invoked
▶	2026-01-14T06:23:49.250Z	Event Type: PAYMENT_FAILED
▶	2026-01-14T06:23:49.250Z	Message: Hi Anita, your payment has failed. Please retry.
▶	2026-01-14T06:23:49.250Z	Priority: HIGH
▶	2026-01-14T06:23:49.529Z	END RequestId: 4810bbfd-a583-4aab-a86f-357e148ac3dc
▶	2026-01-14T06:23:49.529Z	REPORT RequestId: 4810bbfd-a583-4aab-a86f-357e148ac3dc Duration: 261.71 ms Billed Duration: 704 ms Memory Size: 128 MB Max Memory Use...
		No newer events at this moment. <i>Auto retry paused.</i> Resume

Expected Outputs

1. Lambda test status: Succeeded
2. Email received from SNS with correct message
3. CloudWatch logs showing event type and generated message

Cleanup

1. Delete Lambda function: NotificationSimulator
2. SNS → delete topic NotificationTopic (subscriptions removed automatically)
3. CloudWatch → delete log group for Lambda
4. (Optional) Delete IAM role if not used elsewhere

DATE: 26-12-25

Exercise 26: Analyze a CSV File in S3 Using Amazon Athena (No Server)

Upload a small CSV to S3, then use **Athena** to run SQL queries like **count, average, max, group by**.

Part A — Create Sample CSV (on your PC)

1. Open **Notepad**
2. Paste this data and save as: **students.csv**

StudentID,Name,Dept,Marks,Result

101,Anita,MCA,85,Pass

102,Ravi,MCA,72,Pass

103,Meera,MBA,64,Pass

104,John,MCA,35,Fail

105,Sneha,MBA,91,Pass

106,Arun,MCA,49,Fail

107,Kiran,MBA,58,Pass

108,Divya,MCA,77,Pass

Part B — Create S3 Bucket and Upload CSV

1. Go to **AWS Console** → **S3**
2. Click **Create bucket**
3. Bucket name: athena-lab-<yourname>-<number> (must be globally unique)
4. Keep defaults → Click **Create bucket**
5. Open the bucket → Click **Upload**
6. Upload **students.csv**
7. Click **Upload**

Now your CSV is in S3.

Part C — Open Athena and Set Query Result Location (IMPORTANT)

1. Go to **AWS Console** → **Amazon Athena**
2. If it shows “Get started” / “Query editor”, open it.
3. It will ask to set a **Query result location**
4. Click the link/button like **Settings / Manage / Edit**
5. Set query result location to something like:
 - s3://your-bucket-name/athena-results/
6. Click **Save**

This is required so Athena can store query outputs.

Part D — Create a Database in Athena

In Athena query editor, run:

```
CREATE DATABASE labdb;
```

Then on the left side, select:

- **Data source:** AwsDataCatalog
- **Database:** labdb

Part E — Create a Table for the CSV in S3

Run this (replace YOUR_BUCKET_NAME with your actual bucket name):

```
CREATE EXTERNAL TABLE IF NOT EXISTS students (  
  StudentID int,  
  Name string,  
  Dept string,  
  Marks int,  
  Result string  
)  
ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'  
WITH SERDEPROPERTIES (
```

```
'separatorChar' = ',',
'quoteChar'   = '"',
'escapeChar'  = '\\'
)
```

STORED AS TEXTFILE

LOCATION 's3://YOUR_BUCKET_NAME/'

TBLPROPERTIES ('skip.header.line.count'='1');


This tells Athena: “My CSV is in S3, treat it like a table.”

Part F — Run Simple Analysis Queries (Core Part)

1) View all rows

SELECT * FROM students;

Expected: 8 records.




The screenshot shows the Amazon Athena Query Editor interface. The query editor displays the SQL query `SELECT * FROM students;`. The query results are shown in a table with 8 rows and 7 columns: `studentid`, `name`, `dept`, `marks`, and `result`. The results are as follows:

#	studentid	name	dept	marks	result
1	101	Anita	MCA	85	Pass
2	102	Ravi	MCA	72	Pass
3	103	Meera	MBA	64	Pass
4	104	John	MCA	35	Fail
5	105	Sneha	MBA	91	Pass
6	106	Arun	MCA	49	Fail

2) Total students

SELECT COUNT(*) AS total_students FROM students;

Expected: 8



The screenshot shows the Amazon Athena Query Editor interface. The query editor displays the SQL query `SELECT COUNT(*) AS total_students FROM students;`. The query results are shown in a table with 1 row and 1 column: `total_students`. The result is as follows:

#	total_students
1	8


3) Pass vs Fail count

SELECT Result, COUNT(*) AS cnt

FROM students

GROUP BY Result;

Expected (based on data): Pass = 6, Fail = 2



4) Average marks overall

SELECT AVG(Marks) AS avg_marks

FROM students;

5) Department-wise average marks

SELECT Dept, AVG(Marks) AS avg_marks

FROM students

GROUP BY Dept;

6) Top scorer

SELECT Name, Dept, Marks

FROM students

ORDER BY Marks DESC

LIMIT 1;

Expected: Sneha (91)


7) List failed students

SELECT StudentID, Name, Dept, Marks

FROM students

WHERE Result = 'Fail';

Expected: John, Arun



Cleanup (to avoid any charges)

1. In Athena, you can keep DB/table (no cost by itself), but clean S3:
2. Go to **S3 bucket**
3. Delete:
 - students.csv
 - athena-results/ folder contents (query outputs)
4. Delete the bucket (must be empty to delete)

Cost Note (Simple)

- **S3 storage:** tiny (almost negligible for this)
- **Athena:** charges based on **data scanned**; with this tiny CSV it's usually minimal, but **always delete outputs and bucket** after lab.

DATE: 31-12-25

Exercise 27: Cloud-Based IoT Data Ingestion and Processing using AWS

Cloud-Based IoT Data Ingestion and Processing using AWS - To understand how IoT device data is sent to the cloud, processed automatically, and logged using AWS managed services.


In this lab, a simulated IoT device sends sensor data (such as temperature and humidity) in JSON format to the AWS cloud.

The data is received by AWS IoT Core, which acts as the central message broker for IoT devices.

Whenever new sensor data arrives:

- AWS IoT Core detects the event
- The data is processed automatically
- The processed output is stored or logged in the cloud for monitoring and verification

This exercise demonstrates how real-time device data can be handled without managing servers, highlighting the principles of IoT, cloud computing, and event-driven architecture.



Concepts Covered

- Internet of Things (IoT)
- Device-to-Cloud communication
- Event-driven processing
- JSON data format
- Serverless cloud services

- Real-time data handling

Input

Simulated sensor data (example: temperature and humidity values)

Output

- IoT message successfully received by AWS
- Automatic processing triggered
- Cloud logs showing received and processed data

Learning Outcome

After completing this lab, students will be able to:

- Explain how IoT devices communicate with the cloud
- Understand the role of AWS IoT Core in IoT solutions
- Describe event-based data processing in cloud environments
- Relate IoT concepts to real-world smart systems (smart homes, healthcare, industry)

STEP 1: Open IoT Service

0. Core dashboard

STEP 2: Create an IoT Thing (Device)

1. Go to Manage → Things
2. Click Create things
3. Choose Create single thing
4. Give a name (example: TempSensor01)
5. Skip advanced settings
6. Create the thing

This represents a sensor device (even though no real hardware is used).

STEP 3: Create Device Certificate

1. Choose Auto-generate certificate
2. Activate the certificate
3. Download:
 - Device certificate
 - Private key
 - Root CA
4. Attach the certificate to the Thing

Certificate = identity of the device

The screenshot shows the AWS IoT console interface for a device named 'TempSensor01'. The breadcrumb navigation at the top indicates the path: AWS IoT > Manage > Things > TempSensor01. The device's connectivity indexing is not enabled. Action buttons for 'Create secure tunnel', 'Edit', and 'Delete' are visible. The 'Thing details' section shows the name 'TempSensor01', its ARN, and its type and billing group. The 'Certificates' tab is active, displaying a table with one certificate. The certificate is active and associated with the device via a 'Non-Exclusive' association. The footer of the console shows the page number '3' and links to CloudShell, Feedback, and the Console Mobile App, along with a copyright notice for Amazon Web Services, Inc.

Certificate ID	Status	Certificate to thing association
3e3802db39f2e8c25c59296bec7d45aaa8e4b2a86482a1ca04d6934ab91e351f	Active	Non-Exclusive

STEP 4: Create and Attach IoT Policy

1. Go to Secure → Policies
2. Create a new policy
3. Allow:
 - Connect
 - Publish

- Subscribe
 - Receive
4. Use * (for lab simplicity)
 5. Attach the policy to the certificate

Policy = permission for device to talk to AWS

The screenshot shows the AWS IoT console interface for a policy named 'TempSensorPolicy'. The breadcrumb navigation at the top is 'AWS IoT > Security > Policies > TempSensorPolicy'. On the right, there are buttons for 'Edit active version' and 'Delete'. The 'Details' section shows the Policy ARN as 'arn:aws:iot:ap-northeast-1:49949767753:policy/TempSensorPolicy', the Active version as '1', and both 'Created' and 'Last updated' timestamps as 'January 16, 2026, 10:52:27 (UTC+05:30)'. Below this, the 'Versions' tab is selected, showing 'Active version: 1' with a 'Builder' button and a 'JSON' link. A table lists the policy actions: 'iot:Connect', 'iot:Publish', 'iot:Subscribe', and 'iot:Receive', all with an 'Allow' effect and a '*' resource. At the bottom, the 'All versions (1)' section shows a single version with buttons for 'Delete', 'Set as active', 'Edit version', and 'View JSON'. The footer includes 'CloudShell', 'Feedback', 'Console Mobile App', and copyright information for Amazon Web Services, Inc. or its affiliates.

STEP 5: Test Device Messages (No Hardware)

Go to Test → MQTT test client

Subscribe to a topic

Example: sensor/temperature

Publish a message:

```
{
  "deviceId": "TempSensor01",
  "temperature": 32,
  "humidity": 65
}
```


Verify message appears instantly

IoT data successfully reached the cloud

STEP 6: Create a Rule (Automation)

1. Go to Act → Rules
2. Create a rule
3. Rule query example:
4. `SELECT * FROM 'sensor/temperature'`
5. Choose action:
 - Send to CloudWatch Logs
(or *Lambda / DynamoDB if required*)

Rule = “when data arrives, do something”



STEP 7: Verify Output

1. Open CloudWatch → Logs
2. Check log group created by IoT Rule
3. Confirm sensor data entries

Automatic processing confirmed

CloudWatch > Log management > IoT > IoT_Rule-470603366

Log events

You can use the filter bar below to search for and match terms, phrases, or values in your log events. [Learn more about filter patterns](#)

Filter events - press enter to search

Clear 1m 30m 1h 12h Custom UTC timezone Display

Timestamp	Message
	No older events at this moment. Retry
2026-01-16T05:49:15.625Z	{ "message": "Hello from AWS IoT console" }
	No newer events at this moment. <i>Auto retry paused.</i> Resume

Sensor data is sent to AWS IoT Core, where rules automatically process and store the data without using any server.

DATE: 02-01-26

Exercise–28: Accelerate an S3 Static Website Using Amazon CloudFront (CDN)


Pre-requisite

- You already have an **S3 bucket** with at least:
 - index.html
 - (optional) error.html

Part A - Ensure S3 is ready

- Open **S3** → **your bucket**
- Go to **Properties** → **Static website hosting**
- Enable it, set:
 - Index document: index.html
 - Error document: error.html (optional)

Note down the S3 Website endpoint shown there.



The screenshot shows the Amazon S3 console interface. The breadcrumb navigation at the top indicates the path: Amazon S3 > Buckets > cdn-exercise-28. The main header shows the bucket name 'cdn-exercise-28' with an 'info' link. Below this, there are tabs for 'Objects', 'Metadata', 'Properties', 'Permissions', 'Metrics', 'Management', and 'Access Points'. The 'Properties' tab is selected. In the 'Properties' section, the 'Static website hosting' option is enabled. It shows the 'Index document' set to 'index.html' and the 'Error document' set to 'error.html'. The 'Website endpoint' is listed as 's3-website-us-east-1-111111111111.s3.amazonaws.com'. Below the properties, there is a section for 'Objects (2)' showing a list of files: 'error.html' (1.0 KB) and 'index.html' (782.0 B), both with a storage class of 'Standard'.

Part B - Create CloudFront Distribution

- Go to **CloudFront** → **Create distribution**
- **Origin domain**
 - Select your **S3 bucket** (not the website endpoint)

- **Origin access**
 - Choose **Origin access control (OAC)** (recommended)
 - Click **Create control setting** if asked
- **Viewer protocol policy**
 - Choose **Redirect HTTP to HTTPS**
- **Default root object**
 - Set: index.html
- Click **Create distribution**


Copy the **Distribution domain name** (looks like dxxxxx.cloudfront.net)

Part C - Allow CloudFront to read the bucket (important)

After creating distribution, CloudFront usually shows a message like “**S3 bucket policy needs update**”.

- Click **Copy policy**
- Go to **S3 → Bucket → Permissions → Bucket policy**
- Paste and **Save**

This step ensures **public users cannot directly access S3**, but CloudFront can.




Part D - Test

- Wait until Distribution status becomes **Enabled** (and “Deployed”)
- Open in browser:
 - <https://dxxxxx.cloudfront.net>

Expected:

- Your index.html loads via CloudFront.



Viva

- **S3** stores the website files
- **CloudFront** caches them at edge locations (faster)
- **OAC** ensures S3 is not public; CloudFront accesses it securely

DATE: 02-01-26

Exercise–29 Mini Project – Global Static Website Delivery using S3 + CloudFront with Cache Update Demo

Host a static website in S3, deliver it through CloudFront (CDN), then do a content update and observe caching behavior (old page still showing). Finally, force the latest version using CloudFront Invalidation.

A) Prerequisites

- AWS Account
- An S3 bucket
- 2 small HTML files ready:
 - index.html
 - about.html

index.html (Version 1)

```
<html>

<body>

  <h1>My CloudFront Mini Project</h1>

  <h2>Version: V1</h2>

  <p>Loaded via CloudFront CDN</p>

</body>

</html>
```

about.html

```
<html>

<body>

  <h1>About Page</h1>

  <h2>Version: V1</h2>

</body>

</html>
```

B) Step-by-step Execution

Step 1: Create S3 Bucket

Go to S3 → Create bucket

Bucket name: cf-mini-project-<yourname>

Region: (keep default)

Block Public Access: Keep ON

Click Create bucket


Step 2: Upload Website Files

Open your bucket → Upload

Upload:

0. index.html
1. about.html

Click Upload



Step 3: Create CloudFront Distribution (with secure access)

Go to CloudFront → Create distribution

Origin domain

- Select your S3 bucket (from dropdown)

Origin access

- Choose Origin access control (OAC)
- Click Create control setting (if asked) → Create

Viewer protocol policy

- Select Redirect HTTP to HTTPS

Default root object

- Enter: index.html

Click Create distribution

Step 4: Add Bucket Policy for OAC (important)

After distribution creation, CloudFront page will show a message like:

- “S3 bucket policy needs to be updated”

Click Copy policy

Go to S3 → your bucket → Permissions → Bucket policy

Paste policy → Save changes

Now: S3 is NOT public, but CloudFront can fetch objects.

Step 5: Test the Website Using CloudFront

Go back to CloudFront


Open your distribution

Copy the Distribution domain name, like:

- dxxxxxx.cloudfront.net

Open in browser:

- <https://dxxxxxx.cloudfront.net>



My CloudFront Mini Project

Version: V1


Loaded via CloudFront CDN

Also test:

- <https://dxxxxxx.cloudfront.net/about.html>

Screenshot evidence:

- CloudFront distribution details page (domain + status)
- Browser output showing **Version V1**



C) Cache Demo (Version Update + See Old Content)

Step 6: Update Website Content in S3 (V2)

Edit your index.html and change:

From:

```
<h2>Version: V1</h2>
```

To:

```
<h2>Version: V2</h2>
```

Now re-upload updated index.html:

S3 → bucket → Upload → add updated index.html

Upload

Step 7: Observe CloudFront Caching

Immediately refresh:

- <https://dxxxxxx.cloudfront.net>

Many times you may still see **V1** (cached), even though S3 has V2.

This is the caching behavior demonstration.

Tip for students:

- Do a hard refresh: Ctrl + Shift + R
- Still might show V1 because cache is on CloudFront edge.

Take screenshot if it still shows V1.

D) Force Latest Content Using Invalidation

Step 8: Create CloudFront Invalidation

CloudFront → your distribution

Go to Invalidation tab

Click Create invalidation

In “Object paths”, enter:

- /index.html
(or use /* to clear everything)

Click **Create invalidation**

Wait until status becomes **Completed**.


Step 9: Verify Latest Version

Refresh the CloudFront URL again: <https://dxxxxx.cloudfront.net>

Now it should show **Version: V2**

Screenshot evidence:

- Invalidation status “Completed”
- Browser showing V2



About Page

Version: V2

Cleanup (to avoid charges)

1. CloudFront: Disable distribution → wait → Delete distribution

2. S3: Empty bucket → Delete bucket

Mini-Project Output Checklist:

S3 bucket file list (index.html, about.html)

CloudFront distribution domain name

Browser output **V1**

S3 updated file showing **V2**

Browser still showing **V1** (optional if observed)

CloudFront invalidation completed

Browser output **V2**

Viva Questions

1. Why use CloudFront instead of accessing S3 directly?
2. What does “cache” mean?
3. Why was invalidation needed?
4. What is OAC and why is it useful?

DATE: 07-01-26

Exercise–30: Mini Project – Deploying a Containerized App on AWS using Amazon EKS

- Create an EKS cluster on AWS
- Add worker nodes (managed node group)
- Connect using kubectl
- Deploy Nginx and expose it

Region

Use **Mumbai (ap-south-1)** (or keep consistent with your account).

Prerequisites

- Logged in with root email + password
- **AWS CloudShell**

Method: CloudShell + eksctl

Step A1: Open CloudShell

AWS Console (Mumbai region) → click **CloudShell** icon (terminal opens in browser)

Step A2: Check tools


Run:

aws --version

kubectl version --client

eksctl version

If eksctl is not found, install it (CloudShell often has it).



```
CloudShell
ap-northeast-1 +
~ $ aws --version
aws-cli/2.32.34 Python/3.13.11 Linux/6.1.158-180.294.amzn2023.x86_64 exec-env/CloudShell exe/x86_64.amzn.2023
~ $ kubectl version --client
Client Version: v1.34.2-eks-ecaa3a6
Kustomize Version: v5.7.1
~ $ eksctl version
0.221.0
~ $
```

Step A3: Create EKS cluster + node group (single command)

Run:

```
eksctl create cluster \
  --name eks-lab-cluster \
  --region ap-south-1 \
  --nodes 2 \
  --node-type t3.medium \
  --managed
```

If eksctl asks / fails due to IAM (rare, but in case)

Run: `aws sts get-caller-identity`

If it shows:


- Your **root account ID**
- ARN ending with `:root`

Then IAM is 100% fine.

This automatically:

- Creates a **new VPC** with correct subnets/routes
- Creates the **EKS cluster**
- Creates the **Managed Node Group (worker nodes)**
- Configures kubeconfig for you

Wait until it completes (it will print “cluster created”)



Step A4: Verify nodes

`kubectl get nodes`

Expected: 2 nodes in **Ready** state.

Step A5: Deploy Nginx

`kubectl create deployment webapp --image=nginx`

`kubectl get pods`

Step A6: Expose using LoadBalancer

`kubectl expose deployment webapp --type=LoadBalancer --port=80`

`kubectl get svc`

Wait until EXTERNAL-IP becomes a value (not <pending>).

Step A7: Open in browser

Copy the EXTERNAL-IP and open:

http://<EXTERNAL-IP>

Expected: Nginx welcome page

CloudShell

Actions

ap-south-1 X ap-south-1 X +

```
No resources found in default namespace.
~ $
~ $ kubectl create deployment webapp --image=nginx
deployment.apps/webapp created
~ $ kubectl expose deployment webapp --type=LoadBalancer --port=80
service/webapp exposed
~ $ kubectl get pods

NAME                READY   STATUS    RESTARTS   AGE
webapp-c89577d45-2c2z5 0/1     Pending   0           9s
~ $
~ $ kubectl get pods

NAME                READY   STATUS    RESTARTS   AGE
webapp-c89577d45-2c2z5 0/1     Pending   0           13s
~ $ kubectl get pods

NAME                READY   STATUS    RESTARTS   AGE
webapp-c89577d45-2c2z5 0/1     Pending   0           17s
~ $ kubectl get pods

NAME                READY   STATUS    RESTARTS   AGE
webapp-c89577d45-2c2z5 0/1     Pending   0           28s
~ $ kubectl delete deployment webapp
deployment.apps "webapp" deleted from default namespace
~ $ kubectl create deployment webapp --image=nginx
deployment.apps/webapp created
~ $ kubectl delete deployment webapp
deployment.apps "webapp" deleted from default namespace
~ $ kubectl delete service webapp
service "webapp" deleted from default namespace
~ $ kubectl get pods
No resources found in default namespace.
~ $ kubectl get services

NAME                TYPE                CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
kubernetes          ClusterIP           10.100.0.1    <none>         443/TCP    41m
~ $ kubectl create deployment webapp --image=nginx
deployment.apps/webapp created
~ $ kubectl get deployments

NAME    READY   UP-TO-DATE   AVAILABLE   AGE
webapp  0/1     1            0           10s
~ $ kubectl expose deployment webapp --type=LoadBalancer --port=80
service/webapp exposed
~ $ kubectl get svc -w
```

Feedback Console Mobile App

© 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

CloudShell

Actions

ap-south-1 X ap-south-1 X +

```
~ $ kubectl get services

NAME                TYPE                CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
kubernetes          ClusterIP           10.100.0.1    <none>         443/TCP    41m
~ $ kubectl create deployment webapp --image=nginx
deployment.apps/webapp created
~ $ kubectl get deployments

NAME    READY   UP-TO-DATE   AVAILABLE   AGE
webapp  0/1     1            0           10s
~ $ kubectl expose deployment webapp --type=LoadBalancer --port=80
service/webapp exposed
~ $ kubectl get svc -w

NAME                TYPE                CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
kubernetes          ClusterIP           10.100.0.1    <none>         443/TCP    43m
webapp              LoadBalancer       10.100.192.219 <pending>      80:31531/TCP 8s
^C~ $ kubectl port-forward deployment/webapp 8080:80
Forwarding from 127.0.0.1:8080 -> 80
^C~ $ ^C
~ $ ^C
~ $ kubectl delete service webapp
service "webapp" deleted from default namespace
~ $ kubectl expose deployment webapp --type=NodePort --port=80
service/webapp exposed
~ $ kubectl get svc

NAME                TYPE                CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
kubernetes          ClusterIP           10.100.0.1    <none>         443/TCP    45m
webapp              NodePort            10.100.146.24 <none>        80:32646/TCP 8s
~ $
~ $ kubectl get svc

NAME                TYPE                CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
kubernetes          ClusterIP           10.100.0.1    <none>         443/TCP    45m
webapp              NodePort            10.100.146.24 <none>        80:32646/TCP 12s
~ $
~ $ kubectl get nodes -o wide

NAME                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE                                     KERNEL-VERSION   CONTAINER-RUNTIME
i-0227cb7c13d4b6d77 Ready    <none>   35m   v1.34.2-eks-b3126f4 172.31.41.207 13.234.32.123 Bottlerocket (EKS Auto, Standard) 2025.12.18 (aws-k8s-1.34-standard) 6.12.58          containerd://2.1.5+
i-054f4becbb867f1bb Ready    <none>   26m   v1.34.2-eks-b3126f4 172.31.7.85  13.203.157.75 Bottlerocket (EKS Auto, Standard) 2025.12.18 (aws-k8s-1.34-standard) 6.12.58          containerd://2.1.5+
bottlerocket
~ $ kubectl port-forward deployment/webapp 8080:80
Forwarding from 127.0.0.1:8080 -> 80
Handling connection for 8080
^_ ^_ ^_ ^_
```

Feedback Console Mobile App

© 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences